

CASE BRIEFING • CLASS 2

DBS' AI Journey

Zhu, F., Zhu, H. & Wong, A. (2024, rev. 2026). Harvard Business School, case 625-053.

Read this before class. You don't need the full Harvard case beforehand — we brief it live in Session 2 and release the case study afterwards. But the timeline, facts and decisions below are what your team will debate, so skim them first; coming in cold will slow you down. The discussion questions are at the end.

THE SITUATION

May 2024. Piyush Gupta, CEO of DBS — Singapore's largest bank — has just watched OpenAI's demo of GPT-4o, days after DBS posted record first-quarter results that its own leaders credit partly to AI. DBS is an early adopter of generative AI (GenAI) and already sees promising internal results. But the technology is nascent: to scale it safely, Piyush must decide which use cases to pursue and whether the bank's decade-old AI governance can keep pace. Ahead of a GenAI Steering Committee meeting, he weighs how to steer DBS toward GenAI more safely and responsibly — and how far to let it go.

CHRONOLOGY — DBS'S AI JOURNEY

- **1968** — Founded as the Development Bank of Singapore to finance the nation's industrialisation; grows into one of Southeast Asia's largest banks.
- **Early 2000s** — Poor customer-satisfaction scores; nicknamed "Damn Bloody Slow."
- **2009** — Piyush Gupta becomes CEO and launches a deep digital transformation.
- **2012** — A meeting with Alibaba/Alipay's Jack Ma convinces Piyush that fintechs are reimagining banking. DBS accelerates around three pillars — 'digital to the core' (cloud + open source + APIs), customer-obsession, and a start-up culture — and coins "GANDALF" (benchmarking against Google, Amazon, Netflix, LinkedIn, Facebook, with DBS the 'D').
- **2014** — First real AI bet: IBM Watson for real-time wealth advice — abandoned after 18 months (~80% accuracy, "not good enough to take to market"). DBS also opens a joint AI lab with Singapore's A*STAR.
- **2016** — Invests in Kasisto to build Asian-language NLP; chatbot errors push DBS to safer "guided conversations."
- **2018** — Consolidates 12 data warehouses into ADA (Advancing DBS with AI) — one 'source of truth' with tiered data and role-based, auto-approved access (what took months now takes seconds). Introduces the P-U-R-E responsible-data framework and helps shape the regulator's FEAT principles. Adopts 'two-in-a-box' (business + tech co-own projects).
- **2020** — Doubles down on machine learning over structured data; formalises the RDU (Responsible Data Use) model — regulatory alignment + P-U-R-E + model governance — run by an RDU Committee that 'white-lists' approved use cases.
- **2021** — "Managing through Journeys": cross-functional teams organised around customer journeys — ~60 journeys and ~400 'mini-CEOs' by May 2024.
- **Nov 2022** — ChatGPT lands; DBS begins exploring generative AI on the modern stack it has built since 2014.
- **2023** — Forms GSTAR (Group Strategy, Transformation, Analytics & Research) and a GenAI Steering Committee; sets GenAI guardrails (internal-only, human-in-the-loop). Launches DBS-GPT on a 'one horizontal, three verticals' plan.
- **May 2024** — Days after a record quarter, Piyush watches GPT-4o and weighs how far to push GenAI — the decision this case turns on.

THE STACK THEY BUILT (our Class 2 lens)

- **Data layer — ADA.** A central platform (a data 'backbone,' a 'laboratory' sandbox, and a 'factory' for production), tiered T1/T2/T3, with automatic, role-based access.
- **Governance / eval layer — P-U-R-E, inside RDU.** Every data and AI use case must be Purposeful, Unsurprising, Respectful and Explainable; the RDU Committee 'white-lists' approved cases so teams move fast within guardrails.
- **The split.** DBS builds and owns the data and governance layers; it rents the models.

THE GENAI BET — "one horizontal, three verticals"

- **Horizontal — DBS-GPT.** A proprietary assistant that plugs third-party models (GPT-4, Claude, BART) into the bank, paired with an Enterprise Knowledge Base (EKB) for retrieval over DBS's own documents — i.e. retrieval-augmented generation (RAG). 'Inch-deep, mile-wide.'
- **Three verticals.** A call-centre co-pilot (projected +20% productivity), a relationship-manager co-pilot (client research, AML/KYC screening), and software development (coding, testing).

- **The guardrail.** GenAI is internal-employee-only (not customer-facing), always human-in-the-loop, and every use case is reviewed by the RDU Committee — because a GenAI hallucination violates P-U-R-E's 'Explainable.'

KEY FACTS & FIGURES (as reported; point-in-time)

- **Scale.** 600+ AI/ML models across ~300 use cases (2022); since grown toward 2,000+ models / 400+ use cases.
- **Reach.** 45 million hyper-personalised 'nudges' a month to ~5 million customers.
- **Speed.** "DBS Quick Finance" — one minute to apply, one second to approve (SME loans up to SGD 300k).
- **Risk.** SME credit early-warning flagged 95% of non-performing loans ~3 months ahead; ~80% of at-risk borrowers averted.
- **Value.** ~SGD 180M of measured economic value in 2022; projected ~SGD 700M for 2024 (case) and ~SGD 1bn over five years (DBS).
- **People.** ~2,500 data scientists / analysts / engineers (more technologists than bankers); 9,000+ staff on AI learning pathways.

THE DECISION ON THE TABLE, MAY 2024

- Can DBS's existing governance (P-U-R-E + RDU + human-in-the-loop) scale GenAI with minimal human intervention — or does it need more oversight?
- Should GenAI stay internal-only, or is it time to let it face customers — accepting rare but consequential errors? "When the music is on, you have to get on the dance floor."
- Should DBS take centralised ownership of GenAI use cases, and how will the push reshape the workforce?

DISCUSSION QUESTIONS — we will work these in groups

There are no wrong answers. These are about how you reason, not what you recall.

1. DBS built its own data platform (ADA) and governance (P-U-R-E) but rents its foundation models (GPT-4, Claude). Go layer by layer: what did DBS own versus rent — and why does that split make sense for a bank?
2. Every DBS model must pass P-U-R-E — Purposeful, Unsurprising, Respectful, Explainable — before deployment. Which of the four is hardest for a generative-AI system to satisfy, and why?
3. DBS-GPT plugs the same models anyone can rent into the bank's own Enterprise Knowledge Base. What is the part a competitor cannot copy — and what does that tell you about where value sits in the stack?

Bring it home: where do you see the same pattern — own data + governance, rent the model, a data moat — in a bank or institution from your own country? (China: ICBC • CCB • WeBank • Ant • Ping An • India: HDFC • SBI • ICICI • Jio Financial)

GLOSSARY OF ACRONYMS

DBS — Development Bank of Singapore

HBS — Harvard Business School

AI / ML / LLM — Artificial Intelligence / Machine Learning / Large Language Model

GenAI — Generative AI

NLP — Natural Language Processing

RAG — Retrieval-Augmented Generation — retrieve relevant text, then generate from it

ADA — 'Advancing DBS with AI' — DBS's internal data-and-AI platform (its single 'source of truth')

P-U-R-E — DBS's responsible-data framework: Purposeful, Unsurprising, Respectful, Explainable

RDU — Responsible Data Use — DBS's governance model; the RDU Committee approves ('white-lists') AI use cases

EKB — Enterprise Knowledge Base — DBS's internal knowledge store that DBS-GPT retrieves from

DBS-GPT — DBS's proprietary internal GenAI assistant (rented models + the EKB = RAG)

GSTAR — DBS's Group Strategy, Transformation, Analytics & Research unit

FEAT — Fairness, Ethics, Accountability & Transparency — the Singapore regulator's AI principles

AML / KYC — Anti-Money Laundering / Know Your Customer

SME — Small & Medium-sized Enterprise

NOTES

1. This is a briefing, not the case. It summarises a copyrighted Harvard Business School case (625-053) for classroom discussion. For the full case, exhibits and financials, use the copy released after class — do not rely on this summary as a source.
2. The case is a snapshot of May 2024 (published 2024, revised 2026). Figures are point-in-time and have since grown; a great deal has changed, and noticing what has and has not is part of the exercise.
3. This briefing was generated and fact-checked using Anthropic models against the case text. Errors are still possible — flag anything that looks wrong.