



DAVID YOFFIE
ANNIKA Z. WEI
MAHEK MODI
DHRUV JAIN
SWAPNIL LAD

Anthropic's Next Step: From LLM to Agents?

In the fall of 2025, Dario Amodei, cofounder and CEO of Anthropic, contemplated the rapidly evolving landscape of artificial intelligence (AI). Just a few months earlier, in May 2025, the company released its newest generation of large language models (LLMs) under the Claude 4 family. Claude 4 Opus, its most advanced model, had set industry standards for software development and autonomous task execution.¹ With this launch, Anthropic surpassed OpenAI to become the leading provider of enterprise AI solutions, claiming a 32% market share.² Its post-money valuation grew from \$60 billion in March to over \$170 billion in September.³ Anthropic's cautious, safety-first ethos positioned the company at the forefront of building ethical and reliable LLMs.⁴

Yet a major industry shift was underway, one that threatened to upend Anthropic's position. Rather than focusing on LLM refinements, the AI industry was turning its attention to developing AI agents capable of performing tasks autonomously on behalf of users. In September 2024, Salesforce introduced Agentforce, a platform designed to help businesses deploy agents to automate sales workflows.⁵ Following in January 2025, Anthropic's biggest competitor, OpenAI, introduced Operator, an AI agent built to perform browser-based tasks such as form-filling and online shopping.⁶ Then, in May, Microsoft unveiled plans for "Build 2025," a company-wide initiative to develop an open agentic web where AI systems could act independently across organizational platforms.⁷ By this time, over \$2 billion had flowed into the AI agent ecosystem.⁸ Although current products were far from fully autonomous, one in four companies was predicted to have AI agents running pilot programs by year's end, with enterprise adoption expected to double by 2027.

In the quickly evolving race to develop agentic AI, Amodei stood at a crossroads. From its founding, Anthropic prioritized safety and long-term societal benefit over speed and commercialization.⁹ The agent market demonstrated enormous potential, but it also raised critical questions around technical readiness, cost, oversight, and user trust. Should Anthropic move up the stack from leading-edge LLMs into developing agentic products? Should it build its own agent platforms and applications, partner with others, or remain focused on providing the safest foundation models? Alternatively, should Anthropic grow its expertise on the Claude models' leading strength, coding, rather than expanding horizontally into new domains? Amodei's decision would be a consequential one.

Professor David Yoffie, Research Associate Annika Z. Wei, and Mahek Modi, Dhruv Jain, and Swapnil Lad (MBAs 2025) prepared this case. This case was developed from published sources. Funding for the development of this case was provided by Harvard Business School and not by the company. HBS cases are developed solely as the basis for class discussion. Cases are not intended to serve as endorsements, sources of primary data, or illustrations of effective or ineffective management.

Copyright © 2025, 2026 President and Fellows of Harvard College. To order copies or request permission to reproduce materials, call 1-800-545-7685, write Harvard Business School Publishing, Boston, MA 02163, or go to www.hbsp.harvard.edu. This publication may not be digitized, photocopied, or otherwise reproduced, posted, or transmitted, without the permission of Harvard Business School.

An Overview of Anthropic

Anthropic was founded in early 2021 by Dario Amodei and Daniela Amodei, alongside a small group of researchers who departed OpenAI.¹⁰ At the time, Dario Amodei had recently stepped down as Vice President of Research at OpenAI, where he grew increasingly concerned about the implications of large-scale commercialization on AI.¹¹ While OpenAI shifted away from its open source philosophy toward a more commercially driven approach, the Amodei siblings envisioned a new kind of AI company, one grounded in long-term safety, transparency, and alignment.¹²

Guiding Principles

As Anthropic took shape, its founders recognized the development of increasingly powerful AI systems required careful, long-term stewardship alongside technical innovation. Rather than the typical corporate architecture, Anthropic was established as a Public Benefit Corporation with a legal charter that bound the company to prioritize the long-term well-being of humanity.¹³ As part of its governance structure, the company created the Long-Term Benefit Trust (LTBT), an independent body of five trustees with backgrounds in AI safety, national security, public policy, and social enterprise.¹⁴ With the authority to appoint board members, the LTBT played a central role in aligning leadership and corporate decisions with Anthropic's values.

Much like the early internet developers had to build firewalls against cyber security threats, Anthropic sought new technical approaches to AI safety. A breakthrough came in the form of Constitutional AI, a method to reduce reliance on human-labeled examples for fine-tuning model behavior.¹⁵ Constitutional AI models were trained to follow a set of human-written guiding principles and refined based on critiques by other models. The approach enabled Anthropic to tailor its models' core principles, particularly when handling harmful or traumatic content, and allowed the systems to engage with sensitive prompts by explaining its objections rather than evading them.¹⁶ By embedding the value system of AI directly, Constitutional AI provided a scalable and transparent oversight.

Anthropic's emphasis on transparency extended beyond model development. The Transparency Hub, available to the general public, consolidated safety evaluations, model reports, usage policies, and research findings in one accessible site.¹⁷ The Hub included technical documentation outlining Claude's capabilities, safety testing results, and deployment safeguards, as well as updates on addressing risks and Anthropic's policies.¹⁸ In July 2025, Anthropic released a transparency framework intended for governments and regulators, proposing disclosure requirements for developers.¹⁹ These included safety protocols and risk mitigation efforts, reflecting the company's view that accountability should be a shared responsibility across both private and public sectors.

Growth & Development

In early 2021, generative AI (GenAI) was in its infancy. The most well-known model, GPT-3, still required additional finetuning just to complete specific, narrow tasks.²⁰ In May 2021, Anthropic launched with a \$124 million Series A round led by Skype cofounder Jaan (see **Exhibit 1** for Anthropic's funding history). The early funding laid the foundation for a small but highly specialized team of only 24 employees focused on interpretability and long-term safety.²¹

By mid-2022, Anthropic had finished training its first product, Claude, a new GenAI chatbot more powerful than any before.²² However, the company continued with internal testing due to unforeseeable risks of unleashing AI upon the world. When Claude was officially released in March 2023, OpenAI's ChatGPT had already secured a reputation as the leading GenAI system in the public

eye.²³ Nevertheless, Claude quickly became competitive in the chatbot market. Backed by tech giants, Claude leveraged Google's infrastructure and AI platform to support model training and expanded deployment through Amazon Bedrock.²⁴ The release of Claude 2 just four months later saw over 350,000 users on the waitlist for the new product.²⁵

Anthropic continued to develop Claude through distribution partnerships. By the end of 2023, Claude was integrated into enterprise workflows through partnerships with Salesforce and Boston Consulting Group.²⁶ Anthropic later formed a collaboration with Accenture, training over 1,400 engineers to implement Claude in client environments.²⁷ Claude also became accessible through consumer-facing tools such as Slack and GitHub Copilot, expanding accessibility for developers and knowledge workers.²⁸

In 2024, following the launch of Claude 3, Anthropic entered the international market. Claude became available across Canada and Europe with multilingual capabilities in French, German, Spanish, and Italian. Within a year, Anthropic's worldwide website userbase more than quadrupled to its peak at 18.8 million users.²⁹ By 2025, Anthropic had grown to over 1,000 employees.³⁰ Claude 4, released in May 2025, became the first "reasoning" model to maintain the speed of previous general LLMs while working through problems in a drawn-out, step-by-step fashion.³¹ The model provided high-quality answers significantly faster than OpenAI's top model ChatGPT o3. Moreover, the company's emphasis on safe and reliable AI development attracted a two-year \$200 million prototype agreement with the U.S. Department of Defense to develop frontier AI for national security.³² In September 2025, Anthropic closed a Series F fundraising of \$13 billion, valuing the company at \$183 billion post-money.³³

Anthropic's growing user base and strategic partnerships led to accelerating sales. By the end of May 2025, Anthropic told investors it had \$3 billion in ARR (annualized recurring revenue), exceeding internal projections of \$2.2 billion and tripling its revenue run rate from just under \$1 billion the previous year.³⁴ In the meantime, Anthropic expected to reduce its annual cash burn from \$5.6 billion in 2024 to \$3 billion in 2025, though estimated costs for computing alone remained high at \$2.5 billion.³⁵ Looking ahead, the company projected continued rapid growth, with base case revenue estimates ranging from \$12 billion to \$34.5 billion by 2027.³⁶ If the high-end forecast held, Anthropic expected to reach cash flow breakeven by 2027.³⁷

The Claude Family

The Claude family of AI models was Anthropic's primary commercial offering. Designed to support a wide range of professional and consumer applications, Claude evolved through successive releases to improve accuracy, reasoning, and ease of use. By mid-2025, the Claude 4 series represented the company's most advanced generation.

Claude Opus 4, the family's flagship model, could handle complex problem-solving and software development tasks, setting new benchmarks in industry evaluations.³⁸ The model could manage extended, multi-step projects with greater accuracy and reliability, while continuing to deliver responses quickly. In its August 2025 update, Opus 4.1 introduced improvements in research, data analysis, and coding precision, with notable gains in tracking fine details.³⁹ Claude Sonnet 4 provided balanced, all-purpose performance, switching easily between reasoning tasks and the use of external tools, making it suitable for a wide range of enterprise applications.⁴⁰ Claude Haiku 3.5, the smallest and fastest in the family, was designed for real-time scenarios such as customer service chats and live content creation, while maintaining skills in summarizing information and retrieving relevant details.⁴¹

Specialized Claude. Anthropic had also developed versions of Claude tailored to specific industries and use cases. Claude Code functioned as an AI software engineer, able to edit codebases, run tests, and execute commands to support development codes.⁴² Claude for Education was designed to enhance student learning and could be integrated with platforms such as Canvas.⁴³ Claude Gov was developed for national security applications in classified environments, and Claude for Financial Services consolidated internal and external financial data into a single interface for analysts and advisors.⁴⁴ Each variant was built to fit into domain-specific workflows, with an emphasis on safety, accuracy, and context-aware performance.

Claude for AI Agents Anthropic's approach to AI agents in 2025 emphasized safe delegation and user oversight rather than full autonomy.⁴⁵ The company promoted the development of composable and transparent agents that could plan tasks, interpret documents, and coordinate actions while keeping the user in control.⁴⁶ In May 2025, Anthropic released a set of API upgrades to support real-world agent applications, adding capabilities such as code execution and short-term memory.⁴⁷ The following month, the company introduced a multi-agent research system in which a central agent designed a research plan based on a user's query and coordinated parallel agents to gather information and use specialized tools simultaneously.⁴⁸

Tools & APIs. To support developers and enterprises, Anthropic built tools and APIs. The software interfaces allowed developers to connect Claude to other programs and services. These capabilities allowed Claude to write and run code, access private knowledge bases, and analyze data in real time.⁴⁹ Together, these tools reduced development time and enabled teams to build agents, applications, and workflows on top of Anthropic's products. In July 2025, Anthropic launched a Tool Directory to showcase recommended integrations for product teams, designers, and developers.⁵⁰ From generating roadmaps in Notion to converting Figma designs into production-ready code, the directory marked a shift toward positioning Claude as a platform, not just a chatbot.

Pricing. Claude was accessible via Claude.ai, where individuals could choose between a free tier or the Claude Pro plan, which offered priority access and higher usage limits for \$17 per month. The Max Plan bundled the most advanced features (i.e., Claude Code, web search, and integrations) for \$100 per month and was aimed at intensive users who relied on Claude for critical workflows.⁵¹ For collaborative work, the Teams plan provided shared workspaces and administrative tools at \$25-\$30/user per month. Organizations could opt into the Enterprise plan, which expanded the context window to 500,000 tokens and included security features.⁵² For developers and enterprise platforms, Claude was also available via API with prices above most competitors (see **Exhibit 2**).⁵³

Usage. Claude's adoption was driven by its reliability in reasoning, document analysis, and code generation tasks.⁵⁴ As of August 2025, Claude held 42% code generation market share, twice that of OpenAI.⁵⁵ As a result, coding remained Claude's largest use category, accounting for 36% of total activity (see **Exhibit 3** for Claude's usage), but educational and scientific applications grew rapidly as Claude's problem-solving and research capabilities improved. Business and financial operations usage, on the other hand, was reduced almost in half in 2025.⁵⁶

AI Agents

As LLMs continued to improve in its general capabilities, the frontier of artificial intelligence began shifting from text generation toward autonomous task completion. Anthropic defined the term "AI agents" as systems in which an LLM directed its own processes and tool use, deciding how to accomplish a task rather than following a fixed, pre-programmed sequence.⁵⁷ This approach

distinguished agentic AI from traditional workflows, which relied on rigid code and predetermined steps (see **Exhibit 4** for AI agent workflow).

The long-term vision for agentic AI involved systems capable of full autonomy. Analogous to Level 4 self-driving, Level 4 agentic AI referred to systems that could independently set goals, adapt plans, and operate across domains without human intervention.⁵⁸ As of late 2025, agentic AI products remained far more limited. The majority of so-called agentic AI products operated at Level 1 or 2, either executing fixed sequences through rule-based automation or navigating dynamic workflows with limited flexibility. Only a few experimental systems had begun to explore Level 3, where agentic AI exhibited partial autonomy to plan, execute, and adjust actions using domain-specific tools under minimal human oversight.⁵⁹

At a functional level, AI agents combined several components that allowed them to perceive, plan, and act.⁶⁰ At the center was a Reasoning Engine, typically powered by an LLM, which interpreted inputs, understood user intent, and generated plans or responses. A Tool Use and Execution Interface connected the agent to external systems, such as APIs, databases, or software tools, enabling it to act on those plans. A Memory System provided short-term and long-term context, helping the agent maintain continuity across interactions and execute multi-step tasks. Additional elements, such as Input Handling and Perception, allowed agents to process different types of information (e.g., text, structured data, or images). At the same time, a Goal and Planning Module broke complex objectives into manageable steps. Together, these components formed a flexible architecture that bridged reasoning and execution, enabling AI agents to adapt to varied environments.

Agents: A Brief A.I. History

The human fascination with autonomous, intelligent systems stretched back centuries, from early literature like Samuel Butler's *Erewhon* to cinematic depictions such as HAL 9000 in *2001: A Space Odyssey*.⁶¹ In 1950, Alan Turing proposed the Turing Test, suggesting any computer that could convincingly mimic human conversation exhibited intelligence.⁶² In 1966, a significant milestone came with the creation of ELIZA, the first chatbot, which simulated a psychotherapist by detecting key phrases in user input and responding with rephrased statements to sustain conversation.⁶³

As artificial intelligence advanced over the decades, systems gradually evolved in sophistication and capability.⁶⁴ In the 1980s and 1990s, the dominant approach was to build systems that could follow predefined rules to solve narrowly defined problems. This gave way in the 1990s and early 2000s when machine learning enabled AI to recognize patterns and make predictions based on limited memory systems. During this same period, Robotic Process Automation (RPA) emerged in enterprise settings for companies to automate rule-based digital tasks in specified, structured environments.⁶⁵ Rather than depending solely on hard-coded logic, systems developed during this era adapted through time, laying the groundwork for modern applications like speech recognition and recommendation engines.

The 2010s marked a turning point, as deep learning unlocked the ability to process unstructured data such as language, images, and audio. By the late 2010s, the emergence of GenAI and LLMs represented the most advanced form of limited memory AI. These systems could generate coherent text, simulate nuanced conversation, and even orchestrate complex digital workflows. While true self-aware AI remained a theoretical concept, GenAI systems demonstrated behaviors once considered science fiction, serving as the foundation for the emerging age of agentic AI.

AI Readiness: The Foundation for Agents

AI readiness encompassed the full spectrum of infrastructure, processes, and organizational capacities required to support the successful deployment of autonomous systems such as AI agents.⁶⁶ Readiness reflected how well organizations and their surrounding ecosystems could integrate intelligent agents into real-world environments. In the case of AI agents, readiness was particularly multifaceted. It included external technological infrastructure, internal organization, and the broader system-level considerations of trust, safety, and regulation that governed adoption at scale. Together, these layers formed the foundation of agent viability in practice.⁶⁷

External Infrastructure External technological readiness hinged on the maturity of the infrastructure supporting autonomous systems. LLMs, including multimodal systems like Anthropic's Claude, provided essential capabilities such as contextual reasoning and dynamic task execution.⁶⁸ However, models alone were not sufficient. Agents required secure and robust execution environments to run code, make API calls, and interface safely with external systems. Retrieval-Augmented Generation (RAG) emerged as a key infrastructure for grounding outputs by connecting LLMs with external knowledge bases.⁶⁹ By supplementing model reasoning with real-time access to external data, RAG improved accuracy and provided context for specialized domains without extensive fine-tuning. While early implementation relied on single data sources, RAG systems were developed to support agentic AI. Still in its early stages of development, agentic RAG would enable an agent to retrieve, evaluate, and apply information from multiple sources and APIs.⁷⁰

Complementing this infrastructure were orchestration frameworks, which structured agent behavior.⁷¹ At the highest level, sophisticated control mechanisms, including monitoring systems and safety guardrails, were required to ensure agents operate efficiently, reliably, and within predefined ethical and operational boundaries.⁷² One serious challenge was latency, which could become a critical bottleneck. Arising from the computational complexity of LLMs, latency could undermine the responsiveness and overall effectiveness of agentic AI systems.⁷³ Ultimately, the overall effectiveness of any agentic system was constrained by the weakest link. Without dependable infrastructure at every level, even the most capable model could not function safely or effectively.

Internal Integration Internal technological integration referred to an organization's ability to weave AI agents into its existing systems. A prerequisite was the availability of clean, structured, and accessible data, particularly for task automation workflows where agents needed reliable information to learn and operate effectively. Since many agents required broad access to potentially sensitive information, robust data security policies and technical safeguards were paramount to protect confidential data and ensure compliance.⁷⁴

Equally as important was the organization's ability to manage agent deployment. The cost of implementing and running agentic systems, whether based on usage, output, or hybrid models, required careful budgeting and return on investment (ROI) analysis.⁷⁵ Critical errors made by autonomous systems could carry significant consequences, including reputational damage, compliance breaches, or costly operational disruptions that easily outweighed any direct efficiency gains.⁷⁶ As a result, internal integration increasingly demanded risk-aware planning.

Regulations An additional concern was the rapidly evolving regulatory environment. The increasing autonomy of AI agents raised complex issues of liability, where determining responsibility became ambiguous.⁷⁷ Policymakers around the world struggled to keep pace with agentic technologies, leading to fragmented legal approaches.⁷⁸ As a result, many enterprises hesitated to adopt agents, fearing that deployments could become future liabilities.⁷⁹ Beyond corporate risk,

regulators worried about labor markets. AI agents showed the potential to automate routine white-collar tasks in customer service, sales, and content creation, raising concerns about job displacement.⁸⁰

Meeting Diverse Needs The adoption and perceived value of AI agents varied significantly across different market segments.⁸¹ Enterprises prioritized solutions that integrated seamlessly with existing systems, aligned with compliance requirements, and offered strong governance and monitoring tools.⁸² Small and midsize businesses (SMBs), in contrast, operated under different pressures. Their primary driver for technology adoption was typically a clear ROI, with many SMBs piloting tools offering identifiable financial and time savings.⁸³ While SMBs were frustrated with unproductive tasks and desired more automation, their willingness to pay for AI Agents often hinged on successful trials and quickly demonstrating tangible value.⁸⁴

Consumer demand for AI agents was also distinct. User experience was of the highest importance. Agents needed intuitive interfaces, easy access, and natural interaction, often leveraging conversational AI to mirror successful chatbots like ChatGPT.⁸⁵ Privacy was a major concern, given that agent effectiveness often relied on access to personal data. Pricing also played a critical role. Users expected clear, affordable options that reflected the value they received. Finally, consumer satisfaction hinged on agents requiring accuracy, reliability, and utility.⁸⁶

The Future of AI Agents

In early 2025, Jensen Huang, NVIDIA's CEO, described the evolution of artificial intelligence as a four-stage progression: from perception AI systems that could see and hear, to GenAI systems that could create and respond, to agentic AI systems that acted autonomously, and eventually to physical AI robots and embodied agents in the real world.⁸⁷ Technology was at the GenAI stage, with applications introduced to consumers and the workforce (see **Exhibit 5** for GenAI workforce automation). AI development was focused on the third phase, the era of agentic AI.⁸⁸ By 2028, 33% of enterprise software applications were expected to include AI agents, and nearly 90% of executives viewed agentic AI as a potential source of competitive advantage.⁸⁹ With the global agent market projected to grow from \$5.1 billion in 2024 to \$47.1 billion in 2030, the tech world hoped that agentic AI was poised to become a new foundation for how work gets done.⁹⁰

Beyond individual AI agents, developments likely involved multi-agent systems (MAS), where AI agents of different expertise collaborated, competed, or negotiated to complete collective tasks.⁹¹ For example, in a smart city, a multi-agent system could manage traffic in real time through vehicle-to-everything (V2X) communication.⁹² Each traffic light, public transit node, and vehicle would be equipped with its own AI agent, sharing data like location, speed, and road conditions. These agents would work together to adjust signal timing, reroute traffic, notify emergency services, and guide vehicles or pedestrians away from disruptions such as accidents. If a parking lot was full, the system could redirect a vehicle, even overriding its initial preference, to optimize flow and reduce energy use. As these systems grew more interconnected, they could offer a glimpse into the next generation of intelligent environments powered by coordinated AI agents.

Competition in the AI Agent Race

As AI agents moved from research concept to commercial deployment, the competitive landscape faced by Anthropic intensified. Competitors were aggressively investing in data centers. As of Q2 2025, private fixed investment in information processing equipment and software in the United States reached 4.44% of GDP, a level comparable to the peak of the dot-com boom in 2000.⁹³ All tech giants were focused on developing frontier LLMs and driving down the costs per token (words, subwords,

or individual characters). Delivering answers with GenAI (called inference) depended on the number of tokens required for a model to process. The good news was that the price per token was dropping dramatically (see **Exhibit 5**); the bad news was that the number of tokens required for advanced reasoning was exploding.⁹⁴

A growing number of companies, ranging from foundational model labs to consumer tech giants and enterprise cloud providers, were racing to define the agentic AI era. Some players emphasized open-source accessibility while others prioritize multimodal capabilities or enterprise-grade tooling. Across the board, there was growing investment in developing agents that could reason, act, and collaborate in real-world environments with minimal human oversight.

Open- vs Closed-Source Models

Open source software (OSS) encompassed a wide array of software licenses which generally made source code freely available.⁹⁵ OSS emerged as a pushback against the shift of software from a shared resource to a proprietary product in the 1970s.⁹⁶ As OSS progressed, it played an important role in modern computing, with developers relying on open source for everything from basic infrastructure to advanced applications.⁹⁷ Despite its wide use, open source had long struggled with monetization. Linux adoption, for example, depended on corporate donations and support from Red Hat's enterprise services model, while Google captured value from Android through its proprietary layers such as the Google Play ecosystem.⁹⁸ Monetizing open ecosystems often necessitated complementary business models to sustain innovation.

In the agentic AI landscape, competitors were divided between open and closed models. Claude was a closed-source framework, with neither its model weights nor training code publicly available. Closed-source models like Claude, OpenAI's GPT-5, and Google Gemini emphasized proprietary performance gains, safety, and fast enterprise deployment, but limited user transparency and customization.⁹⁹ At the other end of the spectrum, fully open-source models such as OpenLLaMA, and some models developed by France's Mistral AI, allowed developers to self-host, modify, and extend systems with greater affordability and transparency, but requiring more technical capabilities.¹⁰⁰ There was evidence emerging that the performance gap between open and closed models had begun to narrow, with open systems approaching parity with proprietary counterparts in areas such as math and reasoning.¹⁰¹ In specialized verticals, such as biomedicine, law, and defense, some open source models were already taking the lead due to closer alignment with domain-specific requirements.¹⁰²

The potential role of open source was contentious. In a controversial memo on LLMs leaked by a Google engineer in 2023, titled "We Have No Moat," the engineer stated:

Plainly put, [open source] are lapping us...Open source models are faster, more customizable, more private, and pound-for-pound more capable. They are doing things with \$100 and 13B params that we struggle with at \$10 million and 540B. Indeed, in terms of engineering-hours, the pace of Improvement from [specialized] models vastly outstrips what we can do with our largest variants, and the best are already largely Indistinguishable from ChatGPT...The modern Internet runs on open source for a reason. Open source has some significant advantages that we cannot replicate.¹⁰³

A middle ground was open-weight models, where companies disclosed model training weights alongside selective source codes. Meta and China's DeepSeek R1 models were representative of this strategy.¹⁰⁴ Increasingly, companies adopted hybrid strategies, with proprietary premium models or APIs and open lower-tier models.¹⁰⁵ OpenAI, for example, introduced two open-weight models in August 2025, while keeping GPT-4o and GPT-5 closed.¹⁰⁶ Similarly, Mistral AI started offering models

under open-weight licenses but reserved its more capable models as closed-source.¹⁰⁷ This approach aimed to balance transparency and community-driven innovation with safeguards for security and monetization.¹⁰⁸ In contrast with competitors' varied approach, all Anthropic's models remained closed source with an emphasis on safety, control, and enterprise-grade deployment.

OpenAI Among the foundational LLM companies, OpenAI stood out as Anthropic's most direct competitor. The company emerged as a central player in agentic AI by progressively turning its research into mainstream products that empower agents with autonomous action. In January 2025, OpenAI launched Operator, an AI agent designed to perform consumer web tasks. In the same month, OpenAI released a Computer-Using Agent (CUA) that leveraged GPT-4o's vision capabilities to perform tasks without custom APIs.¹⁰⁹

In July 2025, these developments converged in the ChatGPT agent.¹¹⁰ The system combined Operator's autonomy with advanced research tools and coding capabilities. It could plan and execute complex workflows like calendar management, web research, data analysis, and slide generation, all while ensuring user control.¹¹¹ To support developers, OpenAI introduced new tools to make it easier to design and test agents.¹¹² Compared with specialized versions of Claude, the ChatGPT Agent offered general-purpose capabilities that were more versatile but required additional human oversight.¹¹³

In October 2025, OpenAI launched a new generation of apps within ChatGPT. The update allowed users to perform everyday tasks such as creating a weekend playlist through Spotify or browsing homes on Zillow without leaving ChatGPT.¹¹⁴ Other integrated services, including Figma, Expedia, and Booking.com, broadened these capabilities across creative and travel planning workflows. This expanding ecosystem positioned ChatGPT not as a chatbot, but as an operating layer for AI-mediated work, blurring the line between applications and agents.¹¹⁵

The scale of OpenAI's adoption was impressive. With the release of ChatGPT 5 in August 2025, the chatbot accumulated approximately 700 million weekly active users, with users generating over 2.5 billion prompts daily in both work and personal settings (see **Exhibit 6** for ChatGPT usage).¹¹⁶ Additionally, over 90% of Fortune 500 companies used OpenAI products in some capacity: Enterprise subscriptions surged, with ChatGPT Enterprise growing to more than 600,000 paying business users, despite weak reviews for ChatGPT5.¹¹⁷

OpenAI's revenue trajectory was explosive. By mid-year 2025, the company had reached \$12 billion in ARR compared to \$4 billion in revenues in 2024.¹¹⁸ However, OpenAI's projected cash burn increased to \$8 billion for 2025. Meanwhile, infrastructure expenses such as server rentals could potentially exceed projections of \$14 billion.¹¹⁹ Despite its scale, the company remained deeply unprofitable. It did not anticipate becoming cash-flow positive until 2029, when projected revenue would top \$125 billion.¹²⁰

Google Google was investing aggressively to become a leader in agentic AI. DeepMind's Gemini 2.5 family delivered native reasoning, planning, and multimodal capabilities critical for agentic behavior.¹²¹ In March 2024, Google released SIMA, a general-purpose AI agent able to follow language commands and interact with 3D virtual environments.¹²² Unlike traditional systems that required custom code for every action, SIMA could operate directly from natural language instructions.¹²³

Along with individual AI agents, Google announced AI co-scientist, a multi-agent system built on Gemini 2.0 that accelerated hypothesis generation and scientific research workflows.¹²⁴ On the systems side, Google Cloud's Vertex AI Agent Builder enabled developers to design and orchestrate multi-agent workflows with minimal coding. Google announced in 2025 that over 1,000 AI agent use cases

had been built by partners across nearly every industry.¹²⁵ Google also released the Agent2Agent (A2A) protocol designed for agents from different vendors to communicate and collaborate securely across enterprise environments.¹²⁶ Over 50 enterprise technology companies supported the protocol.¹²⁷ Leveraging its global cloud infrastructure, Google maintained a scale and distribution advantage that positioned it as a significant competitor in the AI agent market.

Meta Meta entered the AI agent market with its Llama models, first introduced in February 2023.¹²⁸ Unlike OpenAI, Anthropic or Google, Meta deployed an open-weight strategy in the hope of accelerating broader deployment of its models. While Meta aggressively pushed the advantages of open-weight models, Zuckerberg hinted in 2025 that he considered closing the model in the future.¹²⁹

Like Anthropic, Meta heavily courted the developer community. In March 2025, Meta began piloting Business AI agents on its social platforms, enabling SMBs to deploy customer-service and commerce assistants that autonomously responded to inquiries using product catalogs and FAQs.¹³⁰ With more than 400 million monthly users already interacting through voice, vision, and chat across its ecosystem, Meta could embed agents into everyday social and business environments.¹³¹

Despite Meta's broad user base and significant investments in AI, some tech analysts, as well as Zuckerberg, worried that Meta was falling behind in 2025.¹³² To address this fear, Zuckerberg launched an ambitious initiative in June 2025 to pursue artificial general intelligence through the creation of Meta Superintelligence Labs (MSL).¹³³ Zuckerberg started a war for talent, paying huge sums (e.g., \$100 million signing bonuses) to attract star AI developers.¹³⁴ Meta even made a \$14.3 billion investment in Scale AI, mainly to bring on board Alexandr Wang, Scale AI's 28-year-old CEO.¹³⁵ It was unclear a few months later whether this strategy would deliver on its promise: Meta failed to attract key talent from places like Anthropic, where developers were more mission-driven than compensation-driven.¹³⁶ Huge pay disparities at MSL also generated cultural problems. By late August, a few key developers left their pay packages at Meta to return to OpenAI.¹³⁷

Startups A vibrant startup ecosystem was developing around agentic AI. Venture was investing billions of dollars into AI agentic start-ups.¹³⁸ For example, TinyFish developed enterprise web agents that mimicked human browsing behavior to execute complex, multi-step workflows such as dynamic price surveillance and inventory monitoring.¹³⁹ In customer experience, Decagon scaled conversational agents to more than 10 million end-users via chat, email, voice, and SMS.¹⁴⁰ Its systems promised to deliver concierge-level support with fewer human escalations.¹⁴¹ ThriveAI positioned its product as the world's first AI-powered junior product manager.¹⁴² Its agent integrated with Slack and Microsoft Teams to track issues, synthesize feedback, and monitor competitors.¹⁴³ In legal services, Harvey introduced agents capable of drafting documents, conducting research, reviewing contracts, and supporting compliance and litigation. By mid-2025, Harvey's agents were in use at over a quarter of the top 100 U.S. law firms, strengthened by a partnership with LexisNexis to embed authoritative legal content directly into workflows.¹⁴⁴

International While U.S.-based companies were the dominant players for AI agents in 2025, international players were rapidly emerging with competitive offerings. The Canadian AI company Cohere released its Command model family to optimize for agentic AI in finance, healthcare, and public sector environments.¹⁴⁵ The company's annual revenue rose from approximately \$13 million in 2023 to a projected \$70 million in 2025, with major clients such as Oracle, McKinsey, and the Royal Bank of Canada.¹⁴⁶ In Europe, Mistral AI was at the forefront of open-source agent development with the launch of its Agents API and its enterprise platform Le Chat Enterprise.¹⁴⁷ In 2024, the company raised more than \$600 million at a roughly \$6 billion valuation.¹⁴⁸ In China, leading tech firms were aggressively expanding into the AI agent space, combining open-weight innovation with enterprise-

scale deployment. DeepSeek had introduced a series of LLMs with advanced reasoning capabilities, robust natural language understanding, and efficient performance.¹⁴⁹ Its open-weight model DeepSeek-R1 excelled at multistep reasoning and chain-of-thought logic.¹⁵⁰ For enterprise agentic AI, these capabilities could translate into stronger problem-solving and the ability to automate complex workflows across domains. The Chinese tech giant Tencent presented an Agent Development Platform, a full-stack agent platform to extend “digital companionship” to a population of 1.4 billion users.¹⁵¹ Its AI Arena supported an open research platform, allowing academic and industry teams to test agents in simulated environments mirroring real-world challenges.¹⁵² Similarly, ByteDance open-sourced its Coze platform and launched Trae Agent, a developer tool for automating software engineering tasks.¹⁵³

Agentic AI in Practice: Adoption, Implementation, and Friction

With billions of dollars being invested in agents by small and large competitors around the world, the big question was whether consumers and businesses would adopt the products and agent technologies. As of summer 2025, many companies were experimenting with agentic AI. Organizations across sectors were beginning to embed AI agents into real-world workflows, using them not just for information retrieval or summarization, but for delegated action and coordination. In financial services, agents performed continuous risk monitoring and provided dynamic investment advice based on real-time market conditions. In healthcare, hospitals used agents to automate administrative tasks like scheduling, billing and taking clinical notes in patient appointments. In customer experience, agentic systems could proactively resolve service requests, analyze sentiment, and escalate complex issues without human prompt. Early use cases also emerged in supply chain management, education, agriculture, and mental health, where piloted agents were involved in tasks ranging from demand forecasting to personalized learning and cognitive behavioral support.^{154, 155}

Despite these promising applications, agentic AI remained in the early stages of development. Agents still relied heavily on human input or operated within constrained environments, lacking the robustness and autonomy needed to execute complex, multistep goals reliably.¹⁵⁶ A Gartner poll showed only 19% of organizations had made significant investments in agentic AI, with another 42% making conservative investments, and 39% taking a wait-and-see approach or making no investments.¹⁵⁷ The lack of clarity around what constituted a true “agent” complicated adoption. Many offerings labeled as agentic were, in practice, advanced GenAI tools or simple task automation bots. Thus, agentic AI propositions often lack value and ROI, with use cases positioned as agentic, not requiring agentic implementations. As a result, one analysis predicted that 40% current agentic AI projects would face cancellation by 2027.¹⁵⁸ A different study by MIT in July 2025 reported that “despite \$30-40 billion in enterprise investment in GenAI, ... 95% of organizations are getting zero return,” largely due to flawed integration and misaligned expectations.¹⁵⁹

Barriers to Adoption

As a nascent technology, AI agents faced significant headwinds in adoption, including reliability, implementation, and costs (see **Exhibit 7** for barriers to AI adoption).¹⁶⁰ A primary barrier to AI agents’ adoption was the lack of trust. According to IBM, 45% of business leaders had concerns with data accuracy and bias in AI results.¹⁶¹ Organizations that demanded high levels of reliability, transparency, and control often cited the risks of bias and limited explainability. Most companies trusted AI agents to make decisions only within a defined framework under human supervision, with less than 3% trusting AI agents to make judgment calls autonomously.¹⁶²

Some of these concerns were clearly warranted. One of the main limitations of AI was hallucinations, “a phenomenon where... a generative AI chatbot or computer vision tool, perceives patterns or objects that are nonexistent or imperceptible to human observers, creating outputs that were nonsensical or altogether inaccurate.”¹⁶³ Although fine-tuning alignment techniques were expected to reduce hallucinations, hallucination rates stayed high in newer models such as GPT o3 and o4-mini.¹⁶⁴ As a result, widespread executive apprehension about AI accuracy persisted, and hallucination remained a top reason for hesitancy in broadly deploying AI.¹⁶⁵

Enterprise adoption was further hindered by the complexity of integrating agents into systems, a shortage of internal technical expertise, and increasing uncertainty around economic returns. Successful AI agent deployments often required access to eight or more data sources, and over 86% of organizations needed significant upgrades to their existing tech stacks to support agentic capabilities.¹⁶⁶ Agents depended on dynamic reasoning models, which were computationally intensive and demanded substantial energy and hardware resources, further raising infrastructure costs and environmental concerns.¹⁶⁷ Many early pilots also struggled to demonstrate a clear performance impact. For many enterprises, the inability to fully trust AI outputs, combined with uncertain ROI and operational risk, remained a central reason for delaying or scaling back agent deployments.

For individual consumers, a high proportion of users expressed discomfort with the level of personal data collected by AI-driven applications, particularly when agents operated across devices or accessed sensitive information.¹⁶⁸ Similarly, ethical concerns such as fairness and bias also impeded AI adoption. Additionally, many users harbored psychological reservations about interacting with emotionless systems and delegating subjective, human-centered tasks to AI.¹⁶⁹ The growing autonomy of AI furthered these worries. A 2025 Pew survey found that 76% of Americans were apprehensive about riding in self-driving cars, reflecting broader discomfort with fully autonomous systems.¹⁷⁰ The idea of AI agents acting independently, without clear oversight, often amplified those anxieties, creating a persistent gap between technical capability and user trust.

Model Switching and Convergence

In the GenAI era, users increasingly engaged in multi-homing, tapping into multiple LLMs.¹⁷¹ Usage data from LLM marketplaces indicated that individuals often rotated among models depending on pricing, latency, and context. Among the Fortune 500 companies, adoption was just as distributed. Most firms reported relying on anywhere from three to over five different model providers.¹⁷² The trend reflected a need to match models to specific use cases, optimize costs, avoid vendor lock-in, and stay ahead of rapid advancements.

However, as workflows became more complex and LLMs more deeply embedded in enterprise systems, switching costs could rise. Each model interpreted prompts, formatted responses, and structured outputs differently, introducing friction in reliability, tone, and accuracy.¹⁷³ For instance, while most leading models supported context windows of up to 128,000 tokens, Gemini 1.5 Pro extended to 2 million tokens, enabling longer and more nuanced inputs.¹⁷⁴ Seemingly small differences in how models processed inputs or returned outputs could require significant time and effort to adjust code, retrain agents, and maintain consistent performance. As a result, switching from one model to another, even when technically feasible, could incur hidden operational costs.

To address these frictions, a growing number of tools and frameworks emerged. New orchestration platforms like Krista offered a “hot-swap” capability, enabling enterprises to switch models on demand without rewiring prompts or integrations.¹⁷⁵ Similarly, RAG-powered approaches were being developed to enable dynamic selection between different LLMs.¹⁷⁶ Instead of choosing one model, the

system could automatically pick the best one for each situation based on past performance, eliminating the need for developers to perform a full rewrite when switching models.

As LLMs continued to develop, one theory suggested that the capabilities of different foundational models in areas such as efficiency, cost, and output quality could plateau and become increasingly similar.¹⁷⁷ This shift could reduce technical differences between models, making it easier for organizations to switch providers without sacrificing performance. Over time, the convergence of AI model capabilities in performance and efficiency may enable more flexible agent architectures that allow for switching foundational models of AI agents at low costs. Yet for leading LLM developers such as Anthropic, convergence raised questions about long-term differentiation.¹⁷⁸ If performance differences narrowed, would features such as safety, developer ecosystems, and ease of integration become critical?¹⁷⁹

Anthropic at a Crossroads: Strategy for the Agent Era

As Dario Amodei surveyed the rapidly shifting AI agent landscape in 2025, he faced strategic decisions that would define Anthropic's place with AI agents. Amodei viewed himself as visionary: every month, he held his "Dario Vision Quests" with employees, where he would preach about the future of AI.¹⁸⁰ Anthropic's role in the AI agent landscape was an obvious topic of great concern. On the one hand, investment in agents was burgeoning across consumer, SMB, and enterprise segments. Yet, significant hurdles remained, spanning trust, technical reliability, economic viability, liability, and regulation. Anthropic, with its deeply ingrained safety mission, had earned a reputation as a principled innovator. How it navigated this new phase could define its future.


The central dilemma centered on how to best engage with the agent market while remaining true to its core principles. Anthropic's late entry into the chatbot race, despite having a product to launch before OpenAI's ChatGPT, severely impacted its competitive position. One path forward was to adopt a first-mover strategy, aggressively launching its own agent platforms or applications even if it meant accepting higher risks. Another option was to move deliberately, learning from competitors like OpenAI, Google, and a new crop of agent-focused startups as they tested the market to identify successful applications and surface potential pitfalls. Although the latter would enable Anthropic to enter later with more robust, safe, and potentially more differentiated offerings, the market was evolving at breakneck speed.

If Anthropic chose active participation, the target market became crucial. Should it focus on the consumer market, leveraging Claude's conversational strengths for personal assistants and go head-to-head with OpenAI's Operator? Or should it prioritize the SMB market, offering versatile, cost-effective agents to augment smaller teams? Alternatively, should Anthropic target the enterprise and government, where safety and the potential for deep integration might carry weight, but where competition was intensifying?


Furthermore, the mode of entry required careful consideration. Building the entire agent stack, from foundational models to platforms and end-user applications, would give Anthropic maximum control. However, Anthropic was losing money, and a full-stack solution would require significant resources and new technical capabilities. A partnership model, in which Claude powered agent platforms developed by enterprise software providers or startups, would lean on Anthropic's strengths, but cede influence over the final product and user experience. A hybrid approach could combine these paths, developing core agentic capabilities via APIs while partnering selectively for targeted vertical applications. Alternatively, Anthropic could also release open-weight models.

An alternative path was for Anthropic to pursue vertical specialization by focusing on its most widely used capability: coding. Rather than expanding horizontally into general-purpose agents, Anthropic could concentrate on making Claude the leading system for code generation, debugging, and autonomous software development. Claude's established strengths in reasoning and reliability positioned it well to capture a technically demanding and commercially valuable niche. By further developing coding-oriented agentic tools such as code execution, software testing, and multi-agent collaboration, Anthropic could establish itself as the standard for safe and capable coding agents within the broader AI ecosystem.


Underlying all these choices was a fundamental tension between Anthropic's public benefit mandate and its commercial ambitions in a fast-moving, uncertain, and fiercely competitive market. Ensuring that Anthropic's agentic systems stayed true to its constitutional principles while going up against rivals willing to prioritize speed over safety would be no small feat. A strong safety focus could become a lasting competitive edge, especially in regulated industries and among risk-averse enterprises. On the other hand, it could slow Anthropic to the point of losing relevance. As Amodei prepared to guide Anthropic through this next phase, the decision would shape not only the company's future trajectory but also influence the broader development of safe and beneficial AI. The path forward required navigating the delicate balance between pioneering innovation, commercial reality, and unwavering commitment to the long-term well-being of humanity. What should Dario Amodei do?

Exhibit 1 Anthropic Funding History and Projected Valuation

Source: Created by casewriters using data from Pitchbook database, "Anthropic Overview," <https://pitchbook.com/profiles/company/466959-97#funding>, accessed August 2025.


Exhibit 2 Aug 2025 LLM API Cost

Source: Created by casewriters based on data from each individual company's price for LLM APIs. Pricing based on $\leq 200k$ tokens. "Pricing," Anthropic, <https://www.anthropic.com/pricing#api>; "API Pricing," OpenAI, <https://openai.com/api/pricing/>; "Gemini Developer API Pricing," Google AI for Developers, <https://ai.google.dev/gemini-api/docs/pricing>; "Llama 4 Pricing: API Cost vs. Local Hardware (2025)," llamaimodel.com, <https://llamaimodel.com/price/>; "API Pricing," Mistral AI, <https://mistral.ai/pricing#api-pricing>; "Models & Pricing," DeepSeek API Docs, https://api-docs.deepseek.com/quick_start/pricing, accessed August 20, 2025.

Exhibit 3 Claude.ai Usage


Source: Adapted by casewriters from Ruth Appel, Peter McCrory, Alex Tamkin, Miles McCain, Tyler Neylon, and Michael Stern, *The Anthropic Economic Index Report: Uneven Geographic and Enterprise AI Adoption*, Anthropic, September 15, 2025, <https://www.anthropic.com/research/anthropic-economic-index-september-2025-report>, accessed October 10, 2025.

Note: Each category shows the share of sampled conversations on Claude.ai associated with tasks from each SOC major group.

Exhibit 4 GenAI and Agentic AI Workflows

Source: Created by casewriters.

Exhibit 5 LLM Inference Costs for 2 million Tokens (input + output)




Source: Created by casewriters using data from Guido Appenseller, “LLMflation – LLM inference cost is going down fast,” Andreessen Horowitz, November 12, 2024, <https://a16z.com/llmflation-llm-inference-cost/>; and Tanay Jaipuria, “The Plummeting Cost of Intelligence,” Wing Venture Capital, November 4, 2024, <https://www.wing.vc/content/plummeting-cost-ai-intelligence>, accessed August 2025.

Exhibit 6 ChatGPT Usage

	Asking	Expressing	Doing
Work	12.2%	3.0%	19.5%
Non-Work	36.7%	8.3%	20.3%
Total	48.9%	11.3%	39.8%

Source: Adapted by casewriters from Chatterji, Aaron, Thomas Cunningham, David J. Deming, Zoe Hitzig, Christopher Ong, Carl Yan Shan, and Kevin Wadman. *How People Use ChatGPT*. p. 20, Fig. 13, No. w34255. National Bureau of Economic Research, 2025, <https://www.nber.org/papers/w34255>, accessed December 2025.

Exhibit 7 Barriers to AI Adoption

Source: Created by casewriters using data from, "Tracking Firm Use of AI in Real Time: A Snapshot from the Business Trends and Outlook Survey," page 37, Table 7, U.S. Census Bureau October 2024, <https://www2.census.gov/library/working-papers/2024/adrm/ces/CES-WP-24-16R.pdf>, accessed August 2025.

Note: Data represented the share of firms reporting each reason for not planning to use AI in producing goods or services during the six months following data collection.

Endnotes

- ¹ "Introducing Claude 4," *Anthropic*, May 22, 2025, <https://www.anthropic.com/news/claude-4>, accessed July 27, 2025.
- ² "2025 Mid-Year LLM Market Update: Foundation Model Landscape + Economics," *Menlo Ventures*, August 1, 2025, <https://menlovc.com/perspective/2025-mid-year-llm-market-update/>, accessed August 2, 2025.
- ³ George Hammond and Cristina Criddle, "AI start-up Anthropic valued at \$170bn in expanded funding round," *Financial Times*, September 2, 2025, <https://www.ft.com/content/c701ea30-0454-4c7d-a88d-1cf63eb9f008>, accessed September 3, 2025.
- ⁴ "Claude's Constitution," *Anthropic*, May 9, 2023, <https://www.anthropic.com/news/claude-constitution>, accessed July 27, 2025.
- ⁵ "Salesforce Unveils Agentforce—What AI Was Meant to Be," *Salesforce*, September 12, 2024, <https://www.salesforce.com/news/press-releases/2024/09/12/agentforce-announcement/>, accessed August 2, 2025.
- ⁶ "Introducing Operator," *OpenAI*, January 23, 2025, <https://openai.com/index/introducing-operator/>, accessed June 25, 2025.
- ⁷ Frank X. Shaw, "Microsoft Build 2025: The age of AI agents and building the open agentic web," *Official Microsoft Blog*, May 15, 2025, <https://blogs.microsoft.com/blog/2025/05/19/microsoft-build-2025-the-age-of-ai-agents-and-building-the-open-agentic-web/>, accessed August 2, 2025.
- ⁸ Theo Schnitfink, "How AI Agents Are Transforming Business In 2025 And Beyond," *Forbes*, March 23, 2025, <https://www.forbes.com/councils/forbestechcouncil/2025/03/27/how-ai-agents-are-transforming-business-in-2025-and-beyond/>, accessed August 2, 2025.
- ⁹ "Core Views on AI Safety: When, Why, What, and How," *Anthropic*, March 8, 2023, <https://www.anthropic.com/news/core-views-on-ai-safety>, accessed July 27, 2025.
- ¹⁰ Harrison, Kyle, Andriy Mulyar, Brandon Duderstadt, and Alex Liu. "The Openness of AI." *Contrary Research*, July 14, 2023. <https://research.contrary.com/deep-dive/the-openness-of-ai>, accessed February 25, 2025.
- ¹¹ Billy Perrigo, "Inside Anthropic, the AI Company Betting That Safety Can Be a Winning Strategy," *Time*, May 30, 2024, <https://time.com/6980000/anthropic/>, accessed July 28, 2025.
- ¹² Kyle Harrison, Andriy Mulyar, Brandon Duderstadt, and Alex Liu, "The Openness of AI," *Contrary Research*, July 14, 2023, <https://research.contrary.com/deep-dive/the-openness-of-ai>, accessed February 25, 2025.
- ¹³ "At Anthropic, We Build AI Systems That Are Steerable, Interpretable, and Safe," *Anthropic*, <https://www.anthropic.com/company>, accessed March 12, 2025.
- ¹⁴ "The Long-Term Benefit Trust," *Anthropic*, September 19, 2023, <https://www.anthropic.com/news/the-long-term-benefit-trust>, accessed July 27, 2025.
- ¹⁵ Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen et al., "Constitutional Ai: Harmlessness from Ai Feedback," *arXiv Preprint*, December 15, 2022, <https://doi.org/10.48550/arXiv.2212.08073>, accessed December 2025.
- ¹⁶ "Claude's Constitution," *Anthropic*, May 9, 2023, <https://www.anthropic.com/news/claude-constitution>, accessed July 27, 2025.
- ¹⁷ "Introducing Anthropic's Transparency Hub," *Anthropic*, February 27, 2025, <https://www.anthropic.com/news/introducing-anthropic-transparency-hub>, accessed July 28, 2025.
- ¹⁸ "Anthropic's Transparency Hub: Model Report," *Anthropic*, June 3, 2025, <https://www.anthropic.com/transparency/model-report>, and; "Anthropic's Transparency Hub: Platform Security," *Anthropic*, February 23, 2025, <https://www.anthropic.com/transparency/platform-security>, accessed July 28, 2025.
- ¹⁹ "A Framework for AI Development Transparency," *Anthropic*, July 7, 2025, <https://www.anthropic.com/news/the-need-for-transparency-in-frontier-ai>, accessed July 28, 2025.
- ²⁰ Daniel Zhang, Saurabh Mishra, Erik Brynjolfsson, John Etchemendy, Deep Ganguli, Barbara Grosz, Terah Lyons et al. "The AI index 2021 annual report." *arXiv*, March 9, 2021, <https://doi.org/10.48550/arXiv.2103.06312>.

- 21 "Anthropic," *PitchBook*, August 18, 2025, <https://my.pitchbook.com/profile/466959-97/company/profile>, accessed August 18, 2025.
- 22 Billy Perrigo, "Inside Anthropic, the AI Company Betting That Safety Can Be a Winning Strategy," *Time*, May 30, 2024, <https://time.com/6980000/anthropic/>, accessed July 28, 2025.
- 23 Julián Peller, "ChatGPT: Two Years Later", *Medium*, <https://medium.com/data-science/chatgpt-two-years-later-df37b015fd8a>, accessed July 28, 2025.
- 24 "Anthropic Partners with Google Cloud," *Anthropic*, February 3, 2023, <https://www.anthropic.com/news/anthropic-partners-with-google-cloud>, accessed August 11 2025; "Amazon and Anthropic Deepen Their Shared Commitment to Advancing Generative AI," *Amazon News*, March 27, 2024, <https://www.aboutamazon.com/news/company-news/amazon-anthropic-ai-investment>, accessed August 11, 2025.
- 25 Aaron Drapkin, "What Is Claude AI and Anthropic? A Closer Look at ChatGPT's Rival," *Tech.co*, January 2, 2024, <https://tech.co/news/what-is-claude-ai-anthropic>, accessed August 11, 2025.
- 26 "Salesforce Teams Up with Anthropic to Enhance Einstein Capabilities with Claude," *Anthropic*, September 3, 2024, <https://www.anthropic.com/news/salesforce-partnership>, accessed August 11, 2025; "Anthropic Partners with BCG," *Anthropic*, September 14, 2023, <https://www.anthropic.com/news/anthropic-bcg>, accessed August 11, 2025.
- 27 "Anthropic, AWS, and Accenture Team Up to Build Trusted Solutions for Enterprises," *Anthropic*, March 20, 2024, <https://www.anthropic.com/news/accenture-aws-anthropic>, accessed August 11, 2025.
- 28 "Claude, now in Slack," *Anthropic*, March 30, 2023, <https://www.anthropic.com/news/claude-now-in-slack>, and; "Claude 3.5 Sonnet on GitHub Copilot," *Anthropic*, October 29, 2024, <https://www.anthropic.com/news/github-copilot>, accessed August 11, 2025.
- 29 "Claude Statistics: How Many People Use Claude?" *Backlinko*, March 3, 2025, <https://backlinko.com/claude-users>, accessed July 28, 2025.
- 30 TapTwice Digital Team, "7 Anthropic Statistics (2025): Projected Revenue, Market Share, Users, Valuation, Funding," *TapTwice Digital*, May 17, 2025, <https://taptwicedigital.com/stats/anthropic>, accessed July 28, 2025.
- 31 Adam Levine, "Anthropic's Claude 4 Sets New AI Benchmarks Amid ChatGPT Competition," *Barron*, May 22, 2025, <https://www.barrons.com/articles/anthropic-claude-4-chatgpt-ai-eaf1e6a9>, accessed August 11, 2025.
- 32 "Anthropic and the Department of Defense to Advance Responsible AI in Defense Operations," *Anthropic*, July 14, 2025, <https://www.anthropic.com/news/anthropic-and-the-department-of-defense-to-advance-responsible-ai-in-defense-operations>, accessed August 11, 2025.
- 33 "Anthropic raises \$13B Series F at \$183B post-money valuation," *Anthropic*, September 2, 2025, <https://www.anthropic.com/news/anthropic-raises-series-f-at-usd183b-post-money-valuation>, accessed September 3, 2025.
- 34 "Anthropic hits \$3 billion in annualized revenue on business demand for AI," *CNBC*, May 30, 2025, <https://www.cnn.com/2025/05/30/anthropic-hits-3-billion-in-annualized-revenue-on-business-demand-for-ai.html>, accessed August 20, 2025.
- 35 "AI startup Anthropic forecasts revenue growth to \$34.5 billion by 2027- report," *Investing.com*, February 13, 2025, <https://uk.investing.com/news/company-news/ai-startup-anthropic-forecasts-revenue-growth-to-345-billion-by-2027-report-93CH-3922747>, accessed August 20, 2025.
- 36 "Anthropic sees revenue potentially soaring to \$34.5 billion in 2027, The Information reports," *Reuters*, February 12, 2025, <https://www.reuters.com/technology/anthropic-projects-soaring-growth-345-billion-2027-revenue-information-reports-2025-02-13/>, accessed August 20, 2025.
- 37 Ibid.
- 38 "Claude Opus 4", *Anthropic*, May 22, 2025, <https://www.anthropic.com/claude/opus>, accessed July 28, 2025.
- 39 "Claude Opus 4.1," *Anthropic*, August 5, 2025, <https://www.anthropic.com/news/claude-opus-4-1>, accessed August 11, 2025.
- 40 "Claude Sonnet 4," *Anthropic*, February 24, 2025, <https://www.anthropic.com/claude/sonnet>, accessed July 27, 2025.
- 41 "Claude Haiku 3.5," *Anthropic*, October 22, 2024, <https://www.anthropic.com/claude/haiku>, accessed July 27, 2025.

- ⁴² "Claude 3.7 Sonnet and Claude Code," *Anthropic*, February 24, 2025, <https://www.anthropic.com/news/claude-3-7-sonnet>, accessed July 27, 2025.
- ⁴³ "Advancing Claude for Education," *Anthropic*, July 9, 2025, <https://www.anthropic.com/news/advancing-claude-for-education>, accessed July 25, 2025.
- ⁴⁴ "Claude Gov models for U.S. national security customers," *Anthropic*, June 5, 2025, <https://www.anthropic.com/news/claude-gov-models-for-u-s-national-security-customers>, accessed July 28, 2025; "Claude for Financial Services," *Anthropic*, July 15, 2025, <https://www.anthropic.com/news/claude-for-financial-services>, accessed July 28, 2025.
- ⁴⁵ "Building Effective Agents," *Anthropic*, December 19, 2024, <https://www.anthropic.com/engineering/building-effective-agents>, accessed April 5, 2025.
- ⁴⁶ "Make AI agents your unfair advantage," *Anthropic*, August 7, 2025, <https://www.anthropic.com/solutions/agents>, accessed August 10, 2025.
- ⁴⁷ Anthropic, "New Capabilities for Building Agents on the Anthropic API," *Anthropic*, May 22, 2025, <https://www.anthropic.com/news/agent-capabilities-api>, accessed August 11, 2025.
- ⁴⁸ "How We Built Our Multi-Agent Research System," *Anthropic*, June 13, 2025, <https://www.anthropic.com/engineering/multi-agent-research-system>, accessed August 11, 2025.
- ⁴⁹ "Claude Can Now Use Tools," *Anthropic*, May 30, 2024, <https://www.anthropic.com/news/tool-use-ga>; "API Overview," *Anthropic*, August 11, 2025, <https://www.anthropic.com/api>, and; "Claude Code Remote MCP," *Anthropic*, June 18, 2025, <https://www.anthropic.com/news/claude-code-remote-mcp>, accessed August 11, 2025.
- ⁵⁰ "Discover tools that work with Claude," *Anthropic*, July 14, 2025, <https://www.anthropic.com/news/connectors-directory>, accessed July 28, 2025.
- ⁵¹ "Claude Pricing Plans," *Anthropic*, August 11, 2025, <https://www.anthropic.com/pricing>, accessed August 11, 2025.
- ⁵² "Claude Pricing Plans: Team & Enterprise," *Anthropic*, August 11, 2025, <https://www.anthropic.com/pricing#team-&-enterprise>, accessed August 11, 2025.
- ⁵³ "Claude Pricing Plans: API," *Anthropic*, August 11, 2025, <https://www.anthropic.com/pricing#api>, accessed August 11, 2025.
- ⁵⁴ Ruth Appel, Peter McCrory, Alex Tamkin, Miles McCain, Tyler Neylon, and Michael Stern, The Anthropic Economic Index Report: Uneven Geographic and Enterprise AI Adoption, *Anthropic*, September 15, 2025, <https://www.anthropic.com/research/anthropic-economic-index-september-2025-report>, accessed Oct 9, 2025.
- ⁵⁵ AI Engineering, "Anthropic Overtakes OpenAI: The New Landscape of Enterprise LLM Market," *Medium*, August 3, 2025, <https://ai-engineering-trend.medium.com/anthropic-overtakes-openai-the-new-landscape-of-enterprise-llm-market-062898941d87>, accessed October 14, 2025.
- ⁵⁶ "Anthropic Economic Index report: Uneven geographic and enterprise AI adoption," *Anthropic*, September 15, 2025, <https://www.anthropic.com/research/anthropic-economic-index-september-2025-report>, accessed October 14, 2025.
- ⁵⁷ "Building Effective Agents," *Anthropic*, December 19, 2024, <https://www.anthropic.com/engineering/building-effective-agents>, accessed April 5, 2025.
- ⁵⁸ Sri Elaprolu, "The Rise of Autonomous Agents: What Enterprise Leaders Need to Know about the Next Wave of AI," *AWS Insights*, June 13, 2025, <https://aws.amazon.com/blogs/aws-insights/the-rise-of-autonomous-agents-what-enterprise-leaders-need-to-know-about-the-next-wave-of-ai/>, accessed August 21, 2025.
- ⁵⁹ Ibid.
- ⁶⁰ "AI Agents: Components." *MindsDB*, August 29, 2024, <https://www.mindsdb.com/blog/ai-agents-components>, accessed February 23, 2025.
- ⁶¹ Jeremy M. Norman, "History of Information: MYCIN." *History of Information*, <https://www.historyofinformation.com/detail.php?id=3395>, accessed March 26, 2025; Paul Ceruzzi, Museum. "2001: A Space Odyssey, HAL, and the Future of AI." *Smithsonian National Air and Space Museum*, April 10, 2018 <https://airandspace.si.edu/stories/editorial/2001-space-odyssey-hal-and-future-ai>, accessed April 19, 2025.
- ⁶² "The Turing Test," *Stanford Encyclopedia of Philosophy*, April 9, 2003 (revised October 4, 2021), <https://plato.stanford.edu/entries/turing-test/>, accessed July 27, 2025.

- ⁶³ "ELIZA." NJIT, <https://web.njit.edu/~ronkowitz/eliza.html>, accessed April 2, 2025.
- ⁶⁴ Anna Kocsis, "When Did AI Agents Become a Thing? The History & Evolution of Agentic AI," *Mindset AI*, February 11, 2025, <https://www.mindset.ai/blogs/how-have-ai-agents-evolved-over-time>, accessed August 11, 2025.
- ⁶⁵ Ostadick, Nick, "The Evolution of Robotic Process Automation (RPA): Past, Present, and Future." *UiPath*, July 26, 2016, <https://www.uipath.com/blog/rpa/the-evolution-of-rpa-past-present-and-future>, and; "Robotic Process Automation." *UiPath*, <https://www.uipath.com/rpa/robotic-process-automation>, accessed March 30, 2025.
- ⁶⁶ Cameron Hashemi-Pour, "AI Readiness: What is it, and is your business ready?" *Tech Target*, December 3, 2024, <https://www.techtarget.com/searchenterpriseai/definition/AI-readiness>, accessed July 27, 2025.
- ⁶⁷ "Building Effective Agents." *Anthropic*, <https://www.anthropic.com/engineering/building-effective-agents>, accessed April 5, 2025.
- ⁶⁸ "Foundation Models: What's Next for 2025 and Beyond?" *Gradient Flow*, <https://www.gradientflow.com/foundation-models-whats-next-for-2025-and-beyond/>, accessed March 30, 2025.
- ⁶⁹ Rick Merritt, "What Is Retrieval-Augmented Generation, aka RAG?" *NVIDIA*, January 31, 2025, <https://blogs.nvidia.com/blog/what-is-retrieval-augmented-generation/>, accessed August 21, 2025.
- ⁷⁰ Ivan Belcic and Cole Stryker, "What Is Agentic RAG?" *IBM*, March 3, 2025, <https://www.ibm.com/think/topics/agentic-rag>, accessed August 21, 2025.
- ⁷¹ "What Is Agentic Orchestration?" *UiPath*, <https://www.uipath.com/ai/what-is-agentic-orchestration>, accessed April 18, 2025.
- ⁷² "Implementing Effective Guardrails for AI Agents." *GitLab*, <https://www.about.gitlab.com/the-source/ai/implementing-effective-guardrails-for-ai-agents/>, accessed May 1, 2025.
- ⁷³ Ricky Phyllis, "Mitigate High Latency of Generative AI," *Medium*, December 20, 2024, <https://medium.com/@rickyphyllis/mitigate-high-latency-of-generative-ai-cc0074d18420>, accessed October 13, 2025.
- ⁷⁴ "AI in Data Quality Management 2025." *Techment Accessed*, <https://www.techment.com/ai-in-data-quality-management-2025/>, accessed April 9, 2025.
- ⁷⁵ "AI Agent Development Cost." *PixelBrainy*, <https://www.pixelbrainy.com/blog/ai-agent-development-cost>, accessed February 24, 2025.
- ⁷⁶ "Measuring ROI for AI Agents." *Aquiva Labs*, <https://www.aquivalabs.com/measuring-roi-for-ai-agents/>, and; "AI Agents: From AI to ROI." *Biztory*, <https://www.biztory.com/blog/ai-agents-from-ai-to-roi>, accessed April 1, 2025.
- ⁷⁷ "Legal Liability for AI-Driven Decisions: When AI Gets It Wrong, Who Can You Turn To?" *HFW*, <https://www.hfw.com/insights/legal-liability-for-ai-driven-decisions-when-ai-gets-it-wrong-who-can-you-turn-to/>, accessed February 14, 2025.
- ⁷⁸ "Why AI Still Needs Regulation Despite Its Impact." *Thomson Reuters*, <https://legal.thomsonreuters.com/blog/why-ai-still-needs-regulation-despite-impact/>, accessed April 27, 2025.
- ⁷⁹ "China's Harsh Education Crackdown Sends Parents, Businesses Scrambling." *CNBC*, August 5, 2021, <https://www.cnbc.com/2021/08/05/chinas-harsh-education-crackdown-sends-parents-businesses-scrambling.html>, accessed March 23, 2025.
- ⁸⁰ "AI Agents and Jobs: The Impact of Automation." *Fortune*, April 1, 2025, <https://fortune.com/2025/04/01/ai-agents-jobs-automation/>, accessed April 6, 2025.
- ⁸¹ "New Research Uncovers Top Challenges in Enterprise AI Agent Adoption." *Architecture and Governance*, <https://www.architectureandgovernance.com/artificial-intelligence/new-research-uncovers-top-challenges-in-enterprise-ai-agent-adoption/>, accessed March 27, 2025.
- ⁸² "What Is Enterprise AI?" *IBM*, <https://www.ibm.com/think/topics/enterprise-ai>, accessed October 9, 2025.
- ⁸³ "Small Business Productivity Trends 2024." *Salesforce*, <https://www.salesforce.com/news/stories/small-business-productivity-trends-2024/>, Accessed May 2, 2025.
- ⁸⁴ "Cost-Effective AI Solutions for Small Businesses." *Phaedra Solutions*, <https://www.phaedrasolutions.com/blog/cost-effective-ai-solutions-for-small-businesses>, accessed April 25, 2025.
- ⁸⁵ AI in UI/UX Design. <https://www.appventurez.com/blog/ai-ui-ux>, accessed March 20, 2025.

- ⁸⁶ "AI in UI/UX Design." *AppVenturez*, <https://www.appventurez.com/blog/ai-ui-ux>, accessed March 20, 2025.
- ⁸⁷ Brian Caulfield, "CES 2025: Jensen Huang Unveils the Future of AI and Accelerated Computing," *NVIDIA*, January 6, 2025, <https://blogs.nvidia.com/blog/ces-2025-jensen-huang/>, accessed August 11, 2025.
- ⁸⁸ Anna Kocsis, "When Did AI Agents Become a Thing? The History & Evolution of Agentic AI," *Mindset AI*, February 11, 2025, <https://www.mindset.ai/blogs/how-have-ai-agents-evolved-over-time>, accessed August 11, 2025.
- ⁸⁹ Bastian Nominacher, Kiva Allgood, "How to choose the right AI agent for your business," *World Economic Forum*, May 32, 2025, <https://www.weforum.org/stories/2025/05/ai-agents-select-the-right-agent/>, accessed August 11, 2025.
- ⁹⁰ *Ibid.*
- ⁹¹ Anna Gutowska, "What Is a Multi-Agent System?" *IBM*, August 6, 2024, <https://www.ibm.com/think/topics/multiagent-system>, accessed July 25, 2025.
- ⁹² "Navigating the AI Frontier: A Primer on the Evolution and Impact of AI Agents," *World Economic Forum*, December 16, 2024, <https://www.weforum.org/publications/navigating-the-ai-frontier-a-primer-on-the-evolution-and-impact-of-ai-agents/>, accessed July 27, 2025.
- ⁹³ "Private fixed investment in information processing equipment and software / Gross Domestic Product," *Federal Reserve Bank of St. Louis*, <https://fred.stlouisfed.org/graph/?g=GXc>, accessed October 13, 2025.
- ⁹⁴ "Dramatic increase in token consumption #924," <https://github.com/sst/opencode/issues/924>, accessed September 2025.
- ⁹⁵ Alexander Belokrylov, "How Companies Can Monetize Open-Source Software," *Forbes*, May 31, 2022, <https://www.forbes.com/councils/forbestechcouncil/2022/05/31/how-companies-can-monetize-open-source-software/>, accessed October 10, 2025.
- ⁹⁶ Stephanie Susnjara and Ian Smalley, "What Is Open Source Software?," *IBM*, February 5, 2025, <https://www.ibm.com/think/topics/open-source>, accessed October 10, 2025.
- ⁹⁷ D³ Insights & Experiences and D³ Affiliate Research, "Revealing Value: The Economic Power of Open Source Software," *Digital Data Design Institute at Harvard*, October 31, 2024, <https://d3.harvard.edu/revealing-value-the-economic-power-of-open-source-software/>, accessed October 10, 2025.
- ⁹⁸ "Business Value of Red Hat Partners," *Red Hat*, <https://connect.redhat.com/en/partner-with-us/business-value-of-red-hat-partners>, accessed October 10, 2025; "Android App Monetization Strategies That Work in 2025," *Infatica SDK*, July 11, 2025, <https://infatica-sdk.io/blog/app-monetization/android-app-monetization/>, accessed October 10, 2025.
- ⁹⁹ Melissa Malec, "Open-Source LLMs vs Closed: Unbiased Guide for Innovative Companies [2025]," *HatchWorks*, updated August 4, 2025, <https://hatchworks.com/blog/gen-ai/open-source-vs-closed-llms-guide/#case-closed-source>, accessed August 30, 2025.
- ¹⁰⁰ Aruna Kolluru, "Exploring the World of Open Source and Open Weights AI," *Medium*, March 29, 2024, <https://medium.com/@aruna.kolluru/exploring-the-world-of-open-source-and-open-weights-ai-aa09707b69fc>, accessed August 30, 2025.
- ¹⁰¹ Jean-Stanislas Denain, "Models with Downloadable Weights Currently Lag Behind the Top-Performing Models," *Epoch AI*, November 27, 2024, <https://epoch.ai/data-insights/open-vs-closed-model-performance>, accessed October 14, 2025.
- ¹⁰² Barak Or, "Open vs. Closed LLMs in 2025: Strategic Tradeoffs for Enterprise AI," *Medium*, June 6, 2025, <https://medium.com/data-science-collective/open-vs-closed-llms-in-2025-strategic-tradeoffs-for-enterprise-ai-668af30bffa0>, accessed October 14, 2025.
- ¹⁰³ <https://www.semianalysis.com/p/google-we-have-no-moat-and-neither>, accessed July 18, 2023.
- ¹⁰⁴ "How DeepSeek and Open-Source AI Models Are Disrupting Big Tech," *Bloomberg*, August 6, 2025, <https://www.bloomberg.com/news/articles/2025-08-06/how-deepseek-and-open-source-ai-models-are-disrupting-big-tech>,
- ¹⁰⁵ "The Foundation Model Divide: Mapping the Future of Open vs Closed AI Development," *CB Insights*, January 8, 2025, <https://www.cbinsights.com/research/report/future-of-foundation-models-open-source-closed-source/>, accessed August 30, 2025.
- ¹⁰⁶ "gpt-oss-120b & gpt-oss-20b Model Card," *OpenAI*, August 5, 2025, <https://openai.com/index/gpt-oss-model-card/>, accessed September 1, 2025.

- ¹⁰⁷ "Models Overview," *Mistral AI*, https://docs.mistral.ai/getting-started/models/models_overview/, accessed September 1, 2025.
- ¹⁰⁸ Joshua Toenyes, "What's Better? Open-Source versus Closed-Source AI," *Aquent*, April 18, 2025, <https://aquent.com/blog/whats-better-open-source-versus-closed-source-ai>, accessed September 1, 2025.
- ¹⁰⁹ "Introducing Operator," *OpenAI*, January 23, 2025, <https://openai.com/index/introducing-operator/>, accessed August 11, 2025.
- ¹¹⁰ "Introducing ChatGPT agent: bridging research and action," *OpenAI*, July 17, 2025, <https://openai.com/index/introducing-chatgpt-agent/>, accessed August 11, 2025.
- ¹¹¹ "ChatGPT Agent," *OpenAI*, August 8, 2025, <https://help.openai.com/en/articles/11752874-chatgpt-agent>, accessed August 11, 2025.
- ¹¹² "New Tools for Building Agents," *OpenAI*, March 11, 2025, <https://openai.com/index/new-tools-for-building-agents>, accessed August 11, 2025; "New Tools and Features in the Responses API," *OpenAI*, May 21, 2025, <https://openai.com/index/new-tools-and-features-in-the-responses-api>, accessed August 11, 2025.
- ¹¹³ Wendy Xu, "Claude vs ChatGPT Agent: What Australian Banks Need to Know," *Medium*, August 2, 2025, <https://medium.com/@xu0321wen/claude-vs-chatgpt-agent-what-australian-banks-need-to-know-58e7dc6bbd90>, accessed August 10, 2025.
- ¹¹⁴ "OpenAI lets ChatGPT users connect with Spotify, Zillow in-app," *Bloomberg*, October 6, 2025, <https://www.bloomberg.com/news/articles/2025-10-06/openai-lets-chatgpt-users-connect-with-spotify-zillow-in-app?embedded-checkout=true&agent>, accessed October 14, 2025.
- ¹¹⁵ "Introducing apps in ChatGPT and the new Apps SDK," *OpenAI*, October 6, 2025, <https://openai.com/index/introducing-apps-in-chatgpt/>, accessed October 14, 2025.
- ¹¹⁶ Adam Hales, "'Unstoppable?' ChatGPT Surges to 700 million Weekly Users as Rivals Race to Compete," *Windows Central*, August 5, 2025, <https://www.windowscentral.com/artificial-intelligence/chatgpt-is-set-to-hit-700-million-weekly-users-but-can-its-rivals-catch-up>, accessed August 11, 2025.
- ¹¹⁷ Barry Elad, "OpenAI Statistics 2025: Adoption, Integration & Innovation," *SQ Magazine*, July 22, 2025, <https://sqmagazine.co.uk/openai-statistics/>, accessed August 11, 2025.
- ¹¹⁸ Sri Muppidi, "OpenAI Hits \$12 Billion Annualized Revenue, Breaks 700 Million ChatGPT Weekly Active Users," *The Information*, July 31, 2025, <https://www.theinformation.com/articles/openai-hits-12-billion-annualized-revenue-breaks-700-million-chatgpt-weekly-active-users>, accessed August 21, 2025.
- ¹¹⁹ Ibid.
- ¹²⁰ Ece Yildirim, "OpenAI Thinks Its Revenue Will More Than Triple in 2025, Report Says," *Yahoo Finance*, March 26, 2025, <https://finance.yahoo.com/news/openai-thinks-revenue-more-triple-210900736.html>, accessed August 21, 2025.
- ¹²¹ "Gemini 2.5 Pro," *Google DeepMind*, <https://deepmind.google/models/gemini/pro/>, accessed August 8, 2025.
- ¹²² SIMA Team, "A Generalist AI Agent for 3D Virtual Environments," *Google DeepMind*, March 13, 2024, <https://deepmind.google/discover/blog/sima-generalist-ai-agent-for-3d-virtual-environments/>, accessed August 11, 2025.
- ¹²³ AlphaEvolve team, "AlphaEvolve: A Gemini-Powered Coding Agent for Designing Advanced Algorithms," *Google DeepMind*, May 14, 2025, <https://deepmind.google/discover/blog/alphaevolve-a-gemini-powered-coding-agent-for-designing-advanced-algorithms/>, accessed August 11, 2025.
- ¹²⁴ Juraj Gottweis and Vivek Natarajan, "Accelerating Scientific Breakthroughs with an AI Co-Scientist," *Google Research*, February 20, 2025, <https://research.google/blog/accelerating-scientific-breakthroughs-with-an-ai-co-scientist/>, accessed August 8, 2025.
- ¹²⁵ "Building the Industry's Best Agentic AI Ecosystem with Partners," *Google Cloud*, April 9, 2025, <https://cloud.google.com/blog/topics/partners/best-agentic-ecosystem-helping-partners-build-ai-agents-next25>, accessed August 1, 2025.
- ¹²⁶ Rao Surapaneni, Miku Jha, Michael Vakoc, and Todd Segal, "Announcing the Agent2Agent Protocol (A2A)," *Google Developers*, April 9, 2025, <https://developers.googleblog.com/en/a2a-a-new-era-of-agent-interoperability/>, accessed August 11, 2025.

- ¹²⁷ Kevin Ichhpurani, "Building the Industry's Best Agentic AI Ecosystem with Partners," *Google Cloud*, April 9, 2025, <https://cloud.google.com/blog/topics/partners/best-agentic-ecosystem-helping-partners-build-ai-agents-next25>, accessed August 11, 2025.
- ¹²⁸ "The Llama 4 herd: The beginning of a new era of natively multimodal AI innovation," *Meta*, April 5, 2025, <https://ai.meta.com/blog/large-language-model-llama-meta-ai/>, accessed August 5, 2025.
- ¹²⁹ "Zuckerberg signals Meta won't open source all of its 'superintelligence' AI models, July 30, 2025, <https://techcrunch.com/2025/07/30/zuckerberg-says-meta-likely-wont-open-source-all-of-its-superintelligence-ai-models/>, accessed August 27, 2025.
- ¹³⁰ "Meta Introduces a Customer Service AI Agent for Facebook, Instagram, Targets SMBs with Free Trial," *CX Today*, March 7, 2025, <https://www.cxtoday.com/conversational-ai/meta-introduces-a-customer-service-ai-agent-for-facebook-instagram-targets-smb-s-with-free-trial/>, accessed July 28, 2025.
- ¹³¹ "Meta's AI Products Just Got Smarter and More Useful," *Meta*, September 25, 2024, <https://about.fb.com/news/2024/09/met-as-product-news-connect/>, accessed July 28, 2025.
- ¹³² Ryan Deffenbaugh, "Meta Stock Falls as AI, Metaverse Questions in Focus," August 18, 2025, <https://www.investors.com/news/technology/meta-stock-ai-smart-glasses-gen-ai-concerns/>, accessed August 27, 2025.
- ¹³³ Kurt Wagner, "Zuckerberg Announces Meta Superintelligence Effort, More Hires," *Bloomberg*, June 30, 2025, <https://www.bloomberg.com/news/articles/2025-06-30/zuckerberg-announces-meta-superintelligence-effort-more-hires>, accessed August 21, 2025.
- ¹³⁴ Beatrice Nolan, "Meta's \$100 million signing bonuses for OpenAI staff are just the latest sign of extreme AI talent war," *Fortune*, June 18, 2025, <https://fortune.com/2025/06/18/met-as-100-million-signing-bonuses-openai-staff-extreme-ai-talent-war/>, accessed December 2, 2025.
- ¹³⁵ Jonathan Vanian, "Mark Zuckerberg Creating Meta Superintelligence Labs: Read the Memo," *CNBC*, June 30, 2025, <https://www.cnbc.com/2025/06/30/mark-zuckerberg-creating-meta-superintelligence-labs-read-the-memo.html>, accessed August 21, 2025.
- ¹³⁶ "Anthropic AI staff reject Meta's \$100M offer to stay loyal to company culture, August 4, 2025, <https://www.ainvest.com/news/anthropic-ai-staff-reject-meta-100m-offers-stay-loyal-company-culture-2508/>, accessed August 27, 2025.
- ¹³⁷ Zoe Schiffer and Will Knight, "Researchers are already leaving Meta's New Superintelligence Lab," August 26, 2025, <https://www.wired.com/story/researchers-leave-meta-superintelligence-labs-openai/>, accessed August 27, 2025.
- ¹³⁸ Chris Metinko, "Agentic AI is angelic to investors," July 25, 2025, <https://www.axios.com/pro/enterprise-software-deals/2025/07/25/agentic-ai-vc-investment>, accessed August 27, 2025.
- ¹³⁹ "The Web Outgrew the Browser," *TinyFish*, August 20, 2025, <https://blog.tinyfish.ai/the-web-outgrew-the-browser-98a464b5608a>, accessed August 26, 2025.
- ¹⁴⁰ "Decagon Secures \$131 Million in Series C Funding," *StartupHub.ai*, June 24, 2025, <https://www.startuphub.ai/ai-news/funding-round/2025/decagon-secures-131-million-in-series-c-funding/>, accessed August 26, 2025.
- ¹⁴¹ Nick Patience, "Decagon Aims to Simplify AI Agent Development with Natural Language," *Futurum*, August 14, 2025, <https://futurumgroup.com/press-release/decagon-aims-to-simplify-ai-agent-development-with-natural-language/>, accessed August 26, 2025.
- ¹⁴² ThriveAI, <https://www.thriveai.pm/>, accessed August 26, 2025.
- ¹⁴³ Samantha Stokes, "Here's an Exclusive Look at the Pitch Deck Palantir and Google Alums Used to Raise \$1.2 Million to Build an AI Agent That Acts Like a Junior Product Manager," *Business Insider*, May 13, 2025, <https://www.businessinsider.com/pitch-deck-ai-product-management-startup-thriveai-preseed-round-2025-5>, accessed August 26, 2025.
- ¹⁴⁴ Melia Russell, "LexisNexis Is Teaming Up with Harvey. Here's Why That Matters for the Legal Tech Market," *Business Insider*, June 18, 2025, <https://www.businessinsider.com/lexisnexis-harvey-form-strategic-alliance-amid-growing-competition-2025-6>, accessed August 26, 2025.
- ¹⁴⁵ Aidan Gomez, "Introducing Command R+: A Scalable LLM Built for Business," *Cohere*, April 4, 2024, <https://cohere.com/blog/command-r-plus-microsoft-azure>, accessed July 29, 2025.

- ¹⁴⁶ Zachary Cefaratti and Dalma Capital, "Cohere AI: Well Positioned for the Coming Wave of Enterprise AI Application and Agentic AI. Why We Continue to Invest," *Medium*, March 28, 2025, <https://medium.com/@xu0321wen/cohere-ai-well-positioned-for-the-coming-wave-of-enterprise-ai-application-and-agentic-ai-8cc52ab02ac8>, accessed July 28, 2025.
- ¹⁴⁷ "Build AI Agents with the Mistral Agents API," *Mistral AI*, May 27, 2025, <https://mistral.ai/news/agents-api>, and; "Introducing Le Chat Enterprise," *Mistral AI*, May 7, 2025, <https://mistral.ai/news/le-chat-enterprise>, accessed July 29, 2025.
- ¹⁴⁸ Sam Schechner, "Fears Over U.S. AI Dominance Boost Business for France's Mistral," *Wall Street Journal*, June 11, 2025, <https://www.wsj.com/tech/ai/fears-over-u-s-ai-dominance-boost-business-for-frances-mistral-14405419>, accessed July 2025.
- ¹⁴⁹ Abdullah Farag, "Building Intelligent AI Agents with DeepSeek Models," *Medium*, January 28, 2025, https://medium.com/@a_farag/building-intelligent-ai-agents-with-deepseek-models-f9a78439e717, accessed August 21, 2025.
- ¹⁵⁰ Mehran Maghousi, Chintan Patel, and Chris Alexiuk, "Build an AI Agent with Expert Reasoning Capabilities Using the DeepSeek-R1 NIM," *NVIDIA*, February 28, 2025, <https://developer.nvidia.com/blog/build-ai-agents-with-expert-reasoning-capabilities-using-deepseek-r1-nim/>, accessed August 21, 2025.
- ¹⁵¹ "Tencent Launches Intelligent Agent Development Platform, Plans to Open Source Multiple AI Models," *TMTPOST*, May 21, 2025, <https://en.tmtpost.com/news/7572679>, and; 36Kr English, "Tencent Ups Its AI Game with a Flurry of Releases at WAIC 2025," *KrASIA*, July 28, 2025, https://kr-asia.com/tencent-ups-its-ai-game-with-a-flurry-of-releases-at-waic-2025?utm_source=chatgpt.com, accessed July 29, 2025.
- ¹⁵² "AI Arena: Open Research Platform for Multi-Agent Intelligence," *Tencent AI Arena*, <https://aiarena.tencent.com/aiarena/en>, accessed August 8, 2025.
- ¹⁵³ Asif Razzaq, "ByteDance Just Released Trae Agent: An LLM-Based Agent for General-Purpose Software Engineering Tasks," *MarkTechPost*, July 7, 2025, <https://www.marktechpost.com/2025/07/07/bytedance-just-released-trae-agent-an-llm-based-agent-for-general-purpose-software-engineering-tasks/>, accessed July 28, 2025; "ByteDance Open-Sources AI Agent Platform 'Coze' to Empower Intelligent Agent Development," *Albase*, August 11, 2025, <https://news.aibase.com/news/19991>, accessed August 11, 2025.
- ¹⁵⁴ Molly Hayes and Amanda Downie, "AI Agent Use Cases," *IBM*, April 15, 2025, <https://www.ibm.com/think/topics/ai-agent-use-cases>, accessed August 22, 2025.
- ¹⁵⁵ Erik Pounds, "What Is Agentic AI?" *NVIDIA*, October 22, 2024, <https://blogs.nvidia.com/blog/what-is-agentic-ai/>, accessed August 22, 2025.
- ¹⁵⁶ Keith Kirkpatrick, "Are Today's AI Agents Actually Agentic?," *The Futurum Group*, July 25, 2025, <https://futurumgroup.com/press-release/are-todays-ai-agents-actually-agentic/>, accessed August 22, 2025.
- ¹⁵⁷ "Gartner Predicts Over 40 Percent of Agentic AI Projects Will Be Canceled by End of 2027," *Gartner*, June 25, 2025, <https://www.gartner.com/en/newsroom/press-releases/2025-06-25-gartner-predicts-over-40-percent-of-agentic-ai-projects-will-be-canceled-by-end-of-2027>, accessed August 24, 2025.
- ¹⁵⁸ Ibid.
- ¹⁵⁹ Sheryl Estrada, "MIT Report: 95 Percent of Generative AI Pilots at Companies Are Failing," *Fortune*, August 18, 2025, <https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/>, accessed August 22, 2025.
- ¹⁶⁰ "The Biggest AI Agent Adoption Challenges and How to Overcome Them. Accessed." *Digital Workforce*, <https://www.digitalworkforce.com/rpa-news/the-biggest-ai-agent-adoption-challenges-and-how-to-overcome-them/>, April 8, 2025.
- ¹⁶¹ "AI Adoption Challenges." *IBM*, <https://www.ibm.com/think/insights/ai-adoption-challenges>, accessed March 5, 2025.
- ¹⁶² "Trust Emerges as Main Barrier to Agentic AI Adoption in Finance and Accounting," *Deloitte*, July 29, 2025, <https://www.deloitte.com/us/en/about/press-room/trust-main-barrier-to-agentic-ai-adoption-in-finance-and-accounting.html>, accessed August 22, 2025.
- ¹⁶³ IBM, "AI Hallucinations." *IBM*, <https://www.ibm.com/think/topics/ai-hallucinations>, accessed March 12, 2025.
- ¹⁶⁴ "O3 and O4 Mini System Card." *OpenAI*, <https://cdn.openai.com/pdf/2221c875-02dc-4789-800b-e7758f3722c1/o3-and-o4-mini-system-card.pdf>, accessed March 18, 2025.

- ¹⁶⁵ “C-Suite Research,” *Salesforce*, May 13, 2024, <https://www.salesforce.com/news/stories/c-suite-research/>, and; “Why Agents Are the Next Frontier of Generative AI,” *McKinsey & Company*, July 24, 2024, <https://www.mckinsey.com/capabilities/mckinsey-digital/our-insights/why-agents-are-the-next-frontier-of-generative-ai>, accessed April 5, 2025.
- ¹⁶⁶ Holt Hackney, “New Research Uncovers Top Challenges in Enterprise AI Agent Adoption,” *Architecture & Governance*, January 20, 2025, <https://www.architectureandgovernance.com/artificial-intelligence/new-research-uncovers-top-challenges-in-enterprise-ai-agent-adoption/>, accessed August 22, 2025.
- ¹⁶⁷ Kim, Jiin, Byeongjun Shin, Jinha Chung, and Minsoo Rhu, “The Cost of Dynamic Reasoning: Demystifying AI Agents and Test-Time Scaling from an AI Infrastructure Perspective,” *arXiv* preprint arXiv:2506.04301 (2025), <https://arxiv.org/abs/2506.04301>, accessed December 2025.
- ¹⁶⁸ Anthonette Adanyin, “Ethical AI in Retail: Consumer Privacy and Fairness,” *arXiv* (preprint), October 20, 2024, <https://arxiv.org/abs/2410.15369>, accessed August 22, 2025.
- ¹⁶⁹ Julian De Freitas, “Why People Resist Embracing AI,” *Harvard Business Review*, January 1, 2025, <https://hbr.org/2025/01/why-people-resist-embracing-ai>, accessed August 22, 2025.
- ¹⁷⁰ Lee Rainie, Cary Funk, Monica Anderson, and Alec Tyson, “Americans Cautious about the Deployment of Driverless Cars,” *Pew Research Center*, March 17, 2022, <https://www.pewresearch.org/internet/2022/03/17/americans-cautious-about-the-deployment-of-driverless-cars/>, accessed August 22, 2025.
- ¹⁷¹ Andrey Fradkin, “Demand for LLMs: Descriptive Evidence on Substitution, Market Expansion, and Multi-Homing,” *arXiv*, April 21, 2025, <https://arxiv.org/abs/2504.15440>, accessed August 22, 2025.
- ¹⁷² Sarah Wang and Shangda Xu, “16 Changes to the Way Enterprises Are Building and Buying Generative AI,” *Andreessen Horowitz*, March 21, 2024, <https://a16z.com/generative-ai-enterprise-2024/>, accessed August 22, 2025.
- ¹⁷³ Lavanya Gupta, “Swapping LLMs Isn’t Plug-and-Play: Inside the Hidden Cost of Model Migration,” *VentureBeat*, April 16, 2025, <https://venturebeat.com/ai/swapping-llms-isnt-plugin-and-play-inside-the-hidden-cost-of-model-migration/>, accessed August 22, 2025.
- ¹⁷⁴ “Switching LLM Providers: Why It’s Harder Than It Seems,” *Requesty*, January 15, 2025, <https://www.requesty.ai/blog/switching-llm-providers-why-it-s-harder-than-it-seems>, accessed August 22, 2025.
- ¹⁷⁵ Scott King, “Avoiding LLM Lock-In: Why Enterprises Need ‘Hot Swap’ Capabilities,” *Krista AI*, February 3, 2025, <https://krista.ai/avoiding-llm-lock-in-why-enterprises-need-hot-swap-capabilities/>, accessed August 24, 2025.
- ¹⁷⁶ Hitesh MS, “Can You Switch Between LLMs??” *Medium*, October 25, 2024, <https://medium.com/@hitesh.ms24/can-you-switch-between-llms-09d41040ca92>, accessed August 24, 2025.
- ¹⁷⁷ Maxim Saplin, “Convergence of LLMs: 2024 Trend Solidified by Llama 3.1 Release,” *DEV*, July 25, 2024, <https://dev.to/maximsaplin/convergence-of-llms-2024-trend-solidified-by-llama-31-release-3h54>, accessed August 24, 2025.
- ¹⁷⁸ Victor Holmin, “Breaking Through the Performance Plateau: Rethinking LLM Strategies for Enterprises,” *LinkedIn*, November 18, 2024, <https://www.linkedin.com/pulse/breaking-through-performance-plateau-rethinking-llm-victor-r2are/>, accessed August 24, 2025.
- ¹⁷⁹ Sarah Lee, “Convergence in Antitrust Law,” *Number Analytics*, June 24, 2025, <https://www.numberanalytics.com/blog/ultimate-guide-convergence-antitrust-law>, accessed August 24, 2025.
- ¹⁸⁰ Steven Levy, “If Anthropic succeeds, a nation of benevolent AI geniuses could be born,” March 28, 2025, https://www.wired.com/story/anthropic-benevolent-artificial-intelligence/?utm_campaign=article_email&utm_content=article-15661&utm_medium=email&utm_source=sg, accessed August 27, 2025.