

CASE PRE-READ 1 OF 1 • CAPSULE 3

It works — but should you build it? Anthropic's Next Step: From LLM to Agents?

Yoffie, D., Wei, A. Z., Modi, M., Jain, D., Lad, S. "Anthropic's Next Step: From LLM to Agents?" Harvard Business School case 9-726-407, rev. 9 May 2026 (29 pp., developed from published sources). Licensed to PE6201 for Jul 2026 – Jan 2027. We brief and discuss this **live in Capsule 3** — this note prepares you for that discussion; it does not replace the case.

ASSIGNED READING

Read the case itself — your own licensed copy — with particular attention to the sections **"AI Agents"**, **"AI Readiness"**, **"Barriers to Adoption"** and **"Anthropic at a Crossroads."** Study **Exhibits 1, 2, 5 and 7** (funding and valuation; August 2025 API prices; inference cost for 2 million tokens; barriers to adoption). Come with a view on the closing question — *what should Dario Amodei do?* — and with the reason you hold it. (≈ 45 min)

» case ~45 min • this brief ~20 min

Licence. The case is copyright President and Fellows of Harvard College and is licensed for **this course only, July 2026 – January 2027**. Do not forward, post or reproduce it, or this brief, outside PE6201. This brief carries only figures printed in the case body and is a study aid, not a substitute for reading it.

KEY TAKEAWAY

Week 4 spends two capsules establishing that an agent is a model plus tools plus a loop, and that most things should not be one. This case asks the same question with a company attached. In late 2025 Anthropic is winning the enterprise **model** market — a claimed 32% share, US\$3bn in annualised recurring revenue, a US\$183bn post-money valuation — and losing time in the **agent** market, where OpenAI, Google, Microsoft and Salesforce have already shipped products. Dario Amodei has to choose whether to move first and accept risk, or move deliberately and accept the loss; whether to sell to consumers, small businesses, enterprises or government; whether to build the full stack, partner, or specialise vertically in coding. Underneath every version of the choice sits one tension the case never resolves for you: **Anthropic is a Public Benefit Corporation whose safety-first identity is simultaneously its strongest differentiator in regulated markets and the thing most likely to make it late.** The economics reading you did in Class 3 applies directly — this is the same "it works; can it pay?" question, asked one layer up, about a company rather than a query.

WHY IT MATTERS FOR THIS COURSE

The five foundation pre-reads give you the engineering view of agents: patterns, loops, failure modes, evals, prices. This case supplies the strategic view, and it is the one your **A2** and your end-of-course project ultimately answer to — someone has to decide whether the thing is worth building, for whom, and at what risk. It is also a rare chance to test the doctrine against its author: Reading 1 is Anthropic's own published argument for restraint, and this case puts that company in a market that rewards speed. Whether restraint is a moat or a handicap is a live question, not a settled one, and you will be arguing both sides in class.

Feeds Class 4: C3 (agent economics • the live case block) • and forward to Week 5 (strategy, use-case selection, build-vs-buy at scale) and Week 6 (governance as a commercial position).

MAIN POINTS

Every figure, quotation and option below is from the case itself. Anything marked **Our reading** is this brief's interpretation — the arithmetic to run, the tensions to notice and the links back to the pre-reads are ours, not Yoffie's.

The company, in numbers the case prints

- **Structure and identity.** Founded 2021; a **Public Benefit Corporation** with a charter binding it to humanity's long-term wellbeing, and a **Long-Term Benefit Trust** of five independent trustees with authority to appoint board members. **Constitutional AI** — training a model against written principles, with critique by other models rather than human labels — is both the technical method and the brand.
- **Position in 2025.** Claude 4 released May 2025; Opus 4 sets the bar in software development and autonomous task execution; Anthropic passes OpenAI in enterprise AI with a claimed **32% share**. Valuation runs from **US\$60bn (March 2025)** to a **US\$13bn Series F in September 2025 at US\$183bn post-money**. Over 1,000 employees; a two-year **US\$200m** US Department of Defense prototype agreement.
- **The money — read this pair together.** **US\$3bn ARR** by end-May 2025 (against an internal projection of US\$2.2bn, and just under US\$1bn a year earlier), with cash burn falling from **US\$5.6bn (2024)** to a planned **US\$3bn (2025)**, of which compute alone is about **US\$2.5bn**. Projected 2027 revenue of **US\$12bn–34.5bn**, with cash-flow breakeven only at the top of that range. **Our reading.** The arithmetic to do before class: triple the

revenue line and hold the compute line, then ask what has to be true about inference cost per task for the two lines to cross.

- **The competitor's version of the same page.** OpenAI: **US\$12bn ARR** by mid-2025 (from US\$4bn in 2024), **~700m weekly users**, ~2.5bn prompts a day, 600,000+ ChatGPT Enterprise seats, >90% of the Fortune 500 — and cash burn of **US\$8bn** in 2025, server rentals potentially above US\$14bn, with cash-flow positive not expected until **2029** at US\$125bn of revenue. Both firms are growing fast and losing money; they differ in scale, in burn, and in what they are buying with it.

What the case means by “agent” — and how far the market actually is

- **The definition is Reading 1's.** The case adopts Anthropic's own: a system in which the model **directs its own processes and tool use**, deciding how to accomplish a task rather than following a fixed sequence — explicitly contrasted with workflows built on rigid, predetermined steps (Exhibit 4).
- **Components.** A reasoning engine (the model), a tool-use and execution interface, a memory system for short- and long-term context, input handling and perception, and a goal-and-planning module. Map these against the six modules in Pre-read 3 and note which one the case does not have: a predictive model of the environment.
- **Autonomy levels, borrowed from self-driving.** Level 4 = independently sets goals, adapts plans, works across domains with no human. The case's assessment of late 2025 is deflating and important: **most “agentic” products are Level 1–2** — fixed sequences or lightly dynamic workflows — with only a few experimental systems at Level 3. Hold that next to the market forecasts below.
- **Readiness has three layers.** External infrastructure (models, secure execution environments, retrieval, and the emerging idea of **agentic RAG** — retrieving, evaluating and combining from several sources and APIs); internal integration (clean accessible data, security policy, deployment cost and ROI discipline, and the fact that an autonomous error can cost more than the efficiency it bought); and the system layer of trust, liability and regulation. **Our reading — Readiness is set by the weakest link**, which is rarely the model.

The market — forecast versus friction, on the same pages

- **The forecast.** Agent market from **US\$5.1bn (2024) to US\$47.1bn (2030)**; **33%** of enterprise software applications expected to include agents by 2028; nearly **90%** of executives seeing agentic AI as a potential source of competitive advantage; over **US\$2bn** already invested in the ecosystem; Jensen Huang's four-stage framing — perception → generative → agentic → physical AI — with the industry now betting on stage three.
- **The friction, from the same case.** **Less than 3%** of firms trust an agent to make judgement calls autonomously; **45%** of business leaders report concerns about data accuracy and bias; hallucination remains a headline blocker, and rates did not fall monotonically with newer reasoning models; successful deployments often need **eight or more data sources**, and over **86%** of organisations need significant stack upgrades; **76%** of Americans are uneasy even about riding in a self-driving car. **Our reading.** Put this bullet next to the one above it. The forecast assumes adoption that the friction says is not yet possible, so something has to give — and the order in which it gives is the interesting part. That is focus question 2.
- **The cost scissors.** Price per token is falling sharply (Exhibits 2 and 5) while the number of tokens advanced reasoning consumes is exploding. You met exactly this in Class 3 with the AI Wars case; here it is the industry-level version, and it is why nobody can yet quote a stable price for “an agent”.
- **Competitive shape.** OpenAI (Operator Jan 2025, the Computer-Using Agent, the ChatGPT agent in July 2025, then in-chat apps in October 2025 — an operating layer rather than a chatbot), Google (Gemini reasoning models, SIMA, AI co-scientist, Vertex AI Agent Builder, 1,000+ announced agent use cases), Microsoft's open agentic web, Salesforce Agentforce — plus open-weight players. US private fixed investment in information-processing equipment and software hit **4.44% of GDP** in Q2 2025, comparable to the dot-com peak in 2000.

Open, closed, and the switching question

- **Anthropic is entirely closed.** No public weights, no training code — justified on safety, control and enterprise deployment. Rivals hedge: OpenAI shipped two open-weight models in August 2025 while keeping its frontier models closed; Meta and DeepSeek publish weights; Mistral splits its range.
- **The moat argument, from the case's own quotation.** The 2023 leaked Google memo — “We Have No Moat” — argues open models are “faster, more customizable, more private, and pound-for-pound more capable”, doing with US\$100 and 13B parameters what large labs struggle to do at US\$10m and 540B. The case also notes open models approaching parity in maths and reasoning, and leading in some verticals. Note what the memo is and is not: an internal argument about open models closing the capability gap, written before agents were the question. The case borrows it to aim at a different target — whether a company that sells only models has a moat at all.
- **Multi-homing and switching costs.** Most Fortune 500 firms use **three to five or more** model providers, rotating on price, latency and fit. But agents raise switching costs: prompts, output formats and context limits differ, so swapping a model can mean re-tuning an entire system — hence the emergence of orchestration layers offering model “hot-swap”. **Our reading.** The strategic tension to notice: if switching is easy, models are a commodity

and the value sits in the agent layer; if switching is hard, whoever owns the agent layer owns the customer. Anthropic's choice of where to play is a bet on which of those is true.

The decision Amodeli faces

- **Timing.** First-mover — launch agent platforms or applications aggressively, accepting higher risk — versus deliberate entry, learning from competitors' failures and arriving later with something safer and more differentiated. The case is pointed about the precedent: Anthropic had a chatbot ready before ChatGPT and did not ship it, and paid for that in market position.
- **Segment.** Consumer (Claude's conversational strength, head-to-head with Operator), SMB (versatile, cost-effective, ROI-led, but price-sensitive and trial-dependent), or enterprise and government (where safety and integration are rewarded and competition is fiercest).
- **Mode of entry.** Build the full stack — maximum control, but Anthropic is loss-making and lacks some of the capability; partner and let others build platforms on Claude — plays to its strengths but cedes the user relationship and the product experience; a hybrid of agentic APIs plus selective vertical partnerships; or release open-weight models.
- **Or specialise.** Concentrate on coding — code generation, debugging, testing, autonomous software development, multi-agent collaboration — and aim to be the standard for safe, capable coding agents rather than a general-purpose agent company. Note how well this fits the evidence in Pre-read 5 about where agents actually work.
- **The tension under all four.** A public benefit mandate against commercial ambition in a market rewarding speed. Safety may be a durable edge in regulated industries and among risk-averse buyers — or it may slow Anthropic into irrelevance. The case ends on the open question and gives you no answer.

READING FOCUS QUESTIONS — bring your answers to class

Bring a position, not a summary. Clues point you at what to compare or compute; there are no worked answers here, because the reasoning is the class. Expect to be asked to defend a choice and then argue the opposite.

- Anthropic reports about US\$3bn of annualised recurring revenue against roughly US\$3bn of planned 2025 cash burn, of which about US\$2.5bn is compute. Using only figures in the case, work out what has to happen for the two lines to cross by 2027 — and say which single assumption your answer is most sensitive to.

Clue. Do the arithmetic before you form a view: take the case's 2027 revenue range of US\$12bn–34.5bn, hold compute flat, then grow it in proportion to usage and see how differently the two versions end. Then find the assumption doing the work — it is not the growth rate, it is what happens to **cost per task** when agentic workloads multiply tokens per task while price per token falls. Cost per task, not price per token, is the number that decides this.

- The case gives you a booming forecast (US\$5.1bn → US\$47.1bn; 33% of enterprise software by 2028) and, a few pages later, the reason it may not happen (under 3% trust autonomous judgement; 86% need stack upgrades; most products are Level 1–2). Reconcile them: what specifically has to become true, and roughly in what order?

Clue. Resist choosing the optimistic or the pessimistic page — both are in the same case for a reason. Build a short chain of preconditions instead, and ask which link is technical, which is organisational and which is regulatory. Notice that trust is downstream of demonstrated reliability, that reliability is downstream of evaluation, and that evaluation is downstream of someone having written down what good looks like — which is the whole of your Capsule 2.

- First-mover versus deliberate entry. Anthropic had a chatbot before ChatGPT and held it back. Argue the case that the same caution is now correct — and then argue that it is precisely the mistake being repeated. Which argument survives better when you apply it to a *regulated* buyer?

Clue. Argue both sides properly before deciding; a position you cannot argue against is not a position. Separate what late entry costs — distribution, defaults, data, mindshare — from what it buys: someone else's failures, a clearer specification, and a safety record. Then change the buyer from a consumer to a bank or a ministry and see which costs and benefits change sign. Your Class-1 and Class-2 cases both had governance-constrained buyers; use them.

- Pick **one** segment (consumer, SMB, enterprise/government) and **one** mode of entry (full stack, partner, hybrid, open weights, or coding-only). Defend the pair as a package, and name the specific capability Anthropic would have to buy or build that it visibly does not have today.

Clue. Judge the pair, not the parts — a consumer play through partners and an enterprise play through the full stack are different companies. For each candidate, ask what the case says about competitive intensity, willingness to pay and integration burden in that segment, then ask which of Anthropic's stated strengths actually transfers. The missing capability is usually not a model; it is distribution, a support organisation, or a product surface.

- Is “safety-first” a moat or a handicap in the agent market? Take a side, and state the observable evidence over the next two years that would prove you wrong.

Clue. A moat has to be something rivals cannot cheaply copy — so first ask whether safety is a **capability**, a **reputation** or a **constraint**, because only one of those is hard to imitate once regulation forces everyone to do it. Then commit to falsifiable evidence: enterprise win rates in regulated verticals, whether procurement standards start naming safety artefacts, whether a rival’s public incident moves share. Pre-read 3 asks what would have to be true for two opposite bets to both be reasonable. This asks the harder version — what would prove you wrong.

- Read the switching-cost passage against the “We Have No Moat” quotation. If models are converging and multi-homing is normal, where does durable value sit in the agent stack — and what does that imply for a company that sells only models?

Clue. Trace the value down the stack layer by layer — weights, inference, orchestration, tools and integrations, workflow ownership, the customer relationship — and ask at each layer what a customer would find painful to replace. Then check your answer against Week 2’s build-vs-buy analysis of the same stack. If you conclude the value has moved up the stack, you have just written the argument for one of Amodeli’s four options, and you should be able to say which.

APPENDIX — CASE VOCABULARY, EXPLAINED

Our reading — The one distinction that carries this case: **capability versus adoption**. Anthropic’s models are demonstrably capable; almost every obstacle in the case is about whether an organisation is ready, willing and legally able to let a system act.

Public Benefit Corporation (PBC) and the Long-Term Benefit Trust (LTBT)

- **Plain.** A company legally required to weigh a public mission alongside profit — with an outside body holding some control of the board.
Technically. Anthropic is chartered as a PBC bound to humanity’s long-term wellbeing; the LTBT is five independent trustees (AI safety, national security, public policy, social enterprise) with authority to appoint board members. Governance is therefore a **structural** constraint on strategy, not merely a stated value.

Constitutional AI

- **Plain.** Train the model against a written set of principles, and have models critique each other, instead of relying on human labelling.
Technically. Reduces dependence on human-labelled examples for behavioural fine-tuning; lets the system explain an objection rather than simply refuse. Presented in the case as both a safety method and a scalable, transparent oversight mechanism.

Annualised recurring revenue (ARR) vs cash burn

- **Plain.** What the current run-rate of subscriptions annualises to, versus how much cash the company consumes to produce it.
Technically. Anthropic: ~US\$3bn ARR (end-May 2025) against ~US\$3bn planned 2025 burn, ~US\$2.5bn of it compute. ARR is a run-rate, not booked revenue — for an AI company the informative pair is ARR **and** cost of inference, because the second scales with the first rather than being amortised away.

Autonomy Levels 1–4

- **Plain.** A self-driving-style ladder for how much an agent decides for itself.
Technically. L1–L2 = fixed rule-based sequences or lightly dynamic workflows; L3 = partial autonomy to plan, execute and adjust with domain tools under minimal oversight; L4 = sets its own goals and works across domains unaided. The case places most late-2025 products at L1–L2 — which is Reading 1’s workflow/agent distinction wearing a market analyst’s clothes.

AI readiness

- **Plain.** Everything other than the model that has to be in place before an agent can safely do real work.
Technically. Three layers: external infrastructure (models, secure execution environments, retrieval); internal integration (clean accessible data, security policy, deployment cost and ROI discipline); and trust, liability and regulation. **Determined by the weakest link** — which is why capability improvements do not translate directly into adoption.

Agentic RAG

- **Plain.** Retrieval that an agent steers — deciding what to look up, judging what came back, and going again.
Technically. Extends single-source retrieval-augmented generation to multiple sources and APIs with evaluation between hops. Early-stage in the case; the grounding mechanism from Pre-read 2, applied at system level.

Multi-homing and switching costs

- **Plain.** Customers use several model providers at once — but moving a whole agent from one to another is much harder than swapping a chat window.

Technically. Most Fortune 500 firms use three to five-plus providers, rotating on price, latency and fit. Agents raise the cost of switching because prompt behaviour, output formats and context limits differ; orchestration layers offering hot-swap exist precisely because that friction is real.

Open source vs open weights vs closed

- **Plain.** Everything published; only the trained weights published; nothing published.

Technically. Anthropic is fully closed. OpenAI, Meta, DeepSeek and Mistral publish weights selectively while holding frontier models back. The case reports open models approaching parity in maths and reasoning and leading in some regulated verticals — the empirical core of the “no moat” argument.

The cost scissors

- **Plain.** Each token is getting cheaper; each task is using far more tokens.

Technically. Exhibits 2 and 5 give price per token and inference cost per 2 million tokens. Falling unit price alongside rising tokens-per-task means cost per completed task — the number that actually matters — can rise while the price list falls. The same structure you priced in the Class-3 case.

SOURCE

Yoffie, D., Wei, A. Z., Modi, M., Jain, D., Lad, S. “**Anthropic’s Next Step: From LLM to Agents?**” HBS case 9-726-407, rev. 9 May 2026. Developed from published sources; funded by Harvard Business School, not by the company. Your licensed copy is on NTU Learn. Companion readings: **Pre-read 1** (Anthropic’s own agent doctrine), **Pre-read 4** (a rival shipping the consumer version), **Pre-read 5** (what the market learned from the last “first AI agent” launch).

CAVEATS

1. This brief is a study aid, not a substitute for the case. It carries only material printed in the case body, and it deliberately omits the exhibits’ detail — you are expected to read them yourself.
2. It gives you no verdict. The options, the tensions and the arithmetic are here; the strategic judgement is the class discussion, and arriving with a position you can defend and then attack is the point.
3. The case is set in late 2025 and revised May 2026. Several facts have moved since — read it as a decision made with the information available then, which is the only fair way to judge any strategic choice. Pre-read 5’s distinction is the tool here: the numbers are calibrations and will be stale; the mechanisms — the cost scissors, readiness set by the weakest link, switching costs deciding where the value sits — are not. We close Capsule 3 on exactly that split.
4. Note the conflict of interest running through the whole week: this summary was generated and fact-checked using Anthropic models, about a case concerning Anthropic. Treat the framing sceptically and say so in class if you think it shows.
5. Licence: copyright President and Fellows of Harvard College, authorised for PE6201 from July 2026 to January 2027. Not to be digitised, reproduced, posted or forwarded outside this course.