

PRE-CLASS READING 2 OF 4

Adoption is not value — 88% are using it, 39% can find it in the P&L

McKinsey & Company / QuantumBlack. “*The state of AI in 2025: Agents, innovation, and transformation*”, published 5 November 2025 — online survey of 1,993 participants across 105 nations, fielded 25 June – 29 July 2025, weighted by each nation’s contribution to global GDP. Read together with its companion, “*State of AI trust in 2026: Shifting to the agentic era*” (25 March 2026; ~500 organisations, fielded December 2025 – January 2026). The OBTL+ outline says “the latest”; these two are what that means as of August 2026.

ASSIGNED READING

Read the **November 2025 survey** in full (≈ 20 min — it is mostly charts, and the charts are the argument). Then skim the **March 2026 trust survey** (≈ 8 min) for what changed once agents arrived. Read them the way you would read any survey: **before you look at a number, find out who was asked and what they were asked to report**. We will spend real class time on exactly that habit, because the three most-quoted numbers about enterprise AI in the last year appear to contradict each other and do not.

» **source ~28 min • this summary + reasoning notes ~20 min**

Note on currency. The flagship survey’s fieldwork closed **29 July 2025** — thirteen months before you read this. In a field where a model generation is six months, treat its *agent* numbers as a floor, not a snapshot. The March 2026 trust survey is the fresher instrument but a narrower one (~500 organisations, all of them already governing AI, so its respondents are more mature than the market). Neither is the truth; together they bracket it.

KEY TAKEAWAY

88% of organisations now use AI in at least one business function — up from 78% the year before — and **67% use it in more than one**. That is the adoption story, and it is essentially over: adoption is no longer a variable that distinguishes firms. Then the second set of numbers. **Only 39% report any EBIT impact at all from AI, and most of those attribute less than 5% of total EBIT to it**. Only about 6% qualify as high performers with 5% or more of EBIT attributable. Only 33% have begun scaling AI enterprise-wide; 62% are experimenting with agents but just 23% are scaling them anywhere, and **no single business function has more than 10% of organisations scaling agents in it**. The distance between 88% and 39% is the entire subject of Week 5. It is not a technology gap — everyone has the same models. It is a gap in **use-case selection, operating-model change and measurement**, and the survey names the discriminator: **high performers are nearly three times as likely to have fundamentally redesigned workflows — 55% against 20% — rather than bolting AI onto the existing ones**.

WHY IT MATTERS FOR THIS COURSE

Reading 1 gave you a mechanism and a set of firms it worked for. This reading gives you the **frequency** — how often it works, across two thousand organisations that were not selected for having succeeded. Together they are the honest picture: the mechanism is real, and most firms do not get it. For **Capsule 1**, this is the empirical backing for why use-case selection and data readiness are the first questions, not the last ones. For **Capsule 2**, it supplies the measurement discipline: “what would count as this having worked, and when would we know?” is a question you should be able to answer *before* you build the cost-to-serve model, not after. And in **Capsule 1**, immediately after the numbers above, it supplies the professional habit this course keeps returning to — when someone quotes you a number about AI, ask **who was asked, what threshold counted as success, and what the denominator was**. You will use that habit again on the case in Capsule 3, and again in your **project**.

Feeds Class 5: C1 (use-case selection • data readiness • why workflow redesign is the discriminator • reading vendor and consultancy numbers critically) • C2 (what to measure, and when) • C3 (the same habit, applied to the case).

MAIN POINTS

The adoption numbers — the part that is finished

- **88% use AI in at least one business function**, up from 78%; **67% in more than one**; **50% across three or more**. Adoption is no longer a differentiator, and any strategy that treats “we are using AI” as an achievement is describing table stakes.
- **Generative AI specifically: 79% now use it regularly**, up from 65% in the previous wave (Exhibit 1). Generative AI is now the majority of the adoption story, not a subset of it.
- **Where it is used matters more than whether**. Function-level self-reported benefits split cleanly: **56% cite cost savings** in software engineering and manufacturing; **67% report revenue gains** in marketing and sales. Cost and revenue use cases are different projects with different sponsors, different metrics and different failure modes — Capsule 1’s framework starts here.

The value numbers — the part that is not

- **39% report any EBIT impact from AI**, and among those, most attribute **less than 5%** of total EBIT to it. Read that sentence twice: “any impact” and “less than 5% of EBIT” are both in it, and both are doing work.
- **About 6% are high performers** — 5% or more of EBIT attributable to AI. Six per cent, after three years of the most intense enterprise technology adoption since the web.
- **33% have begun scaling enterprise-wide**; roughly two-thirds are still experimenting or piloting. Company size matters: **48% of firms above \$5bn revenue** have reached scaling, against **29% of firms under \$100m** — the AI-factory fixed-cost argument from Reading 1, visible in survey data.
- **Broader, softer benefits are widely reported**: 64% say AI enables innovation, and nearly half report improved customer satisfaction (45%) and competitive differentiation (45%). These are real and they are also the kind of benefit that does not have to survive a finance review — notice how much larger they are than the EBIT number, and ask why.
- **The discriminator: workflow redesign**. High performers are **nearly three times as likely** to have fundamentally redesigned workflows rather than inserting AI into existing ones — **55% against 20%**. This is Reading 1’s operating-model argument arriving as an empirical finding rather than a theory.

What this does to Reading 1 — the mechanism is real, and it is rare

Reading 1 builds its case from **selected exemplars**: Ant, Amazon, Google, Netflix, Ping An — firms picked because they succeeded, then examined for what they had in common. That is a legitimate way to find a **mechanism** and a useless way to establish a **frequency**, which is why Reading 1’s own criticism section flags it and points here. This survey is the corrective, because its two thousand respondents were not selected for having succeeded at anything. So here is the frequency, as close as the available evidence gets: **33%** have begun scaling enterprise-wide, **39%** can find any EBIT impact, and about **6%** can attribute 5% or more of EBIT to AI. **Reading 1’s five companies are drawn from that last band**. Nothing in this survey contradicts the mechanism — it says the mechanism is **rare**, and that reading the exemplars without the base rate is exactly how a strategy deck gets written that quietly assumes your firm is Ant.

Be precise about what has and has not been answered, because this is the kind of join that goes wrong. **McKinsey never asked “have you built an AI factory?”** Its 6% is defined by an **outcome** — EBIT attribution — while Reading 1’s exemplars are defined by a **mechanism**. The two populations overlap; they are not the same set. So 6% is a *proxy* for Reading 1’s success rate, not a measurement of it: some firms in the 6% got there with nothing resembling an AI factory, and some firms that built a real one are sitting inside the 61% because nobody captured a baseline. What the survey does supply — and it is the strongest link available — is the **discriminator**. The **55% against 20%** workflow-redesign gap is Reading 1’s operating-model claim appearing as a measurable, population-level difference rather than as an inference from five famous companies. That is the honest join: a mechanism proposed from exemplars, and one of its predictions confirmed on a sample that was not selected on the outcome.

Take the two together: Reading 1 tells you the mechanism exists. Reading 2 tells you it is a tail, not a trend. So the question to bring to Capsule 1 is not “does AI create value?” but “what would have to be true for us to be in the six per cent?” — and the only population-scale answer this survey offers is: redesign the workflow, do not decorate it.

Agents — the 2026 shape of the same gap

- **62% experimenting with or scaling agents; 23% scaling anywhere; a maximum of 10% scaling in any individual function**. The pilot-to-production ratio for agents in 2025 looks like the pilot-to-production ratio for generative AI in 2023 — which is either reassuring or alarming depending on how quickly you think the second curve moves.
- **What stops them, from the March 2026 companion**. **64% cite security and risk concerns as the top barrier** to fully scaling agentic AI (Exhibit 6). Separately, asked which risks are *highly relevant*, **74% say inaccuracy** and **72% say cybersecurity** (Exhibit 7); **60% name knowledge gaps** as the primary responsible-AI implementation barrier. Map those against Week 4 and they are not new: they are the agent failure modes, restated by the people who have to sign off on deployment.

Two of those figures are not siblings, and it matters — 64% and 72% answer different questions. They come from different exhibits with different answer lists, so you cannot add them, rank one against the other, or treat “security” and “cybersecurity” as a fine distinction the survey drew. It did not. The distinction worth drawing is this. **Cybersecurity assumes an adversary**: somebody is deliberately trying to make your system do the wrong thing — prompt injection through content the agent reads, a poisoned tool, data exfiltrated through a tool the agent is allowed to call, the agent’s own credentials misused. That is Week 6’s material. **“Security and risk”, as a barrier to scaling, mostly assumes nobody is attacking you at all** — it is the compound of things a risk function says no to: data governance and residency, lawful basis and consent, least-privilege access, auditability, third-party risk, regulatory compliance, model risk management. **An agent with write access and no audit trail is a security problem on a completely quiet day**. Both numbers are proxies for one anxiety: this thing can take an action, and we cannot yet demonstrate that we control it. Q4 below is about exactly that.

- **Governance maturity is improving slowly**. Overall responsible-AI maturity **2.3 out of 5, up from 2.0** the previous year; only about **one-third** reach level 3 or above on strategy, governance and agentic controls; **8%**

experienced an AI-related incident, and almost 60% of those rated their own response as only satisfactory or worse. That last figure is Week 6's subject and worth noticing now.

- **Workforce expectations are less dramatic than the discourse.** 32% expect workforce reductions of 3% or more; 43% expect no change; 13% expect increases of 3% or more.

HOW TO READ A NUMBER LIKE THIS — three surveys, three answers, no liars

Three widely-quoted findings about the same question, published within months of each other. They look incompatible. They are not, and taking them apart is the most transferable skill in this reading. Read the table **down a column** for what one study actually did, and **across a row** for why the three answers differ.

	MIT NANDA	McKinsey / QuantumBlack	Wharton × GBK
The study	<i>The GenAI Divide: State of AI in Business 2025</i> · July 2025 · non-peer-reviewed working paper	<i>The state of AI in 2025: Agents, innovation, and transformation</i> · 5 Nov 2025 · fielded 25 Jun – 29 Jul 2025	<i>Accountable Acceleration: Gen AI Fast-Tracks Into the Enterprise</i> · Oct 2025 · fielded 26 Jun – 11 Jul 2025
The number you will hear	“95% of organisations are getting zero return” on \$30–40bn of enterprise spend	39% report any EBIT impact · about 6% at 5% of EBIT or more	74% report positive ROI to date
Who was asked	52 structured interviews + 153 survey responses, recruited at four industry conferences, plus 300+ publicly disclosed initiatives	1,993 respondents across 105 nations, weighted by each nation's share of global GDP	~800 senior decision-makers at US firms with 1,000+ staff and \$50m+ revenue — the people who approved the spend , reporting on their own decision
What counted as success	A marked and sustained P&L or productivity impact within about six months . The harshest bar of the three.	Any impact on EBIT — a defined line in a real income statement. A finance judgement.	“Positive ROI” , with the threshold left to the respondent. A manager's judgement.
What the denominator was	All firms in the sample — and roughly 80% of them never piloted a custom system at all. Worked out below.	All respondents.	All respondents. Inside the same data, 34% of \$2bn+ firms report a neutral or “too early to tell” outcome.
Interest to declare	The authors build and sell the agentic infrastructure the paper recommends.	A consultancy that sells AI transformation work.	Co-published with a commercial marketing-strategy and insights consultancy.
What it actually supports	Among firms that did pilot, roughly one in four succeeded — an ordinary innovation success rate. The headline is largely a statement about firms that never took the shot.	Finance can see AI in the P&L at about four firms in ten , and materially at about one in sixteen .	Managers believe it paid off. Belief is not attribution, and the distance to McKinsey's 6% is exactly that gap.

The denominator, worked — the same five successes, two headlines

MIT's proportions, scaled to a round 100 firms so the arithmetic is visible. Nothing here is invented: the two figures the paper reports — 95% with zero return, and about a quarter of those who piloted succeeding — are consistent with each other, and this is what they look like when you write them out.

Step	Firms	What it means
Start with the sample	100	MIT's reported proportions, scaled to 100 for clarity.
Never piloted a custom system	80	About 80% of the sample. These firms cannot succeed or fail at something they never attempted.
Did pilot	20	The only firms of whom the question “does this work?” can honestly be asked.
...and got a return	5	The numerator . It does not change anywhere below.
Denominator = all firms	$5 \div 100 = 5\%$	→ “95% of organisations are getting

Step	Firms	What it means
		zero return. True — and mostly a measure of non-participation.
Denominator = firms that piloted	$5 \div 20 = 25\%$	→ “ three in four pilots fail. ” Also true, from the same five successes, and an unremarkable innovation success rate.
The version you will actually hear	“95% of pilots fail ”	Implies one success in twenty. The data say about one in four — the misquotation understates pilot success roughly fivefold.

The numerator never moved. Five firms succeeded in both readings. Only the choice of what to divide by changed, and it changed the answer by a factor of five — and then a further factor when the sentence was repeated with the word “pilots” swapped in for “organisations”. That is one number, one dataset, three different claims, and no dishonesty required at any step. When you next see a striking statistic about AI, the fastest useful question is not “is that true?” but “**a percentage of what?**”

They do not contradict. They ask different populations about different thresholds using different denominators. Look at the fieldwork dates: McKinsey and Wharton were in the field in the same three weeks of mid-2025, so 39% against 74% is not the world changing between two studies — it is the question changing. The gap between Wharton’s 74% saying “positive ROI” and McKinsey’s 6% showing 5% or more of EBIT is the distance between a manager’s judgement and a finance function’s attribution, and closing that distance is what Week 5 is for.

And one source with nothing to sell. All three studies above have an interest to declare — one sells the remedy, one sells the transformation work, one is co-published with a consultancy. The **Stanford HAI AI Index** does not: it is compiled by a university institute, it publishes its methodology and its sources, and it reports adoption and investment rather than ROI. **You already have it — it was Pre-read 2 in Class 1.** Go back to it when you need a baseline figure you can quote without a caveat, and notice the trade: it will tell you accurately how much is being spent and how many firms are using AI, and it will not tell you whether any of it paid off. Nobody neutral is measuring that, which is itself worth knowing.

Three questions to ask of any figure in this field, and they are the same three every time. **Who was asked?** — decision-makers who chose the investment are not neutral observers of it. **What counted as success?** — “positive ROI”, “any EBIT impact” and “sustained P&L impact within six months” are three different bars, and moving between them changes the answer by an order of magnitude. **What was the denominator?** — all firms, firms that piloted, or firms that deployed? MIT’s 95% and a “95% of pilots fail” headline differ only in the denominator, and the difference is the whole claim.

READING FOCUS QUESTIONS — bring your answers to class

There are no wrong answers here — these surface how you think, not what you recall. A clue sits under each question to give you a direction to push on; the fuller reasoning is in the section after.

- 88% are using AI; 39% can find any EBIT impact; about 6% can find more than 5%. Where exactly does the value leak out between those three numbers?

Clue. Draw it as a funnel and name each stage-to-stage loss separately rather than blaming “execution”. Between using and scaling, between scaling and value, between value and attributable value — three different losses with three different causes, and at least one of them is a measurement problem rather than a value problem. Then ask which stage your own organisation would fall out at, and be specific about why.

- MIT says 95% get zero return, Wharton says 74% get positive ROI, McKinsey says 39% see any EBIT impact. Reconcile them without accusing anyone of dishonesty.

Clue. Do not argue about which is right. For each one, write down three things only: **who was asked, what threshold counted, what the denominator was.** Then check whether any two of the three claims are even about the same population — and notice that the MIT figure changes meaning entirely depending on whether the denominator is “all firms” or “firms that piloted”.

- High performers are nearly three times as likely to have fundamentally redesigned workflows — 55% against 20%. Is redesign the cause of their performance, or a symptom of something else — and what would it take to tell?

Clue. Ask what a cross-sectional survey can and cannot establish, then generate at least two rival explanations that fit the same correlation — for example, that firms already capable of redesigning workflows are unusual firms, or that success came first and paid for the redesign. Then say what evidence would separate them. Whether you can still act on the finding while it remains ambiguous is the interesting part.

- 62% are experimenting with agents; no function has more than 10% scaling them. Name the specific thing that stops an agent pilot from becoming an agent deployment.

Clue. You already have the taxonomy — this is Week 4, seen from the buyer's side. Line the reported concerns up against the six mechanisms from the Week-4 close: (1) thought → action → observation, and the observation is the only door to reality; (2) **errors compound across steps** — $0.95^{20} \approx 0.36$; (3) cost grows with the square of the turns; (4) cost per task, not price per token; (5) **the first write is the governance cliff**; (6) readiness is set by the weakest link. Add the brownfield finding from the Devin evaluation — 2 of 8 greenfield tasks completed, **0 of 8 brownfield**. Sort them first into **adversarial** (someone is trying to make it misbehave) and **non-adversarial** (nobody is attacking and you still cannot show you are in control) — the two are usually collapsed into the word “security” and they need completely different answers. Then ask which of them is genuinely technical, and which is really an accountability question about who is responsible when the agent is wrong.

- You have budget for three AI use cases next year. Write down what you will measure to know whether you were right, and when you will know it.

Clue. Design the measurement **before** the use case, not after — that ordering is the whole exercise. For each candidate, name the baseline you would have to capture this month, the single metric that moves if it works, the time horizon before that metric can move at all, and the confound most likely to fool you. If you cannot name the baseline, you have found the reason 39% is not higher.

HOW TO THINK ABOUT THESE — read after you have had a go

These are not model answers to be memorised — they are the line of reasoning we will pick up in class. Class time is short, so they are written to stand on their own.

Q1 — Three separate leaks, and one of them is not a leak at all

The funnel is roughly: 88% using → 33% scaling enterprise-wide → 39% any EBIT impact → 6% material EBIT impact. The **first loss, from using to scaling**, is an operating-model loss, and it is Reading 1's argument in survey form: a pilot runs on a hand-assembled data extract and one enthusiastic team, and scaling means building the pipeline, the permissions and the support model that the pilot deliberately skipped. The firm-size split confirms it — 48% of firms above \$5bn have scaled against 29% below \$100m, because the fixed cost of an AI factory is the same regardless of how much revenue you spread it over. The **second loss, from scaling to value**, is a use-case-selection loss: plenty of systems work perfectly and automate something that was never expensive. This is the loss Capsule 1 is designed to prevent, and the reason “can we do this?” is the wrong first question. The **third gap, from value to attributable value**, is mostly not a loss at all — it is measurement. If you did not capture a baseline before deployment, you cannot attribute the improvement afterwards, and a finance function will not book what cannot be attributed. Some real value is therefore sitting inside the 61% who report no EBIT impact, invisible because nobody instrumented it. That is a genuinely different problem from the first two, and it has a cheap fix that almost nobody applies: **measure the baseline in the month before you start**. The practical takeaway is that these three leaks want three different interventions — plumbing, selection, instrumentation — and that organisations reliably respond to all three by buying another tool.

Q2 — Different populations, different thresholds, different denominators

Take them in order. **Wharton** asked about 800 enterprise decision-makers whether they had seen positive ROI; 74% said yes and 72% said they formally measure it. Notice who is being asked: the people who approved the spend, reporting on their own decision, against a threshold they define. That is not fraud, it is how self-reported ROI works, and the quarter of the largest firms saying “too early to tell” is the honest signal inside the same instrument. **McKinsey** asked a much larger, GDP-weighted sample about EBIT — a defined line in a real income statement — and got 39% for any impact and ~6% for 5% or more. The threshold moved from “I believe this paid off” to “finance can see it,” and the number fell by more than an order of magnitude — 74% down to about 6%. **MIT** applied the harshest bar available — sustained, marked P&L or productivity impact within about six months — to a small conference-recruited sample of 52 interviews and 153 responses, and reported that 95% cleared no return. Then the denominator problem: roughly 80% of that sample never piloted a custom system at all, so the headline is largely counting firms that never took the shot. Restated on the population that did pilot, the figure is closer to one in four succeeding, which is an unremarkable innovation success rate and would have made no headlines. Two further disclosures belong in any honest citation of it: it was not peer-reviewed, and its authors commercialise the infrastructure it recommends. So the reconciliation is complete without anyone lying. **The three numbers measure belief, attribution and a six-month P&L test, on three different populations.** The lesson generalises far beyond this reading, and it is the one to carry into your project: whenever two AI statistics disagree by an order of magnitude, the disagreement is almost always in the definition, not in the world.

Q3 — Probably both, and the survey cannot tell you which

The finding is a cross-sectional correlation, so causation is not available from it, and at least three explanations fit. **Redesign causes performance** — the intended reading, and the mechanically plausible one, because inserting AI into an unchanged workflow means the human handoffs, approval steps and batch cadences that set the cycle time are all still there; you have made one step faster in a process whose speed was never determined by that step.

Capability causes both — firms that can fundamentally redesign a workflow are already unusual: they have executive authority reaching across functions, process documentation, and a tolerance for disruption. Those same properties would predict success at any large technology programme, AI or not. **Success causes redesign** — early wins create the political capital and the budget to attempt something structural, so the arrow may run backwards. Separating them needs something a survey cannot supply: longitudinal data on the same firms, or a natural experiment, ideally with a measure of prior organisational capability held constant. Neither exists here. But — and this is the part worth arguing in class — the finding is still **actionable while ambiguous**, because all three explanations point the same way for you as a manager. If redesign causes performance, do it. If capability causes both, then the capability is what you should be building, and redesigning a workflow is how you build it. If success causes redesign, then sequence deliberately: pick a first use case small enough to win, and spend the credibility on the structural change rather than on a second tool. The honest position is to hold the causal claim loosely and the practical implication firmly, and to be suspicious of any consultancy that presents the correlation as a proven lever.

Q4 — Nothing technical; it is who signs for the mistake

Line the barriers up against the mechanisms and most of them resolve into one thing. **Inaccuracy at 74%** is Week-4 mechanism 2, errors compounding across steps: a system that is 95% right per step is about 36% right over twenty steps, and a pilot with a human watching every step never experiences that because the human is silently the error-correction layer. Remove the human — which is what scaling means — and the failure rate the pilot never measured becomes the deployment's defining property. **Security concerns, however they are labelled**, are the consequence of giving a probabilistic system real write access: a pilot reads, a deployment writes, and the difference in blast radius is the difference between an embarrassing answer and a reversed transaction. Split them before you argue about them. The **adversarial** half is cybersecurity proper — prompt injection, poisoned tools, exfiltration through a tool the agent may call — and it has technical answers, which is Week 6. The **non-adversarial** half is larger and is what actually blocks scaling: lawful basis, least-privilege access, auditability, vendor risk, model risk management. Nobody is attacking you and you still cannot demonstrate control. **Knowledge gaps at 60%** is the brownfield problem from Week 4 — the constraints that govern real work are mostly unwritten, so the agent cannot read them. The Devin evaluation is the calibration: 2 of 8 greenfield tasks completed, 0 of 8 brownfield. But under all four sits something that is not technical at all. In a pilot, when the agent is wrong, the answer to “who is accountable?” is “nobody, it is a pilot.” In a deployment, someone signs. Every one of the mitigations that actually unblocks scaling — confirmation gates on irreversible actions, capped autonomy, single-threaded writes, audit trails, a defined reversal path — is an **accountability** mechanism wearing engineering clothes. That reframes the 10% ceiling usefully. Firms are not stuck because agents cannot do the work; they are stuck because nobody has designed the arrangement under which a human is willing to be responsible for work they did not personally check. Design that arrangement first and the technical questions become tractable; design the agent first and you will do the pilot three times. This is also the most useful thing you can take into A2 and into your project: **ship the gate, the log and the reversal path in version one**, because they are what make version two deployable.

Q5 — Measure the baseline before you build, or you will join the 61%

A workable structure for each of your three candidates, and it fits on one page. **Baseline**: the number as it stands today, captured *this month*, before anything changes — average handling time, cost per case, conversion rate, error rate, days to close. If it does not exist, capturing it is the first sprint, and that is not a delay, it is the project. **Primary metric**: one number that moves if this works, named in advance, owned by someone outside the AI team. One, not five; a dashboard of twelve metrics is a way of guaranteeing that something will have gone up. **Horizon**: the earliest date at which that metric *could* move, which is usually later than the sponsor expects — a cost-to-serve metric can move in a quarter, a retention metric cannot, and holding a retention use case to a quarterly review will kill it before it can succeed. MIT's six-month bar is exactly this error institutionalised. **Confound**: the most likely alternative explanation, written down in advance so you cannot rationalise it later — seasonality, a pricing change, the sales team's new incentive, the fact that the pilot team is your best team. **Counterfactual**: how you will get one, ideally a holdout group or a staged rollout, which is Reading 1's experimentation platform doing its job. And one addition that costs nothing and changes behaviour: **write down in advance what result would make you stop**. Teams that define a kill condition before they start get to reallocate budget from the failures into the winners, and teams that do not, run pilots forever. Notice that none of this is about AI. It is ordinary programme discipline — which is precisely why the gap between 88% and 39% is not a technology gap, and why the firms closing it are not the ones with the best models.

APPENDIX — SURVEY AND MEASUREMENT VOCABULARY, EXPLAINED

The distinction that carries this reading: using AI is an activity; earning from AI is a result measured at the P&L. Almost every disagreement about “does enterprise AI work?” is two people quoting numbers from opposite sides of that line.

EBIT

- **Plain.** Earnings before interest and taxes — the profit the business makes from operating, before financing and tax effects.

Technically. A defined line in the income statement, which is why “% of EBIT attributable to AI” is a much harder claim than “positive ROI”. McKinsey: **39% report any EBIT impact**; most of those below 5%; **~6% at 5% or more**.

ROI (return on investment)

- **Plain.** Did we get more out than we put in — usually as judged by the person who put it in.
Technically. Undefined unless the numerator, denominator and period are stated. Self-reported ROI (Wharton: **74% positive**) and attributed EBIT impact (McKinsey: **39% any**) are not comparable measures, and the gap between them is mostly attribution rigour.

Adoption vs scaling vs transformation

- **Plain.** Someone is using it • it runs across the enterprise • the way work is done has changed.
Technically. Three distinct survey states with very different populations: **88%** adopted, **33%** begun enterprise-wide scaling, and a much smaller group having fundamentally redesigned workflows — the last being the group with the EBIT impact.

Denominator

- **Plain.** What the percentage is a percentage of.
Technically. The single most common source of misleading AI statistics. MIT’s “95% get zero return” is computed over **all firms**, ~80% of which never piloted; reported as “95% of pilots fail” it becomes a different and false claim. Always ask for the base.

Self-reported data and response bias

- **Plain.** People answering about their own decisions tend to answer kindly.
Technically. All three surveys here are self-reported. Direction of bias is predictable: decision-makers over-report success, and non-respondents skew toward firms with nothing to report. Sample weighting (McKinsey weights by national GDP contribution) corrects representativeness, not honesty.

Baseline and counterfactual

- **Plain.** What the number was before you started, and what it would have been if you had done nothing.
Technically. Without a pre-deployment baseline, no attribution is possible; without a holdout or staged rollout, no causal claim is. The absence of both is why real value can sit invisibly inside the 61% reporting no EBIT impact.

Workflow redesign

- **Plain.** Changing how the work is done, not just making one step of it faster.
Technically. The survey’s strongest discriminator — high performers **nearly three times as likely (55% vs 20%)**. Mechanism: inserting AI into an unchanged process leaves the handoffs, approvals and batch cadences that actually set cycle time untouched, so the local speed-up does not reach the outcome.

Shadow AI

- **Plain.** Staff using consumer AI tools for work because the company has not given them one.
Technically. MIT reports about **90% of employees using personal LLM tools for work** against roughly **40% of companies** buying official subscriptions. It is simultaneously evidence of real usefulness and an unmanaged data-governance exposure — a Week-6 problem that starts as a Week-5 procurement decision.

Agentic scaling

- **Plain.** Agents running real work without someone watching each step.
Technically. **62%** experimenting, **23%** scaling somewhere, **≤10%** in any single function. Reported concerns: security and risk as the top barrier (64%, Exhibit 6); inaccuracy (74%) and cybersecurity (72%) as highly relevant risks (Exhibit 7, a different question). In practice they reduce to one question — who is accountable when an autonomous system is wrong.

Cybersecurity vs “security and risk”

- **Plain.** One assumes somebody is attacking you. The other assumes nobody is, and you still cannot prove you are in control.
Technically. **Cybersecurity** is the adversarial subset: prompt injection through content the agent reads, poisoned tools, exfiltration via a tool the agent may call, misuse of the agent’s credentials (Week 6, OWASP). **“Security and risk”** as a scaling barrier is the compound a risk function actually blocks on — data governance and residency, lawful basis and consent, least-privilege access, auditability, third-party risk, regulatory compliance, model risk management. Most of it has no attacker in it. The survey reports them from **two different questions**, so the 64% and the 72% cannot be compared with each other.

SOURCE LINKS

Companion readings in this pack — these are the “Reading N” references used throughout this brief. **Reading 1** — Iansiti & Lakhani, *Competing in the Age of AI* (the operating-model framework this survey is testing in the field). **Reading 3** — *AI unit economics and cost to serve* (why so much of the value that does appear gets eaten before it

reaches EBIT). **Reading 4** — *GraphCast, Climate TRACE and AI in verticals* (the two archetypes of a use case). **Case pre-read 1** — *Ant Group (A), the Capsule-3 case*.

McKinsey / QuantumBlack — “The state of AI in 2025: Agents, innovation, and transformation” (5 Nov 2025) • www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai (the assigned reading)

McKinsey — “State of AI trust in 2026: Shifting to the agentic era” (25 Mar 2026) • www.mckinsey.com/capabilities/tech-and-ai/our-insights/tech-forward/state-of-ai-trust-in-2026-shifting-to-the-agentic-era (the 2026 companion)

Wharton Human-AI Research × GBK — “Accountable Acceleration” (Oct 2025) • ai.wharton.upenn.edu/wp-content/uploads/2025/10/2025-Wharton-GBK-AI-Adoption-Report_Full-Report.pdf (74% self-reported positive ROI)

MIT NANDA — “The GenAI Divide: State of AI in Business 2025” (Jul 2025, working paper) • mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf (the “95%” claim — read the method section)

Stanford HAI — 2026 AI Index Report (13 Apr 2026) • hai.stanford.edu/ai-index/2026-ai-index-report (your Class-1 Pre-read 2 — the non-commercial baseline for adoption and investment)

CAVEATS

1. This summary is not a substitute for the survey. The charts carry information the headline numbers do not — particularly the function-level and firm-size breakdowns, which is where most of the usable guidance sits.
2. All figures here are **self-reported**, including the EBIT ones. Nobody audited a general ledger. Treat every percentage as a report about what executives believe and are willing to state.
3. The flagship survey's fieldwork closed 29 July 2025. Its agent numbers in particular are likely to be a floor by August 2026.
4. The MIT NANDA paper is included because it is the most-quoted statistic in the field and because taking it apart is a teaching exercise, not because it is reliable. It is a non-peer-reviewed working draft, its sample is 52 interviews and 153 responses, and its authors commercialise the remedy it proposes.
5. This summary was generated and fact-checked using Anthropic models — a vendor with a direct commercial interest in enterprise AI adoption being higher than these numbers suggest. The three-question test, the funnel decomposition and the accountability reading of the agent barriers are the instructor's. Flag anything that looks wrong.