

CASE PRE-READ 1 OF 1 • CAPSULE 3

It scaled exactly as the theory said it would — and then it was stopped: Ant Group (A)

Palepu, K. G., Zhu, F., Ma, S. L., Herman, K. "Ant Group (A)." Harvard Business School case 9-122-003, rev. 2 February 2023 (22 pp. including exhibits; developed from published sources). Licensed to PE6201 for Jul 2026 – Jan 2027. We brief and discuss this **live in Capsule 3** — this note prepares you for that discussion; it does not replace the case. Note the coincidence that makes this the right case for this week: *the case's own endnotes cite Iansiti and Lakhani twice*, and Ant is the flagship example of the argument *Reading 1* sets out.

ASSIGNED READING

Read the **case body** (pages 1–8) carefully and work through **Exhibits 8a, 8b, 10 and 11** — the revenue mix, the user counts and the payment volumes are where the argument actually lives, and we will use them live. Exhibits 3, 6 and 7 (competitors, overseas investments, ownership) reward a skim. Read it as a decision case, not as history: the last line asks what Ant faces next, and you are being asked to answer that from **3 November 2020**, with only what was known then.

» case ~45 min • this brief ~23 min

Note on currency. The case ends on **3 November 2020**, two days before a cancelled listing, and was last revised in **February 2023**. Nearly six years have passed. A short, dated factual update is provided at the **end of this brief, deliberately placed after the focus questions** — read it only once you have answered them. Knowing how something turned out makes it very hard to reconstruct how the decision looked at the time, and reconstructing that is the skill the case is for.

KEY TAKEAWAY

Ant is what Reading 1 describes, built and running at national scale. A payments app launched in 2004 to solve an escrow problem became a **cloud-native operating model** that processed transactions at **pennies apiece**, assessed creditworthiness **in seconds** from behaviour it already observed, and redeployed the same data and distribution into money-market funds, lending, wealth management, insurance and credit scoring — each new business cheaper to launch than the last. The numbers the case prints are the AI-factory thesis in its strongest form: **one billion annual Alipay users** by 2020, **729 million** of them using at least one financial product, **¥118 trillion (US\$17 trillion)** of payment volume in the year to June 2020, **¥1.73 trillion (US\$263 billion)** of consumer credit outstanding — of which **Ant funded 2%** — and a delinquency rate of about **2% against roughly 6% for banks**. Revenue mix moved from **54.9% payments in 2017 to 56.2% digital-finance technology platform by 2019**, and by the first half of 2020 the technology platform was **63% of revenue**. The IPO was expected to raise **US\$34 billion at a US\$310 billion valuation**. Then, two days before listing, it did not happen. **The case is not about whether the AI factory works. It plainly worked. The case is about what happens when it works so well that it becomes something a state has to have an opinion about.**

WHY IT MATTERS FOR THIS COURSE

Every Week-5 reading meets here. **Reading 1** gives you the framework, and **Iansiti and Lakhani built that framework on this company**: Ant is *their* exhibit A. What Reading 1 can relay of their evidence for the headline claim is one flattering paragraph — a billion consumers within five years of launch, on **fewer than one-tenth the employees of the largest US banks**. You need neither the article nor the book to sit this class, and the licence does not let us circulate either; Reading 1 is self-contained, and that paragraph is in its **Key Takeaway**. What the framework never showed you is the **full record of the same company** — 22 pages that do not flatter it, and which you are about to hold. **Reading 2** tells you how rare this outcome is, and the point applies to this case as much as to anyone else: 88% of firms use AI, about 6% show material EBIT impact, and **Reading 1's five exemplars are drawn from that last band — Ant is one of the five**. The Capsule-3 case block spends its half hour on a company selected because it succeeded — which is the exact move Reading 2 warns about, and its warning is worth carrying as one line: **studying only the winners tells you the mechanism and nothing about the frequency**. That is how a strategy deck gets written which quietly assumes your firm is Ant. So do it knowingly: any explanation you offer for Ant has to explain **the base rate** too — the roughly **94%** who also use AI and show no material EBIT impact. Whatever you credit Ant with, ask why it did not work for them. Reading 2 also supplies the discriminator — high performers are **nearly three times as likely to have fundamentally redesigned workflows, 55% against 20%** — and Ant is that finding in its limit case, because it did not add a scoring model to a loan-application process; for three-quarters of the population there was no loan-application process. **Reading 3** is the arithmetic under "pennies per transaction" — cost to serve is not a detail of this business, it is the strategy. **Reading 4** hands you the archetypes, and note that this company ran **both**: processing payments at pennies apiece is **archetype A**, a settlement and clearing computation that banks already ran and already trusted, made radically cheaper; scoring creditworthiness for people with no credit file is **archetype B**, a measurement manufactured where none existed. And the discussion runs directly into your **A2** and your project: this is what it looks like when the use case is selected correctly, the data is

ready, the unit economics work, the network effects compound — and the binding constraint turns out to be the one Reading 1's framework does not contain. **Not an internal bottleneck at all, but an externally imposed constraint that arrives precisely because you succeeded.**

Feeds Class 5: C1 (use-case selection • the AI factory • network effects and data flywheels) • C2 (cost to serve at national scale) • C3 (the live case block) • and forward to Week 6 (regulation as a design constraint, not an afterthought).

MAIN POINTS

The company, in the numbers the case prints

- **Origins.** Alipay launched **2004** to hold buyer payments in escrow on Taobao, at a time when only **1% of Chinese citizens had credit cards**. Merchant fee **0.55% per transaction** against a US small-business average of about **2%**. Over **50 million users by 2007**; **650 million** by the 2011 spin-off, prompted by the PBOC's 2010 requirement that non-bank payment providers be licensed.
- **The 2013 inflection.** Yu'e Bao, a money-market fund reachable from a spare balance with a minimum of **¥1 (about 16 cents)**, took **30 million users and US\$16.3 billion in six months**. By 2017 it was the **world's largest money-market fund** with **370 million account holders and US\$211 billion** in assets, yielding over **4%** against **1.5%** on one-year bank deposits.
- **2014.** Alipay becomes **Ant Financial**; **190 million active users**; **82% of the Chinese mobile-payments market**. Ma insists on **"techfin"**, not fintech — a self-description worth interrogating rather than accepting.
- **Valuation ladder.** **US\$60 billion** (April 2016) → **US\$150 billion** after a **US\$14 billion** round (June 2018) → an IPO expected to raise **US\$34 billion at US\$310 billion** (October 2020).
- **Scale at the end of the case.** First half of 2020: revenue **¥72.5 billion (US\$10.5 billion)**, up **38%**; profit **¥21.2 billion (US\$3 billion)** in six months, more than all of 2019. **One billion** annual Alipay users, **729 million** using a financial product, **¥118 trillion (US\$17 trillion)** of payment volume in the year to June 2020.

The operating model — Reading 1's four components, visible in one company

- **Infrastructure built for cost, not prestige.** A cloud-based platform processing transactions at **pennies apiece** on **inexpensive PC server clusters**, holding **99.9% uptime** through synchronised replicas across multiple locations. Elasticity to absorb demand spikes. This is a deliberate cost-to-serve decision at the architecture layer, and it is what made **¥1** products viable.
- **An open architecture with a data layer at the bottom.** The case describes three layers: a **centralised data repository**, an intermediate layer of **pluggable computing engines** with per-engine security policies, and an **API layer** on top serving many business applications. That is Reading 1's software-infrastructure component described almost word for word — and it is what made **scope** cheap.
- **And notice how that data layer was acquired.** Ant's arrived as a **by-product**: it was running a payments network, and creditworthiness fell out of the exhaust. Our **Class-2 DBS case** is the other route — an IBM Watson experiment killed in 2014 at about 80% accuracy, then years of unglamorous consolidation of **twelve separate data warehouses into one platform, ADA, by 2018**, with the governance framework built alongside it. Different countries, different starting points, the same decade. **Neither firm rented it, and no API price list shows it** — which is exactly where Reading 1 says the moat now sits.
- **Decisions in seconds, from data already held.** Huabei (2014, interest-free for 40 days) and Jiebei (2015, larger loans) assessed creditworthiness **in seconds** from prior Alipay behaviour. **MYBank**, in which Ant took a **30% stake** after receiving one of a handful of private banking licences in 2015, approved SME loans in **under one minute, 24/7**, running each application through **3,000 risk-control strategies** with **zero human intervention**, up to **¥5 million (about US\$800,000)**.
- **The learning loop closes by itself.** Risk categories update continuously: *"as consumers with a small initial Huabei credit line transact, repay and build credit histories, we assess their creditworthiness on an ongoing basis and dynamically adjust our credit approval process and limits."* Default is a label the world supplies for free — which is why the delinquency figure (**~2% vs ~6% for banks**) is the number to notice, not the speed.
- **An executive states the whole thesis plainly.** *"Lending institutions need to answer only three basic questions: Should we lend to them, how much should we lend, and at what interest rate? Once sellers on our platforms gave us authorization to analyze their data, we were well positioned to answer those questions."*

Scope — the flywheel, business by business

- **Payments → funds.** Alipay's float and its users made Yu'e Bao possible; Ant earned a fee from the fund manager (Tianhong) as a percentage of assets under management, rather than managing money itself.
- **Payments → credit.** Transaction history made Huabei and Jiebei possible for people with no credit file — the case reports about **72% of Chinese adults financially illiterate** in the mid-2010s and traditional credit reporting covering only about **25%** of the population.

- **Credit → wealth and insurance.** Ant Fortune (2015) offered thousands of products from **100+ fund institutions**, recommended by data analytics. Insurance arrived in 2018 on an explicit principle: *“any activity with risk can be insured if sufficient data is available.”* **Quanminbao** priced a pension annuity at **¥1 (15 cents)**; **Xianghubao** reached **50 million participants** at about **US\$2 a month** with payouts up to **US\$45,000** — **47% migrant workers, 31% rural residents**. In the year to June 2020, over **570 million** Alipay users bought an insurance product.
- **And a scoring product on the side.** **Zhima Credit** (2015, one of eight permitted) scored users on transactions, shopping behaviour, education and employment, public records, partner data — **and the credit scores of their social connections**. High scores unlocked deposit-free hotels and rentals. The case’s footnote is the interesting part: **Ant reportedly never used Zhima to inform its own lending decisions.**

The business-model shift — from balance sheet to technology service fee

- **What changed.** Ant moved from underwriting Huabei and Jiebei itself to **partnering with banks**, which supplied the capital while Ant supplied the credit assessment. Users largely did not know a partner bank was behind the loan.
- **How Ant got paid.** A **“technology service fee”** the case reports as **30% to 40% of the interest income** the partner bank generated. Annualised rates on most loans **14.6%**, reaching **18.25%**. Notice the *shape* of that fee and not only its size: Ant is paid a share of the value the loan actually produces, so its revenue moves with the same thing its costs move with. **That is outcome pricing, a decade before anyone called it that** — Reading 3 derives the same design from first principles for AI products, where cost of goods sold scales with usage and a per-seat fee therefore breaks. A fixed fee per query would have decoupled the two.
- **How big it got.** **44% of 2017 revenues** came from technology service fees; the technology platform was **63% of total revenue** in the first half of 2020. By June 2020, more than **US\$240 billion** of consumer credit could be attributed to Ant, of which **Ant funded 2%**. The prospectus is explicit: *“Our approach is not to take credit risk ourselves.”*
- **What observers said about it.** *“It’s the western bankers’ worst nightmare of what would happen under an open banking system. Essentially, the banks lose their direct relationship with the customer, and all of it is mediated by the platform.”* And, from a bank official, the admission that neither PBOC credit data nor the bank’s own customer data was as strong as Ant’s.

The counter-pressure — the constraint that arrives because you succeeded

- **2010–2011.** PBOC requires non-bank payment providers to be licensed; Alibaba spins off Alipay; Alipay becomes the first to receive the Payment Business Licence. Regulation arrives early and Ant adapts.
- **2017.** Limits on how much of users’ idle balances payment firms may invest for their own benefit — reducing float income and making standalone payments less profitable, though Alipay remained valuable as the entry point. Yu’e Bao and its imitators are required to diversify and prioritise liquidity over yield.
- **Early 2018.** Rules prohibiting independent parties from issuing credit ratings **effectively end Zhima Credit’s business.**
- **Autumn 2020.** New PBOC rules require companies with at least two financial businesses — or whose financial assets are 85%+ of total assets — to hold paid-up capital of **¥5 billion (US\$731 million)**, effective November 2021, and described as aimed at conglomerates and internet companies controlling financial businesses that could pose systemic risk.
- **The last two weeks.** **21 October 2020:** dual listing approved. **24 October:** Ma tells a room of regulators *“we cannot regulate the future with yesterday’s means”* and *“there’s no systemic financial risks in China because there’s no financial system in China.”* Days later he is summoned by four agencies. **3 November:** the IPO is shut down.
- **The tension the case sets up without resolving it.** Ma’s own words, earlier: *“We always step ahead of the regulators — we have to. Otherwise, we go nowhere,”* alongside *“Now we’re big. We probably regulate ourselves much more strictly — much more — than the regulators.”* Both cannot be the operative principle at the same time.
- **And do not file this under “China”.** Reading 1 now makes the general claim and it is the one to carry into class: the term missing from the operating-model framework is **any externally imposed constraint that arrives because you succeeded — regulator, platform, or public**. Social-media recommendation built the purest version of the flywheel Reading 1 describes and no capital rule ever touched it, yet content and platform-transparency duties arrived anyway, **because** ranking at that scale worked as well as it did. Digital advertising looked un-constrainable for two decades until **Apple** changed the tracking rules on its own operating system in 2021 — no regulator involved at all, and Meta’s own estimate of the cost was about **US\$10 billion** of 2022 revenue. Ant is the *state* variant of a three-way category. Read it as a story about one country and you will take the wrong lesson home.

What the case leaves open — worth having a position on before class

- **Was Ant a technology company or a bank?** Ant said technology: it takes fees, not credit risk. Regulators looked at ¥1.73 trillion of credit intermediated by one firm and drew a different conclusion. Almost everything in the case turns on which description you accept, and the answer decides both the capital requirement and the valuation multiple.
- **Who bears the risk when the platform bears 2%?** Partner banks funded the loans; the case notes some were *weaker financial institutions at higher risk if borrowers defaulted*. Ant's ~2% delinquency was achieved on loans mostly sitting on someone else's balance sheet.
- **Whose data was it?** Zhima scored users partly on **their social connections' scores** and on government records. The case reports it without comment; you should not.
- **The ownership picture (Exhibit 7).** Alibaba **32.65%**; Hangzhou Junhan **29.86%** and Hangzhou Junao **20.66%**, both controlled through Hangzhou Yunbo in which Ma held 34%; state-owned or state-affiliated investors together about **5.3%**. Read the exhibit before deciding who could have prevented what.
- **Competition, which the case keeps in the background.** Tencent's Tenpay held about **20%** of mobile payments by 2015 and grew through WeChat; JD.com and Baidu were also present (Exhibit 3). Ant's 82% share in 2014 was not a permanent condition.

READING FOCUS QUESTIONS — bring your answers to class

There are no wrong answers here — these surface how you think, not what you recall. A clue sits under each to give you a direction to push on. **This brief carries no answers section and no verdict:** the strategic judgement is the live discussion in Capsule 3, and arriving with a position you can defend and then attack is the point.

- Using Reading 1's framework, which of **scale**, **scope** and **learning** did Ant actually monetise — and which one did the 2018 and 2020 interventions take away?

Clue. Take the three separately and attach each to a specific business and a specific number in the case rather than arguing in general. For scale, use Reading 1's three shapes rather than a yes or no: the branch-based lender Ant displaced is the **superlinear** case — you cannot hire ten times the underwriters at the same quality of judgement, so supervision, error rates, queueing and training lead time all pile on top of headcount — and Reading 1 flags superlinear as the shape rooms most often **under-claim**. Then run the same test over the regulatory actions: the Zhima ban, the float restrictions and the capital requirement each removed a different one, and only one of them touched the thing that was actually compounding. Notice which of the three is hardest for a regulator to reach, and why that might be the uncomfortable finding.

- Ant intermediated **¥1.73 trillion** of consumer credit while funding **2%** of it, and took **30–40% of the partner bank's interest income** as a technology service fee. Whose risk is that — and what should it do to the multiple you would pay for the business?

Clue. Work out who is holding what, from the case's own numbers, before you reach for a label. Then ask the valuation question in two halves: a technology company earning fees and a lender earning spread are valued on completely different multiples, and the difference between them at Ant's scale is a very large number — so calculate roughly what the label itself was worth. Finally ask whether the arrangement is stable when defaults rise, and who finds out first.

- Use Reading 3's method on the case's own figures: what does one credit decision cost Ant, and what does it cost a branch-based lender? Then say what that ratio bought — and what it did **not** buy.

Clue. Use Reading 3's three layers and keep them apart, because the answer is in which one dominates. **Layer 1, per-decision variable:** "pennies per transaction", 3,000 risk-control strategies, under a minute, zero human intervention. **Layer 2, expected failure cost:** $(1 - \text{repayment rate}) \times \text{loss given default}$ — and on ¥1.73 trillion at roughly 2% delinquency this term is orders of magnitude larger than layer 1. **Layer 3, fixed:** the platform itself, amortised across a billion users. Now put a branch lender's loaded cost per decision beside layer 1: the ratio is spectacular, and you have just priced **the smallest term in the equation** — which is the same lesson Reading 3 teaches by setting the token price to zero and watching almost nothing move. So the second half is the real question. The layer-1 ratio explains which customers became **reachable**, and the case gives you the financial-inclusion numbers to size that. Ask whether it explains the **2% versus 6% delinquency**, and if not, what does. Then ask what Ant could have priced any of this against, given that for three-quarters of the population there was no incumbent process to save money on at all.

- Ant built Zhima Credit and, the case reports, **never used it for its own lending decisions**. Why build a credit-scoring product you do not use — and what exactly did the 2018 ban target?

Clue. Generate several rival explanations before choosing: a product for others, a data-acquisition instrument, a behavioural device that makes users act more creditworthy, a regulatory hedge, a licence worth holding. Check each against what Zhima actually did for users. Then look closely at what the 2018 rule prohibited — **who may issue a rating**, not what may be measured — and ask what that tells you about which part of the arrangement the state considered a public function.

- Would the IPO have been stopped without Ma's 24 October speech? Marshal the evidence in the case on both sides before you answer.

Clue. Build two columns from the case text alone. On one side, the sequence and the timing — the speech, the summons, forty-eight hours from listing. On the other, the trail of measures stretching back to 2010, the capital rules already issued that autumn, and the academic quoted saying “the consensus now is that something must be done.” Then ask the harder question underneath: if the answer is “both”, what did the speech actually change — the decision, or its timing and its visibility? Those are very different lessons for anyone running a regulated business.

- You are the regulator. One platform reaches a billion users and intermediates **US\$263 billion** of consumer credit while carrying **2%** of the risk. Design your intervention — and price what it costs the people the platform was serving.

Clue. Diagnose before you prescribe: write down the specific harm you are preventing, and notice that “systemic risk”, “consumer over-indebtedness”, “data concentration” and “an unlicensed firm performing a licensed function” are four different harms wanting four different instruments. Then do the part most people skip — the case tells you who these borrowers were and what they had access to before, so estimate what your intervention costs them, and be honest about whether you would still choose it.

WHAT HAPPENED NEXT — read only after you have answered the questions above

Factual, dated, and deliberately placed here. Knowing the ending makes it much harder to reconstruct the decision as it looked on 3 November 2020, which is what the case is for. No verdict is offered — these are events, and what they mean is Capsule 3.

- **April 2021.** The PBOC and three other regulators jointly order Ant to restructure as a **financial holding company** with every financial activity separately licensed — a five-point rectification plan. The application was accepted in June 2022; **no public confirmation of a licence being granted has followed as of August 2026.**
- **June 2021 onward.** Huabei and Jiebei are moved into **Chongqing Ant Consumer Finance**, a licensed consumer-finance company in which Ant holds **50%**. Registered capital rises from **¥8 billion** to **¥18.5 billion** (January 2023) and then **¥23 billion** (October 2023), with state-linked shareholders taking roughly **22%**.
- **December 2021 – January 2022. Xianghubao closes.** Ant announced on **28 December 2021** that the mutual-aid health scheme would shut the following month; press reports at the time put it above **100 million participants**. Set that against the principle the case quotes for the whole insurance business — “any activity with risk can be insured if sufficient data is available.” Xianghubao did not close for want of data. Two questions worth holding: what did it close for, and was a mutual-aid scheme with a hundred million members insurance in the first place? The case does not settle the second, and for three years neither did anyone else.
- **January 2023.** Ant announces that **Jack Ma is no longer its controlling person**: voting rights fall from **over 50% to 6.2%**, with the ten major shareholders agreeing to vote independently rather than in concert. Economic interests are unchanged.
- **July 2023.** A fine of **¥7.123 billion (US\$984 million)** covering corporate governance, consumer protection, payment and settlement, anti-money-laundering and fund sales. The PBOC states that supervision of platform financial businesses now moves **from centralised rectification to normal supervision**; Ant says rectification is complete. Days later Ant offers to buy back shares at an implied valuation of about **US\$78–79 billion** — roughly **70% below** the level implied by the cancelled IPO.
- **November 2024.** The PBOC grants a personal credit-reporting licence to **Qiantang Credit Reporting**, in which Ant holds **35%** alongside a state-owned enterprise at **35%**, valid to November 2027 — only the third such licence in mainland China, after a three-year wait. Read against the 2018 Zhima ban, this is the credit-data question being answered not by prohibition but by **licensing, co-ownership and a fixed term**.
- **2024–2026.** Ant separates **Ant International**, **Ant Digital Technologies** and **OceanBase** as independent units. The group states in December 2024 that it has **no plans to go public**; by mid-2026 the live prospect is a possible Hong Kong listing of the **international unit**, not the parent.
- **2026.** New consumer-lending rules cap all-in annualised borrowing costs, with a stated path from **24% to 20% to 18%** and a long-run target near 12% by end-2027 — directly compressing the economics of a loan book Caixin reports at **over ¥2 trillion**. Meanwhile Ant's R&D spending reaches **¥35.03 billion (US\$5.17 billion)** in 2025 and it publishes open-weight **Ling** and **Ring** model families; Alibaba's reported share of Ant's profit fell sharply through 2025–26, attributed to AI investment and fair-value movements.

Two things to notice, and then stop. **The regulatory answer was not to break the company up — it was to license and capitalise each activity, and to take a state stake in the data.** And the tension the case describes did not end in 2020: the 2026 rate caps are the same argument about who bears the cost of consumer credit, arriving a second time.

APPENDIX — CASE VOCABULARY, EXPLAINED

The distinction that carries this case: **a technology company that charges fees for a service, versus a financial institution that intermediates credit and must hold capital against it.** Ant claimed the first; regulators concluded the second. Every term below helps you argue that question precisely.

Techfin vs fintech

- **Plain.** A technology company that does finance, or a finance company that uses technology.
Technically. Ma's preferred self-description was **techfin**. It is not marketing: it determines whether you are supervised as a licensed financial institution, what capital you must hold, and what multiple the market pays. The case's central unresolved question in two words.

Technology service fee

- **Plain.** The platform assesses the borrower, the bank lends the money, and the platform takes a cut of the interest.
Technically. Reported at **30–40% of the partner bank's interest income; 44% of 2017 revenues**. Economically it is a share of lending spread; accounted for and argued as a fee for a technology service. Which of those two descriptions is right is the whole capital-requirement question.

Credit intermediation and who holds the risk

- **Plain.** Someone decides who gets the loan; someone else's money is at stake if it is not repaid.
Technically. **¥1.73 trillion (US\$263 billion)** outstanding at June 2020, of which Ant funded **2%**. The prospectus states the position directly: *"Our approach is not to take credit risk ourselves."* Regulators' concern, per the case, was that partner banks — some of them weaker institutions — carried default risk on assessments they did not make.

Float and float income

- **Plain.** Money sitting in users' accounts between paying and being paid out, which the platform can invest.
Technically. A material revenue source for payment firms until the 2017 rules restricted how much of it could be invested for the platform's own benefit — which is what made standalone digital payments *"less profitable as a standalone business"* while leaving Alipay valuable as the entry point to everything else.

Money market fund (MMF)

- **Plain.** A very low-risk fund you can put spare change into and take out any time.
Technically. Yu'e Bao: **¥1 minimum**, managed by Tianhong with Ant taking a percentage of assets under management. The world's largest MMF by 2017 — **370 million account holders, US\$211 billion** — which is precisely why regulators required diversification and liquidity limits. Size itself became the regulated property.

Zhima Credit (Sesame Credit)

- **Plain.** A private credit score built from shopping, payments, public records — and the scores of your friends.
Technically. One of eight private scoring schemes permitted in 2015. Its business was effectively ended in early 2018 by rules prohibiting **independent parties from issuing credit ratings**. Note the case's footnote that Ant reportedly did not use it for its own lending decisions.

Financial inclusion

- **Plain.** Giving people who never had access to a bank a way to pay, save, borrow and insure.
Technically. The case's framing and the mission statement's claim. Its numbers: **~72%** of Chinese adults financially illiterate in the mid-2010s, traditional credit reporting covering only **~25%** of the population, and **9.6%** of adults accessing credit from a financial institution between 2011 and 2014. Xianghubao's 50 million participants were **47% migrant workers, 31% rural**. Reading 4 names this the commoner flavour of archetype B — not a number somebody shaded, but one that **never existed for most cases at all**, because producing it took a scarce person and was only ever done for a sample.

Systemic risk and paid-up capital

- **Plain.** A firm big enough that its failure damages everyone else — so it is made to hold real money against what it does.
Technically. The autumn-2020 rules required **¥5 billion (US\$731 million)** paid-up capital for firms with two or more financial businesses, or with financial assets above 85% of total assets. Capital requirements are the instrument that converts "too big to fail" from an observation into a constraint — and they price a business model directly.

Total payment volume (TPV)

- **Plain.** The total value of money that moved through the platform.
Technically. **¥118 trillion (US\$17 trillion)** in the year to June 2020 (Exhibit 10). A scale metric, not a revenue metric — Ant's take rate on payments was **0.55%** charged to merchants — and a good example of a number that is quoted to impress and needs a take rate attached before it means anything.

Annual active users (AAU)

- **Plain.** People who used the app at least once in the year.
Technically. 987 million Alipay AAUs in the twelve months to June 2020, of which 729 million used a digital-finance product (Exhibit 10). The gap between the two is the flywheel's conversion rate, and it is the single most informative pair of numbers in the exhibits.

SOURCES AND COMPANION READINGS

Companion readings in this pack — the four foundation briefs referred to by number in class. **Reading 1** — Iansiti & Lakhani, *Competing in the Age of AI* (the operating-model framework; this company is its flagship example, and Reading 1 is the assigned text — the article behind it is optional and not needed here). **Reading 2** — McKinsey, *The State of AI*. **Reading 3** — *AI unit economics and cost to serve*. **Reading 4** — GraphCast, *Climate TRACE and AI in verticals* (Ant ran both: **archetype B** on credit, **archetype A** on payments).

Palepu, K. G., Zhu, F., Ma, S. L., and Herman, K. “**Ant Group (A)**,” Harvard Business School case **9-122-003**, rev. 2 February 2023. Copyright © 2021, 2022, 2023 President and Fellows of Harvard College. Developed from published sources; funded by Harvard Business School and not by the company. Obtain your own copy through the course's licensed link.

CAVEATS

1. This brief is a study aid, not a substitute for the case. It carries material printed in the case body and points you at the exhibits; it deliberately does not reproduce their detail, because working through Exhibits 8a, 8b, 10 and 11 yourself is part of the preparation.
2. It gives you **no verdict**. The numbers, the tensions and the open questions are here; the strategic judgement is the live discussion in Capsule 3. Arrive with a position you can defend, and then be willing to attack it.
3. The case ends on 3 November 2020 and was revised in February 2023. Judge the decisions in it on the information available at the time — the only fair way to assess any strategic choice — and treat the “What happened next” section as context for the discussion, not as a scorecard.
4. The post-2020 update is compiled from press and regulatory reporting, not from the case. Where sources were unclear it says so — notably, there is no public confirmation that the financial-holding-company licence was ever granted. Figures in that section are as reported and may have moved again.
5. This brief was generated and fact-checked using Anthropic models. The framing — that Ant is Reading 1's thesis meeting an externally imposed constraint that arrived **because** the thesis worked, and that such a constraint need not come from a regulator — is the instructor's reading, not the case authors'. Flag anything that looks wrong.
6. Licence: copyright President and Fellows of Harvard College, authorised for use in PE6201 at Nanyang Technological University from **July 2026 to January 2027**. Not to be digitised, reproduced, posted, transmitted or forwarded outside this course.