

PRE-CLASS READING 4 OF 5

Your home markets — Singapore, India and China, and three completely different answers to the same question

Singapore — IMDA / AI Verify Foundation, Model AI Governance Framework for Generative AI (30 May 2024), and the Model AI Governance Framework for Agentic AI (January 2026, updated May 2026). **India** — MeitY's Digital Personal Data Protection Act 2023 with the DPDP Rules 2025, NITI Aayog's Responsible AI for All, and the India AI Governance Guidelines (November 2025). **China** — CAC's Interim Measures for the Management of Generative AI Services (effective 15 August 2023) and the instruments that followed it. All free, all public, none paywalled. **This is the reading that matters most for this cohort**, because these are the markets most of you will actually deploy into, and Capsule 2 gives them the class time that the EU and US frameworks in Reading 3 do not get.

ASSIGNED READING

Read the **nine dimensions** of Singapore's Generative AI framework — the summary table, not the full 20 pages (≈ 10 min). Read the **seven Sutras and six Pillars** of India's AI Governance Guidelines, and note the DPDP Act's commencement dates (≈ 10 min). Read **China's Interim Measures in full** — it is genuinely short, twenty-four articles, and the China Law Translate English version is excellent (≈ 15 min). Do not read the DPDP Act itself; read a summary. **Read them looking for what each is trying to control**, because it is a different thing in each case and that is the whole lesson.

» **sources ~35 min • this summary + reasoning notes ~24 min**

△ **Note on currency — two things that most commentary gets wrong. India: the DPDP Act's substantive provisions are not yet in force.** The Act received assent on 11 August 2023, but nothing operative commenced until the DPDP Rules were notified in **November 2025**, on a three-stage timetable. Stage 1 (13 November 2025) brought only definitions and the Data Protection Board. Stage 2 (13 November 2026) brings Consent Manager registration. **Stage 3 — notice, consent, Data Principal rights, children's data, Significant Data Fiduciary duties and cross-border transfer — does not commence until 13 May 2027.** So as of today India's headline data-protection obligations bind nobody. (*Sources differ on whether the Gazette date is 13 or 14 November 2025; the commencement clock runs from 13 November.*) **Singapore: the agentic framework is the current front edge**, launched at Davos in January 2026 and substantially updated in May 2026 — the Generative AI framework of 2024 is still live and is no longer the newest thing. (*We could not confirm the version string on the agentic PDF; treat "January 2026, updated May 2026" as the description.*) **China has continued to legislate vertically and has still not enacted a comprehensive AI Law** — do not assume one exists.

KEY TAKEAWAY

Reading 3 gave you two philosophies — regulate the process, or regulate the product. **These three are neither, and each is organised around a third thing.** **Singapore regulates the practice and makes it testable:** nine voluntary dimensions, plus open-source tooling — AI Verify for traditional models, Project Moonshot for language models — so that "we govern our AI responsibly" becomes a claim somebody can check. Nothing is binding; the binding layer is the PDPA underneath. **India regulates the data, and is now regulating the content:** the DPDP Act governs AI indirectly by governing the personal data it consumes, with an unusually strict consent regime — and India's only *binding* AI-specific rules to date arrived not as AI law or data law but as **intermediary law**, the IT Amendment Rules 2026, mandating labelling of synthetically generated content since February 2026. **China regulates the output and the licence to operate:** a generative AI service is treated as a public-information channel, so filing, security assessment, content labelling and value alignment are the conditions of serving the Chinese public — and China moves fastest of anyone, shipping a binding regime for AI companions and a national policy framework for AI agents, both effective 15 July 2026. **The sentence to carry into class: Singapore asks whether you can test it, India asks whether you had lawful consent for the data, and China asks who is accountable for what it says.** And the uncomfortable corollary, given what Reading 3 established: **the two regimes actually shipping binding AI rules right now are the Asian ones.**

WHY IT MATTERS FOR THIS COURSE

Capsule 2 was going to walk five regulatory regimes and does not have the time, so it walks **these three** and names the other two. That is a deliberate choice about this room: the cohort is overwhelmingly from China, India and Singapore, and these are the rules you will actually meet — in a compliance review, in a procurement questionnaire, in the conversation where somebody asks whether the model can be deployed at all. It also matters for A2 — though be precise about where it lands. **D3(b)'s ten cases are test cases and A2 fixes their shape;** what an instrument name makes far more defensible is the **argument in report section 6**, because a control someone in your organisation has heard of is a control they will fund. And there is a course-level reason: these three regimes are the strongest available evidence for a claim this course keeps making — **that the constraint on an AI system is almost never the algorithm.** Three governments looked at the same technology and regulated three different things.

Feeds Class 6: C2-A (the regulatory landscape — these three get the class time) · C2-B (China's labelling and filing duties are concrete guardrail requirements) · A2 (naming the instrument behind each control — in report section 6, not inside D3(b)'s ten cases).

MAIN POINTS

Singapore — voluntary, but with tooling, which changes what voluntary means

The **Model AI Governance Framework for Generative AI** was published on **30 May 2024** by IMDA and the AI Verify Foundation, after an international consultation on a January 2024 draft. It extends rather than replaces the 2020 Model AI Governance Framework, which covered traditional AI. Nine dimensions, and they are worth knowing in order because the order is an argument.

- **1 Accountability · 2 Data · 3 Trusted Development and Deployment · 4 Incident Reporting · 5 Testing and Assurance · 6 Security · 7 Content Provenance · 8 Safety and Alignment R&D · 9 AI for Public Good.** Note what a small country put in the list that larger frameworks leave out: **incident reporting** as a named dimension rather than a sub-clause, **content provenance** at a time when few instruments mentioned it, and **AI for Public Good** — an explicitly developmental objective sitting inside a governance framework, which is a policy position, not an oversight.
- **The tooling is the distinctive part, and it is what makes the framework different from a values statement.** **AI Verify** is an open-source testing toolkit — technical tests plus process checks — for traditional AI. **Project Moonshot**, launched 31 May 2024, is the equivalent for large language models: benchmarking and red-teaming, on GitHub, free. Singapore's bet is that governance becomes real when it becomes **testable**, and that publishing the test harness is more effective than publishing another principle. Whether the bet pays off is arguable; that it is a distinct strategy is not.
- **The agentic framework is where Singapore now is.** The *Model AI Governance Framework for Agentic AI* launched on **22 January 2026** at the World Economic Forum in Davos — billed as the first of its kind — and was substantially updated in **May 2026**. Four dimensions: **bound the risk** through appropriate use cases and limits on agent power; **human accountability** through significant checkpoints requiring approval; **technical controls** across the agent lifecycle; and **end-user responsibility** through transparency and training. **Automation bias** was named at the January launch. The May update added access controls, guardrails, logging and monitoring as core agent components, separated the accountability of platform providers from system developers, and added best practices for guarding against automation bias, plus skill degradation and operational continuity as named risks.

Read that agentic framework against your Week 4 build and it is almost a specification of what you were told to do: **bound the agent's power (tool scoping)**, **require approval at significant checkpoints (confirmation gates)**, **log everything (observability)**, and **train the user not to over-trust it (automation bias)**. A regulator and an engineering doctrine arrived at the same list independently, which is the strongest available evidence that the list is right.

- **And the part people miss: the binding layer is the PDPA.** The MGF frameworks are voluntary throughout — the agentic one is presented as a document that will be revised as the technology moves. What actually binds in Singapore is the **Personal Data Protection Act**, and the PDPC's **Advisory Guidelines on the Use of Personal Data in AI Recommendation and Decision Systems (1 March 2024)** interpret existing PDPA obligations for AI. Those guidelines have real force because the statute behind them does. Add sectoral regulation — MAS's FEAT principles for financial institutions — and the picture is a soft AI layer over a hard data-and-sector layer.

India — the data first, then the content, and the AI framework last

- **The DPDP Act 2023** received presidential assent on **11 August 2023** and is consent-centric in a way that matters for AI. It offers consent or a narrow set of "legitimate uses" and — critically — **omits the GDPR's "legitimate interests" and "contractual necessity" bases entirely**. That is a real constraint: the argument most commonly used in Europe to justify training on personal data is not available in India. It creates **Data Fiduciaries** (controllers) and **Significant Data Fiduciaries**, designated by government, who owe a India-resident Data Protection Officer, an independent auditor, periodic **Data Protection Impact Assessments** and **algorithmic due diligence**. Data Principals get access, correction, erasure, grievance redressal and nomination. Purpose limitation and data minimisation are statutory duties. **Children under 18 require verifiable parental consent, and tracking, behavioural monitoring and targeted advertising directed at children are prohibited outright** — a direct constraint on personalisation. Cross-border transfer uses a **blacklist** model: permitted everywhere except countries the government restricts, which is notably more permissive than the localisation-leaning earlier drafts.
- **Penalties are fixed caps, not turnover-based** — a structural contrast with the EU worth noticing. Failure of reasonable security safeguards causing a breach: up to **₹250 crore**. Failure to notify a breach: **₹200 crore**. Children's-data obligations: **₹200 crore**. Significant Data Fiduciary obligations: **₹150 crore**. Everything else, including consent, purpose limitation and cross-border: **₹50 crore**. A fixed cap is proportionally brutal for a mid-size Indian company and trivial for a global platform, which is the standing criticism.

- **And none of it binds yet.** Three-stage commencement from the November 2025 Rules: institutional scaffolding now, Consent Manager registration from **13 November 2026**, and the **entire substantive regime from 13 May 2027**. If you are building for India today, you are building against a regime you can read and that nobody is yet obliged to follow — which is a genuinely unusual and rather comfortable position, and it will not last.
- **NITI Aayog's Responsible AI for All** is the principles layer and predates all of it: **Part 1 (February 2021)** sets out seven principles — Safety and Reliability · Equality · Inclusivity and Non-discrimination · Privacy and Security · Transparency · Accountability · protection and reinforcement of positive human values — and **Part 2 (August 2021)** addresses operationalising them. Read it as India's statement of intent rather than as a live instrument.
- **The India AI Governance Guidelines (5 November 2025)** are the current AI-specific document, released by MeitY under the IndiaAI Mission, and they are **explicitly non-binding and explicitly anti-legislative**: India has decided *not* to create standalone AI legislation, and to layer a lightweight adaptive approach on existing law. Seven **Sutras** — Trust is the Foundation · People First · Innovation over Restraint · Fairness and Equity · Accountability · Understandable by Design · Safety, Resilience and Sustainability — and six **Pillars**: Infrastructure, Capacity Building, Policy and Regulation, Risk Mitigation, Accountability, Institutions. *Innovation over Restraint* is a deliberate signal and should be read as one.
- **The binding AI rules came from somewhere else entirely.** The **Information Technology (Intermediary Guidelines and Digital Media Ethics Code) Amendment Rules, 2026** were notified on **16 February 2026** and took effect on **20 February 2026**. They mandate labelling of **"synthetically generated information"**: visual SGI must carry labels that are prominent, easily noticeable and adequately perceivable; audio SGI requires a prominently prefixed audio disclosure. Significant social media intermediaries must deploy technical measures to verify user declarations. **The draft's proposed 10%-of-surface-area watermark was dropped** in favour of that qualitative standard. This is India's enforceable AI rule, and it arrived as content regulation.

China — the output, the filing, and the fastest legislative cadence in the world

- **The Interim Measures** were promulgated **10 July 2023**, released publicly on 13 July and effective **15 August 2023**, issued jointly by **seven agencies** — CAC, the NDRC, and the ministries of Education, Science and Technology, Industry and Information Technology, Public Security, plus the National Radio and Television Administration. Seven signatures is itself informative: this is not a technology regulator acting alone.
- **The obligations.** **Algorithm filing** with the CAC, and a **security assessment**, for services with "public opinion properties or the capacity for social mobilisation" — a phrase that does a great deal of work and captures most consumer-facing services. **Labelling** of AI-generated content. **Training-data legality**: lawful sources, respect for intellectual property, consent for personal data. **Content must adhere to Core Socialist Values** (Article 4), with prohibitions on content inciting subversion, endangering national security, terrorism, ethnic hatred, violence, obscenity, or false and harmful information. **Tiered and classified supervision** (Article 3) — "tolerant and prudent" graded management by category, pairing innovation encouragement with risk-calibrated oversight. And **extraterritorial reach**: the Measures apply to services provided to the public within mainland China, and Article 20 empowers the CAC to direct technical measures against non-compliant services provided from outside.
- **The narrowing between the April 2023 draft and the final text is the most instructive thing in the document**, and it is rarely taught. Scope was narrowed to services offered **to the public**, expressly excluding purely internal enterprise use and research and development. The training-data duty was **softened** from a duty to ensure the authenticity, accuracy, objectivity and diversity of training data to one to take effective measures to improve its quality (wording follows the China Law Translate rendering; 真实性 is variously translated "authenticity" or "truthfulness") — a change from a guarantee to an effort, and the drafters presumably worked out that the guarantee was impossible. **Dropped entirely**: the three-month deadline to prevent regeneration of illegal content; a standalone fine regime of RMB 10,000–100,000 plus service termination; and restrictions on user profiling from input data. **Added**: the extraterritorial hook and express user-side responsibilities. **A regulator engaged with the technology, discovered what was not achievable, and changed the rule.** Hold that against the common assumption that Chinese AI regulation is uniformly maximalist.
- **What came after, and this is where China outpaces everyone.** The **Measures for Labelling AI-Generated Content**, issued 14 March 2025 and **effective 1 September 2025**, require **both explicit labels** (visible or audible to the user) and **implicit labels** (metadata embedded in the file), backed by a **mandatory national standard, GB 45438-2025**. Then, effective **15 July 2026**, two more: the **Interim Measures for the Administration of Anthropomorphic AI Interaction Services** — a world-first binding regime for AI companions and emotional-interaction chatbots, requiring filing and security assessment, **banning virtual intimate relationships for minors**, parental consent for under-14s, minor modes with time limits, break reminders after two consecutive hours, and **human operator takeover when a user signals self-harm** — and — policy guidance rather than a binding regulation, and worth distinguishing — the **Implementation Opinions on Intelligent Agents** (released 8 May 2026), applying differentiated governance to AI agents in healthcare, transportation, media and public safety. The **TC260 Ethics-Safety Guidelines for AI Applications 1.0** took effect 1 July 2026 as voluntary guidance.

- **And still no comprehensive AI Law.** A scholars' draft has circulated since 2023 and the item has appeared on legislative work plans, but China continues to legislate **vertically, instrument by instrument, service type by service type**, rather than horizontally. This is a strategy rather than a delay, and it is the mirror image of the EU's: a horizontal law must define the artefact in advance and then wait for standards (Reading 3's sixteen-month slip); a vertical approach defines a service category that already exists and regulates it in months.

The worked example — one assistant, three markets, three first questions

Take a single system: **a customer-service assistant, built on a rented foundation model over your company's own support documentation, answering customer questions in a consumer app.** The Week-2 and Week-3 build, essentially. Deploy it in all three markets and watch the first question change.

1. **Singapore asks: can you show it was tested?** Nothing in the MGF obliges you to do anything. But the PDPA binds, so the first hard question is the data one: what personal data is in the retrieval corpus and the conversation logs, is there a lawful basis, and have you met the PDPC's March 2024 guidance on AI recommendation and decision systems? Then the soft-but-real question, the one an enterprise buyer will ask in procurement: have you run it through **Project Moonshot** or an equivalent, and what did it find? In Singapore the framework is voluntary and the *market* enforces it, which is a meaningfully different mechanism from a regulator enforcing it — weaker in some ways, faster in others.
2. **India asks: what is your lawful basis, and did you label the output?** Today, in September 2026, the DPDP answer is uncomfortable and honest: **the substantive obligations do not commence until 13 May 2027**, so the consent regime is a design target rather than a duty. But design for it now, because retrofitting consent is close to impossible — you cannot go back and obtain a basis for data you already trained on, and the absence of a legitimate-interests ground means the European workaround is unavailable. Meanwhile the **binding** obligation today is the **IT Amendment Rules 2026**: if your assistant produces synthetic content, it must be labelled to the qualitative standard. Notice the inversion — the future obligation is about data and the present one is about output.
3. **China asks: who is accountable for what it says, and have you filed?** This is the market where the first question is a permission question, and the sequence matters. If the service has public-opinion properties or the capacity for social mobilisation — which a consumer-facing assistant does — you need **algorithm filing** and a **security assessment** before launch, not after. Generated content must be **labelled explicitly and implicitly** to GB 45438-2025. Content must align with the Article 4 requirements, which means a content-moderation layer that has nothing to do with model quality and everything to do with what is permitted. Training data must be lawfully sourced. And if any part of the experience is companion-like or emotionally interactive, the July 2026 anthropomorphic-services regime engages with a further set of duties including minor protection and self-harm escalation. **This is the most operationally demanding of the three by a wide margin, and it is also the clearest — you know what is required and when.**

The transferable observation: the three markets do not disagree about whether AI should be governed. They disagree about what the governable object is. Singapore points at the practice, India at the data, China at the output. A team that builds for one and assumes the others are a subset will be wrong in a different way each time — and the compliance work that transfers between them is smaller than you would hope, while the *engineering* work (logging, provenance, evaluation, content filtering) transfers almost entirely.

WHERE TO BE SCEPTICAL

Four cautions, and the first is about a story this cohort will hear repeatedly and should be able to correct.

1. **"EU strict, US loose, Asia lagging" is out of date, and you should stop repeating it.** As of September 2026 the EU has postponed its high-risk regime by sixteen months, the US has revoked its AI executive order and renamed its safety institute, while **China has two binding AI instruments that took effect within the last twelve months — plus a national policy framework for AI agents — and India has had mandatory synthetic-content labelling since February.** That is not lagging. The framing survives because it was true in 2023 and because it flatters its authors, and it will be repeated to you in decks all year.
2. **A voluntary framework with excellent tooling is still voluntary.** Singapore's approach is genuinely admirable and it rests on a bet: that market pressure, reputational incentives and a small, well-connected business community will do the enforcement work. That bet is plausible in Singapore and does not obviously travel — a jurisdiction with a large fragmented market and weak procurement discipline gets the framework and none of the enforcement. Be careful about generalising from Singapore's success to "voluntary frameworks work".
3. **Legislating fast is not the same as legislating well, and the anthropomorphic-services rules are the test case.** China shipped a binding regime for AI companions faster than anyone else could convene a working group, and it addresses real and documented harms to minors. It also regulates an intimate category of speech and relationship, with filing requirements, and reasonable people will assess that differently depending on what else they believe about the regulator. **Both observations are true and this brief is not going to resolve them for you**

— but notice that “they moved fast” and “they moved well” are separate claims, and that the same is true of every regime in these two readings.

4. **Enforcement records are thin everywhere, and compliance advice is currently a forecast.** China has the longest operating history and the most opaque enforcement; Singapore has nothing to enforce; India’s regime has not commenced. So almost everything written about how these rules will be applied in practice — including this brief — is inference from text rather than from decisions. **Where a rule matters commercially, get local counsel, and treat any confident statement about enforcement posture, including a confident statement in a course reading, as a hypothesis.**

READING FOCUS QUESTIONS — bring your answers to class

There are no wrong answers here — these surface how you think, not what you recall. A clue sits under each question to give you a direction to push on; the fuller reasoning is in the section after.

- Singapore points at the practice, India at the data, China at the output. For each, work out what that regime catches that the other two miss — and what it misses that they catch.

Clue. Do it with one concrete failure rather than in the abstract: take a system that confidently tells a customer something false and damaging, and ask which regime has anything to say about it. Then swap the failure — a system trained on data nobody consented to, a system that leaks between customers, a system that is accurate and unfair — and run the same test. You should find that each regime is silent on at least one failure that another catches squarely, and the pattern of silences is the answer. Notice which failures **all three** miss; that gap is the interesting finding.

- India’s DPDP substantive provisions do not commence until May 2027, and its binding AI rule today is a content-labelling obligation that arrived through intermediary law. What does that sequence tell you about how AI actually gets regulated?

Clue. Ask why content law moved first, and the answer is probably about what was politically urgent rather than what was intellectually prior — look at what was happening with synthetic media in India in 2025–26. Then connect it to Reading 3’s finding that US employers face binding AI obligations through employment law rather than AI law. The generalisation you are reaching for is about which body of law is already in place and already has enforcement machinery, and what that predicts about where the next binding AI rule in any jurisdiction will come from.

- China’s final Interim Measures softened the training-data duty from “ensure truthfulness and accuracy” to “take effective measures to improve” data quality. Why — and what does the change tell you about regulating a technology you are still learning?

Clue. Ask what it would actually take to comply with the original wording, for a corpus of the size a foundation model needs, and how you would demonstrate compliance to an inspector. Once you have concluded it is impossible, ask the more useful question: what happens to a regulation that mandates the impossible — is it enforced, ignored, or enforced selectively, and which of those is worst? Then look at the other things dropped between draft and final and ask whether they share a property. This is a rare documented case of a regulator learning in public, so treat it as evidence rather than as a footnote.

- Singapore’s agentic framework and Week 4’s engineering doctrine arrived at nearly the same list — bound the power, gate the significant actions, log everything, train the user. What should you make of two independent routes to one answer?

Clue. First check that they really are independent, and that one is not simply reading the other — look at the dates and at who wrote each. If they are, the convergence is evidence about the problem rather than about either author, and the question becomes what property of agents forces this particular list. Try to derive it yourself from Reading 1’s grid: what does an irreversible action require, and what does an invisible one require? Then ask the sceptical version — whether a list that both an engineer and a regulator find obvious is a strong result or merely the easy part, and what is on neither list.

- Pick the market you know best. What would you have to change about your A2 system to deploy it there — and which of those changes are compliance work, and which are engineering you should have done anyway?

Clue. Sort every requirement you find into three buckets and be strict about the boundaries: things that exist only to satisfy a regulator (a filing, a registration, a published notice), things a regulator requires that also make the system better (logging, provenance, evaluation, content filtering), and things you would want even if nobody was watching. The middle bucket is usually much the largest and it is the interesting finding, because it changes how you argue for the work internally. Then ask which bucket the effort in your own team is currently going into.

HOW TO THINK ABOUT THESE — read after you have had a go

These are not model answers to be memorised — they are the line of reasoning we will pick up in class. Class time is short, so they are written to stand on their own.

Q1 — Each regime is blind to a different failure, and all three are blind to one

Run four failures through all three and the pattern is clean. **A confidently false and damaging answer.** China catches this most directly, because it regulates output: false and harmful information is named in Article 4 and the service provider is accountable for what the system says, full stop. Singapore reaches it obliquely through Testing and Assurance and through Incident Reporting, but obliges nothing. India, under DPDP, has nothing to say at all — no personal data is implicated by a wrong answer about a refund policy — though the IT Rules bite if the output is synthetic media. **Training on data nobody consented to.** India is squarely on this and is the strictest of the three, precisely because it removed the legitimate-interests basis that the rest of the world leans on. China addresses it through the training-data legality provision, though softened. Singapore reaches it through the PDPA rather than through the AI framework. **Cross-tenant leakage in a shared vector store.** Reading 2's LLM09. India catches it as a security-safeguards failure with the largest penalty in the Act — ₹250 crore. Singapore's Security and Data dimensions cover it, non-bindingly, and the PDPA behind them does bind. China covers it through the broader cybersecurity and personal-information regime rather than through the Interim Measures. **A system that is accurate and unfair** — the CV screener from Reading 3, ranking well and disadvantaging a group. And here is the finding worth bringing to class: **none of the three catches it squarely.** India's DPDP is a data-protection statute and says nothing about the fairness of a lawful inference. Singapore names fairness in its principles and obliges nothing, though the PDPC guidance and MAS FEAT reach parts of it. China's regime is oriented to content and public order, not to distributive outcomes. **The EU AI Act would catch it — that is exactly what the Annex III high-risk regime is for — and it does not bind until December 2027.** So the honest summary is that in September 2026, across the three markets this cohort will mostly work in, an accurate and discriminatory automated decision system faces no AI-specific obligation anywhere, and is governed only by whatever sector or anti-discrimination law happens to exist. That is a real gap, it is not an accident — it reflects what each government was most worried about — and it is the strongest available argument for Reading 3's point that you should start from the activity rather than from the AI framework.

Q2 — Binding AI rules come from whichever body of law already has enforcement machinery

The Indian sequence is almost a controlled experiment, so it is worth stating precisely. Parliament passed a comprehensive data-protection statute in August 2023 and its substantive provisions will not bind until May 2027 — nearly four years. MeitY published AI governance guidelines in November 2025 and made them explicitly non-binding, and explicitly declined to legislate. And in February 2026, binding rules on AI-generated content took effect in **four days** from notification. The difference is not importance and it is not drafting skill. It is that the IT Rules amend an existing instrument under an existing statute with an existing regulator, existing intermediaries who are already subject to compliance obligations, and existing takedown machinery. **Nothing had to be built. The rule was a modification to a control surface that already worked.** The DPDP Act, by contrast, requires a Data Protection Board that did not exist, Consent Managers that did not exist, a registration system, and a compliance capability inside every organisation in India — which is why it needs three stages and four years. Generalise it and you get a decent predictive rule: **the next binding AI obligation in any jurisdiction will attach to whichever regulator already has a working enforcement path to the people you want to bind.** Reading 3 gave you the American version — employment law, through the EEOC and a municipal audit requirement, while federal AI legislation stalls. The European version is the same story from the other end: the AI Act had to build notified bodies, harmonised standards and a new AI Office, and the parts that needed none of that (the Article 5 prohibitions, which are conduct rules) commenced on time while the parts that did have slipped sixteen months. The Chinese version is the strongest confirmation of all — the CAC regulates through **filing and security assessment**, machinery that already existed for algorithmic recommendation systems and was simply pointed at a new service category, which is exactly why China can ship a binding regime for AI companions in under four months from draft to promulgation. **For you, the practical consequence is a change of habit.** When you want to know what will bind your system next, do not read AI policy commentary. Ask which regulator already has a relationship with your organisation, what they already require, and what the cheapest extension of that requirement to cover AI would look like — because that is what will arrive, and it will arrive years before anything called an AI Act.

Q3 — Because the original was impossible, and a rule that mandates the impossible is worse than no rule

Start by trying to comply with the draft language. A duty to ensure the authenticity, accuracy, objectivity and diversity of training data, for a corpus in the trillions of tokens scraped from the open web. **Truthfulness** would require adjudicating the truth of every proposition in the corpus, including contested ones. **Accuracy** raises the same problem with a date attached, since a true statement from 2019 may be false now. **Objectivity** is not a property a document has. **Diversity** requires a target distribution nobody has specified. There is no procedure, no measurement and therefore no possible demonstration of compliance — and the regulator, having engaged with practitioners between April and July 2023, evidently worked this out. The revision from a **duty of result** to a **duty of effort** — “take effective measures to improve” — is precisely the right repair, because it is auditable: you can show your filtering pipeline, your deduplication, your provenance records, your removal process. The other things dropped share the property. The three-month deadline to prevent regeneration of illegal content assumed you could reliably remove a capability from a trained model in a fixed window, which the machine-unlearning literature says you cannot. The specific fine schedule was dropped in favour of existing law, which is ordinary legislative tidiness. The

scope narrowing to public-facing services removed research and internal enterprise use, which is a judgement that the risk lives in the public channel. **The general lesson has two halves and the second is the one people skip.** First: a regulation that mandates the impossible does not produce compliance, it produces one of three worse outcomes — universal technical non-compliance that everyone ignores, which corrodes the whole instrument; selective enforcement, which becomes an arbitrary instrument; or theatrical compliance, where firms produce documents attesting to something nobody has done. All three are worse than a weaker rule that is actually met. Second, and this is the part worth arguing about in class: **a regulator who softens a requirement after consultation is usually doing their job, and it will be reported as capture.** You cannot distinguish learning from lobbying by looking at the direction of the change; you have to look at whether the original was achievable. Here it demonstrably was not. Apply the same test to the EU's Omnibus delay from Reading 3 and you get the same answer for the same reason — the harmonised standards did not exist, so conformity could not be demonstrated — and the same accusation was made. **The test to carry: when a rule is softened, ask what compliance with the original version would have required, concretely, and whether anyone could have produced it.** That question distinguishes the two cases far better than motive-reading does.

Q4 — The convergence is real, and what is on neither list is more interesting

First establish the independence, because the claim is worthless otherwise. Anthropic's agent doctrine is engineering guidance from December 2024 with follow-ups through 2026, written by people building agents. Singapore's agentic framework is public policy from January 2026, updated May 2026, written by a regulator and a foundation after industry consultation. The May update added automation bias, skill degradation and operational continuity, which are human-factors risks an engineering doctrine would not naturally produce — so the influence is not one-directional, and there is real independent content on the policy side. Granting that, the convergence is evidence about the **problem**, and you can derive the list yourself from Reading 1's grid, which is the useful exercise. An agent takes **irreversible** actions, and Reading 1 says the control for irreversibility is a gate placed before the action — hence confirmation at significant checkpoints, on both lists. An agent's reasoning is **invisible**, so the control is monitoring — hence logging as a core component, on both lists. An agent **compounds** error over steps, so the control is bounding the blast radius — hence limits on agent power and tool scope, on both lists. And a human supervising an agent is subject to **automation bias**, so the control is training and design that resists it — which is on the policy list and, notably, not on most engineering ones. So the list is not a coincidence and not a fashion; it is what the failure structure of an agent forces. **That is genuinely reassuring**, because it means an engineering team doing careful work will land close to the regulatory requirement without reading the regulation, and it means the requirement is unlikely to be reversed by the next framework. Now the sceptical half, which is where the class discussion should go. **The convergence is on the easy part.** Both lists tell you to gate significant actions and neither tells you what "significant" means — and that is where the entire cost sits, because a gate on everything is not an agent and a gate on nothing is not a control. Both say bound the agent's power and neither says how much power the task needs. Both say log, and neither says what to do with the logs, which is the difference between observability and storage. Reading 1 named the deeper gap: **nobody has a good method for evaluating an agent's trajectory**, so neither list can tell you whether your agent is safe — only whether you have put the recognised controls around it. **That distinction is worth landing hard, because it generalises to every governance framework in these two readings.** A framework can tell you that you have done the recognised things. It cannot tell you that your system is safe, and the gap between those two claims is where incidents live.

Q5 — Most of it is engineering you should have done anyway, and that changes how you argue for it

Sorted honestly, the three buckets come out very lopsided, and discovering that is the point of the exercise. **Pure compliance overhead** is a short list: algorithm filing and security assessment in China, registration and published notices, appointing a Data Protection Officer resident in India, a specific consent-collection flow, an independent audit. Real cost, real calendar time — filing in particular is a launch dependency rather than a task, which teams discover late — and no engineering benefit. **Required and genuinely useful** is much the longest list and it is where the argument gets won: logging of inputs, outputs and tool calls; content provenance and labelling; evaluation and red-teaming before release; access controls and tenant isolation; incident detection and reporting; data lineage; human review at consequential points; retention limits. Every one of those is something Readings 1 and 2 told you to build for reasons that had nothing to do with a regulator — logging because agent failures are invisible until you look, evaluation because you cannot manage what you have not measured, gates because irreversible actions cannot be undone, provenance because Reading 2's indirect injection is a provenance failure at heart. **Wanted regardless** is short only because the middle bucket absorbed most of it. The strategic consequence is the reason to do this sorting before you write your A2 report rather than after. Compliance work is usually argued for defensively — we must, or we cannot ship — and defensive arguments get minimum budget and a checkbox implementation, which produces the worst of both worlds: cost incurred, benefit missed. **If you can show a stakeholder that most of the regulatory requirement is the same work as making the system reliable, the conversation changes from cost to sequencing**, and you get to build it properly. The concrete version for A2, and be exact about which artefact takes it. **D3(b)'s ten are test cases, and A2 fixes their shape: each names the wrong behaviour it exists to catch and states the observed result.** An instrument name is not part of that, and a case citing a regulation instead of a result does not satisfy the deliverable. **Where the mapping belongs is report section 6 and the trade-off argument** — the

place you say why each control exists at all. Written there it reads like this: *we log every tool call — it prevents undetected agent failure (Reading 1), and it is required by Singapore’s agentic framework and by NIST MANAGE. We label generated content — it prevents downstream over-trust, and it is required by China’s GB 45438-2025 and India’s IT Amendment Rules 2026.* Reading 3’s Rubric-1 point applies here too: an argument that explains why each control exists is a trade-off analysis, and one that only cites regulations is a compliance artefact. If you want the instrument as a third column on the checklist as well, nothing stops you — just never let it displace the observed result. **And notice the last thing this exercise reveals, which is usually where the effort actually is going.** Ask which bucket your team is currently spending on. Teams under regulatory pressure tend to over-invest in bucket one, because it is legible and has a deadline, while bucket two — the work that would actually prevent the incident — has no deadline and no external prompt. That is precisely the failure Reading 3 described from the other direction: **an exemplary governance programme wrapped around a system nobody has validated.**

APPENDIX — VOCABULARY, EXPLAINED

The distinction that carries this reading: what each regime treats as the governable object — the practice, the data, or the output. It is not a difference of strictness; it is a difference of target, and it decides which failures each regime can even see. Every term below belongs to one of the three.

Algorithm filing and security assessment (China)

- **Plain.** Registering your service with the regulator, and passing a review, before you launch it to the public.
Technically. Required under the Interim Measures for generative AI services with “**public opinion properties or the capacity for social mobilisation**” — in practice most consumer-facing services. Runs through the existing algorithm-recommendation filing system, which is why it could be extended to generative AI quickly. **A launch dependency, not a post-launch task,** and the single most commonly underestimated item in a China deployment plan.

Explicit and implicit labelling (China)

- **Plain.** A label the user can see or hear, and a marker inside the file that machines can read.
Technically. Measures for Labelling AI-Generated Content, issued 14 March 2025, **effective 1 September 2025**, with mandatory national standard **GB 45438-2025**. Both layers are required. Compare India’s IT Amendment Rules 2026, which mandate a **qualitative** standard — “prominent, easily noticeable and adequately perceivable” — after dropping a proposed 10%-of-surface-area watermark. Related to the EU AI Act’s Article 50 machine-readable marking; three regimes, three specifications, one engineering problem.

Data Fiduciary and Significant Data Fiduciary (India)

- **Plain.** Anyone deciding why and how personal data is processed; and the larger or riskier ones, who owe more.
Technically. DPDP Act 2023. A **Data Fiduciary** is broadly the GDPR controller. **Significant Data Fiduciaries** are designated by government on volume, sensitivity and risk criteria, and additionally owe an India-resident Data Protection Officer, an independent data auditor, periodic **Data Protection Impact Assessments** and **algorithmic due diligence**. Substantive obligations commence **13 May 2027**.

Consent, and the missing lawful bases (India)

- **Plain.** In India you generally need consent. The two arguments Europe uses to avoid needing it do not exist.
Technically. The DPDP Act permits processing on consent or a narrow set of “**legitimate uses**” (employment, legal obligation, emergency, State functions). It **omits the GDPR’s “legitimate interests” and “contractual necessity” bases**, which is the constraint that matters most for AI training on personal data. Children under 18 need verifiable parental consent, and tracking, behavioural monitoring and targeted advertising to children are prohibited outright. Penalties are **fixed caps** (to ₹250 crore), not turnover-based.

The nine dimensions (Singapore, generative AI)

- **Plain.** Singapore’s checklist for governing a generative AI deployment. Voluntary, and paired with free tools to test against it.
Technically. Accountability • Data • Trusted Development and Deployment • Incident Reporting • Testing and Assurance • Security • Content Provenance • Safety and Alignment R&D • AI for Public Good. Published **30 May 2024** by IMDA and the AI Verify Foundation; extends the 2020 Model AI Governance Framework rather than replacing it.

The four dimensions (Singapore, agentic AI)

- **Plain.** The same idea for agents: bound what it can do, keep a human at the checkpoints, control the lifecycle, and train the user.
Technically. Launched **22 January 2026** at Davos, updated **May 2026**. Risk assessment and bounding via appropriate use cases and agent power limits • human accountability through significant checkpoints requiring approval • technical controls across the agent lifecycle • end-user responsibility through transparency and training. **Automation bias** was named at the January 2026 launch. The May update added access controls, guardrails, logging and monitoring as core components, separated platform-provider from system-developer accountability, and added best practices against automation bias plus skill degradation and operational continuity.

AI Verify and Project Moonshot

- **Plain.** Free, open-source test harnesses — one for traditional models, one for language models.
Technically. **AI Verify:** technical tests plus process checks for traditional AI, with published crosswalks to ISO/IEC 42001 and the NIST AI RMF. **Project Moonshot** (launched 31 May 2024): benchmarking and red-teaming for large language models. Both governed by the AI Verify Foundation. **The strategy is that governance becomes real when it becomes testable** — publish the harness, not another principle.

Core Socialist Values requirement (China)

- **Plain.** What the system says has to be consistent with a specified set of values, and the provider is accountable for it.
Technically. Article 4 of the Interim Measures, with prohibitions on content inciting subversion, endangering national security, terrorism, ethnic hatred, violence, obscenity, or **false and harmful information**. Operationally this means a content-moderation layer independent of model quality. It is the clearest instance of a regime treating a generative service as a **publishing channel** rather than as a product or a process.

Tiered and classified supervision (China)

- **Plain.** Different rules for different kinds of service, rather than one rule for all AI.
Technically. Article 3 — “tolerant and prudent” graded management by category. The organising principle behind China’s **vertical** legislative strategy: regulate service categories that already exist (recommendation algorithms, deep synthesis, generative AI, anthropomorphic interaction, intelligent agents) instead of defining AI horizontally. The mirror image of the EU’s approach, and the reason China can ship a binding regime in months while a horizontal law waits on standards.

SOURCE LINKS

Singapore — Model AI Governance Framework for Generative AI (30 May 2024) · aiverifyfoundation.sg/resources/mgf-gen-ai/

Singapore — Model AI Governance Framework for Agentic AI (Jan 2026, updated May 2026) ·

www.imda.gov.sg/-/media/imda/files/about/emerging-tech-and-research/artificial-intelligence/mgf-for-agentic-ai.pdf

Singapore — Project Moonshot, the open-source LLM testing toolkit · aiverifyfoundation.sg/project-moonshot/

India — DPDP Rules 2025 notification and phased commencement (PIB) · www.pib.gov.in/PressReleasePage.aspx?PRID=2190014®=3&lang=2

India — AI Governance Guidelines, MeitY (5 November 2025) · www.pib.gov.in/PressReleasePage.aspx?PRID=2186639

India — NITI Aayog, Responsible AI for All, Part 2 (August 2021) · www.niti.gov.in/sites/default/files/2023-02/Part2-Responsible-AI-12082021.pdf

China — Interim Measures for the Management of Generative AI Services (English translation) · www.chinalawtranslate.com/en/generative-ai-interim/

China — draft vs final comparison of the Interim Measures (Future of Privacy Forum) · fpf.org/blog/chinas-interim-measures-for-the-management-of-generative-ai-services-a-comparison-between-the-final-and-draft-versions-of-the-text/

- **Plain.** Related but not assigned: Singapore’s PDPC Advisory Guidelines on the Use of Personal Data in AI Recommendation and Decision Systems (1 March 2024) — the binding layer beneath the voluntary frameworks; MAS FEAT principles for financial institutions; China’s Measures for Labelling AI-Generated Content (effective 1 September 2025) and national standard GB 45438-2025; India’s IT (Intermediary Guidelines and Digital Media Ethics Code) Amendment Rules 2026 (in force 20 February 2026). Reading 3 covers NIST and the EU AI Act.

CAVEATS

1. **This summary is not a substitute for the frameworks, and two of the three are moving fast.** China’s cadence in particular is the fastest of any major jurisdiction; a rule described here as current may have been superseded between writing and reading.
2. **Chinese instruments are read in translation.** Titles, thresholds and the Core Socialist Values requirement are rendered from official Chinese-language texts and widely used English translations. Where a distinction turns on a single word, verify against the original before relying on it — several people in this room can do that better than this brief can.
3. **Commencement dates are the part most likely to be wrong somewhere.** India’s DPDP stages and the IT Amendment Rules dates are stated as published; staged commencement has slipped before in this jurisdiction and may slip again. **Check the date before you put it in a plan.**
4. **Singapore’s frameworks are voluntary and the PDPA is not.** Where this brief says an obligation “binds” in Singapore it means the PDPA or a sectoral regulator, never the IMDA framework — the distinction is the entire teaching point and it is easy to blur when summarising.
5. **Generated and fact-checked using Anthropic models.** The three-first-questions comparison, the governable-object framing and the three-bucket sorting exercise are the instructor’s. **Flag anything that looks wrong.**